

C interfaces to GALAHAD NLS

Generated by Doxygen 1.8.17

1 GALAHAD C package nls	1
1.1 Introduction	1
1.1.1 Purpose	1
1.1.2 Authors	1
1.1.3 Originally released	1
1.1.4 Terminology	2
1.1.5 Method	2
1.1.6 Reference	3
1.1.7 Call order	3
1.1.8 Unsymmetric matrix storage formats	4
1.1.8.1 Dense storage format	4
1.1.8.2 Dense by columns storage format	4
1.1.8.3 Sparse co-ordinate storage format	4
1.1.8.4 Sparse row-wise storage format	4
1.1.8.5 Sparse column-wise storage format	5
1.1.9 Symmetric matrix storage formats	5
1.1.9.1 Dense storage format	5
1.1.9.2 Sparse co-ordinate storage format	5
1.1.9.3 Sparse row-wise storage format	5
1.1.9.4 Diagonal storage format	5
1.1.9.5 Multiples of the identity storage format	5
1.1.9.6 The identity matrix format	6
1.1.9.7 The zero matrix format	6
2 File Index	7
2.1 File List	7
3 File Documentation	9
3.1 nls.h File Reference	9
3.1.1 Data Structure Documentation	10
3.1.1.1 struct nls_subproblem_control_type	10
3.1.1.2 struct nls_control_type	12
3.1.1.3 struct nls_time_type	17
3.1.1.4 struct nls_subproblem_inform_type	17
3.1.1.5 struct nls_inform_type	18
3.1.2 Function Documentation	19
3.1.2.1 nls_import()	19
3.1.2.2 nls_information()	21
3.1.2.3 nls_initialize()	22
3.1.2.4 nls_read_specfile()	22
3.1.2.5 nls_reset_control()	22
3.1.2.6 nls_solve_reverse_with_mat()	23
3.1.2.7 nls_solve_reverse_without_mat()	26

3.1.2.8 nls_solve_with_mat()	28
3.1.2.9 nls_solve_without_mat()	31
3.1.2.10 nls_terminate()	34
4 Example Documentation	37
4.1 nlst.c	37
4.2 nlstf.c	44
Index	53

Chapter 1

GALAHAD C package nls

1.1 Introduction

1.1.1 Purpose

This package uses a **regularization method to find a (local) unconstrained minimizer of a differentiable weighted sum-of-squares objective function**

$$f(\mathbf{x}) := \frac{1}{2} \sum_{i=1}^m w_i c_i^2(\mathbf{x}) \equiv \frac{1}{2} \|\mathbf{c}(\mathbf{x})\|_{\mathbf{W}}^2$$

of many variables \mathbf{x} involving positive weights w_i , $i = 1, \dots, m$. The method offers the choice of direct and iterative solution of the key regularization subproblems, and is most suitable for large problems. First derivatives of the *residual function* $c(x)$ are required, and if second derivatives of the $c_i(x)$ can be calculated, they may be exploited—if suitable products of the first or second derivatives with a vector may be found but not the derivatives themselves, that can also be used to advantage.

1.1.2 Authors

N. I. M. Gould, STFC-Rutherford Appleton Laboratory, England.

C interface, additionally J. Fowkes, STFC-Rutherford Appleton Laboratory.

1.1.3 Originally released

October 2016, C interface August 2021.

1.1.4 Terminology

The *gradient* $\nabla_x f(x)$ of a function $f(x)$ is the vector whose i -th component is $\partial f(x)/\partial x_i$. The *Hessian* $\nabla_{xx} f(x)$ of $f(x)$ is the symmetric matrix whose i, j -th entry is $\partial^2 f(x)/\partial x_i \partial x_j$. The Hessian is *sparse* if a significant and useful proportion of the entries are universally zero.

The algorithm used by the package is iterative. From the current best estimate of the minimizer x_k , a trial improved point $x_k + s_k$ is sought. The correction s_k is chosen to improve a model $m_k(s)$ of the stabilised objective function $f_{\rho,p}(x_k + s)$ built around the objective function $f(x_k + s)$ built around x_k . The model is the sum of two basic components, a suitable approximation $t_k(s)$ of $f(x_k + s)$, another approximation of $(\rho/r)\|x_k + s\|_r^r$ (if $\rho > 0$), and a regularization term $(\sigma_k/p)\|s\|_{S_k}^p$ involving a weight σ_k , power p and a norm $\|s\|_{S_k} := \sqrt{s^T S_k s}$ for a given positive definite scaling matrix S_k that is included to prevent large corrections. The weight σ_k is adjusted as the algorithm progresses to ensure convergence.

The model $t_k(s)$ is a truncated Taylor-series approximation, and this relies on being able to compute or estimate derivatives of $c(x)$. Various models are provided, and each has different derivative requirements. We denote the m by n *residual Jacobian* $J(x) \equiv \nabla_x c(x)$ as the matrix whose i, j -th component

$$J(x)_{i,j} := \partial c_i(x)/\partial x_j \text{ for } i = 1, \dots, m \text{ and } j = 1, \dots, n.$$

For a given m -vector y , the *weighted-residual Hessian* is the sum

$$H(x, y) := \sum_{\ell=1}^m y_\ell H_\ell(x), \text{ where } H_\ell(x)_{i,j} := \partial^2 c_\ell(x)/\partial x_i \partial x_j \text{ for } i, j = 1, \dots, n$$

is the Hessian of $c_\ell(x)$. Finally, for a given vector v , we define the *residual-Hessians-vector product matrix*

$$P(x, v) := (H_1(x)v, \dots, H_m(x)v).$$

The models $t_k(s)$ provided are,

1. the first-order Taylor approximation $f(x_k) + g(x_k)^T s$, where $g(x) = J^T(x)Wc(x)$,
2. a barely second-order approximation $f(x_k) + g(x_k)^T s + \frac{1}{2}s^T Ws$,
3. the Gauss-Newton approximation $\frac{1}{2}\|c(x_k) + J(x_k)s\|_W^2$,
4. the Newton (second-order Taylor) approximation $f(x_k) + g(x_k)^T s + \frac{1}{2}s^T [J^T(x_k)WJ(x_k) + H(x_k, Wc(x_k))]s$, and
5. the tensor Gauss-Newton approximation $\frac{1}{2}\|c(x_k) + J(x_k)s + \frac{1}{2}s^T \cdot P(x_k, s)\|_W^2$, where the i -th component of $s^T \cdot P(x_k, s)$ is shorthand for the scalar $s^T H_i(x_k)s$, where W is the diagonal matrix of weights w_i , $i = 1, \dots, m$.

Access to a particular model requires that the user is either able to provide the derivatives needed ('*matrix available*') or that the products of these derivatives (and their transposes) with specified vectors are possible ('*matrix free*').

1.1.5 Method

An adaptive regularization method is used. In this, an improvement to a current estimate of the required minimizer, x_k is sought by computing a step s_k . The step is chosen to approximately minimize a model $t_k(s)$ of $f_{\rho,p}(x_k + s)$ that includes a weighted regularization term $(\sigma_k/p)\|s\|_{S_k}^p$ for some specified positive weight σ_k . The quality of the resulting step s_k is assessed by computing the "ratio" $\% (f_{\rho,p}(x_k) - f_{\rho,p}(x_k + s_k))/(t_k(0) - t_k(s_k))$. $(f(x_k) - f(x_k + s_k))/(t_k(0) - t_k(s_k))$. The step is deemed to have succeeded if the ratio exceeds a given $\eta_s > 0$, and in this case $x_{k+1} = x_k + s_k$. Otherwise $x_{k+1} = x_k$, and the weight is increased by powers of a given increase factor up to a given limit. If the ratio is larger than $\eta_v \geq \eta_d$, the weight will be decreased by powers of a given

decrease factor again up to a given limit. The method will terminate as soon as $f(x_k)$ or $\|\nabla_x f(x_k)\|$ is smaller than a specified value.

A choice of linear, quadratic or quartic models $t_k(s)$ is available (see the [Terminology](#) section), and normally a two-norm regularization will be used, but this may change if preconditioning is employed.

If linear or quadratic models are employed, an appropriate, approximate model minimizer is found using either a direct approach involving factorization of a shift of the model Hessian B_k or an iterative (conjugate-gradient/Lanczos) approach based on approximations to the required solution from a so-called Krylov subspace. The direct approach is based on the knowledge that the required solution satisfies the linear system of equations $(B_k + \lambda_k I)s_k = -\nabla_x f(x_k)$ involving a scalar Lagrange multiplier λ_k . This multiplier is found by uni-variate root finding, using a safeguarded Newton-like process, by the GALAHAD packages RQS. The iterative approach uses the GALAHAD package GLRT, and is best accelerated by preconditioning with good approximations to the Hessian of the model using GALAHAD's PSLS. The iterative approach has the advantage that only Hessian matrix-vector products are required, and thus the Hessian B_k is not required explicitly. However when factorizations of the Hessian are possible, the direct approach is often more efficient.

When a quartic model is used, the model is itself of least-squares form, and the package calls itself recursively to approximately minimize its model. The quartic model often gives a better approximation, but at the cost of more involved derivative requirements.

1.1.6 Reference

The generic adaptive cubic regularization method is described in detail in

C. Cartis, N. I. M. Gould and Ph. L. Toint, "Adaptive cubic regularisation methods for unconstrained optimization. Part I: motivation, convergence and numerical results", *Mathematical Programming* 127(2) (2011) 245-295,

and uses "tricks" as suggested in

N. I. M. Gould, M. Porcelli and Ph. L. Toint, "Updating the regularization parameter in the adaptive cubic regularization algorithm". *Computational Optimization and Applications* 53(1) (2012) 1-22.

The specific methods employed here are discussed in

N. I. M. Gould, J. A. Scott and T. Rees, "Convergence and evaluation-complexity analysis of a regularized tensor-Newton method for solving nonlinear least-squares problems". *Computational Optimization and Applications* 73(1) (2019) 1–35.

1.1.7 Call order

To solve a given problem, functions from the nls package must be called in the following order:

- [nls_initialize](#) - provide default control parameters and set up initial data structures
- [nls_read_specfile](#) (optional) - override control values by reading replacement values from a file
- [nls_import](#) - set up problem data structures and fixed values
- [nls_reset_control](#) (optional) - possibly change control parameters if a sequence of problems are being solved
- solve the problem by calling one of
 - [nls_solve_with_mat](#) - solve using function calls to evaluate function, gradient and Hessian values
 - [nls_solve_without_mat](#) - solve using function calls to evaluate function and gradient values and Hessian-vector products

- `nls_solve_reverse_with_mat` - solve returning to the calling program to obtain function, gradient and Hessian values, or
- `nls_solve_reverse_without_mat` - solve returning to the calling program to obtain function and gradient values and Hessian-vector products
- `nls_information` (optional) - recover information about the solution and solution process
- `nls_terminate` - deallocate data structures

See Section 4.1 for examples of use.

1.1.8 Unsymmetric matrix storage formats

The unsymmetric m by n Jacobian matrix $J \equiv \nabla_x c(x)$ and the residual-Hessians-vector product matrix $SP(x,v)$ may be presented and stored in a variety of convenient input formats. Let A be J or P as appropriate.

Both C-style (0 based) and fortran-style (1-based) indexing is allowed. Choose `control.f_indexing` as `false` for C style and `true` for fortran style; the discussion below presumes C style, but add 1 to indices for the corresponding fortran version.

Wrappers will automatically convert between 0-based (C) and 1-based (fortran) array indexing, so may be used transparently from C. This conversion involves both time and memory overheads that may be avoided by supplying data that is already stored using 1-based indexing.

1.1.8.1 Dense storage format

The matrix A is stored as a compact dense matrix by rows, that is, the values of the entries of each row in turn are stored in order within an appropriate real one-dimensional array. In this case, component $n * i + j$ of the storage array `A_val` will hold the value A_{ij} for $0 \leq i \leq m - 1$, $0 \leq j \leq n - 1$.

1.1.8.2 Dense by columns storage format

The matrix A is stored as a compact dense matrix by columns, that is, the values of the entries of each column in turn are stored in order within an appropriate real one-dimensional array. In this case, component $m * j + i$ of the storage array `A_val` will hold the value A_{ij} for $0 \leq i \leq m - 1$, $0 \leq j \leq n - 1$.

1.1.8.3 Sparse co-ordinate storage format

Only the nonzero entries of the matrices are stored. For the l -th entry, $0 \leq l \leq ne - 1$, of A , its row index i , column index j and value A_{ij} , $0 \leq i \leq m - 1$, $0 \leq j \leq n - 1$, are stored as the l -th components of the integer arrays `A_row` and `A_col` and real array `A_val`, respectively, while the number of nonzeros is recorded as `A_ne = ne`.

1.1.8.4 Sparse row-wise storage format

Again only the nonzero entries are stored, but this time they are ordered so that those in row i appear directly before those in row $i+1$. For the i -th row of A the i -th component of the integer array `A_ptr` holds the position of the first entry in this row, while `A_ptr(m)` holds the total number of entries plus one. The column indices j , $0 \leq j \leq n - 1$, and values A_{ij} of the nonzero entries in the i -th row are stored in components $l = A_ptr(i), \dots, A_ptr(i+1)-1$, $0 \leq i \leq m - 1$, of the integer array `A_col`, and real array `A_val`, respectively. For sparse matrices, this scheme almost always requires less storage than its predecessor.

1.1.8.5 Sparse column-wise storage format

Once again only the nonzero entries are stored, but this time they are ordered so that those in column j appear directly before those in column $j+1$. For the j -th column of A the j -th component of the integer array A_ptr holds the position of the first entry in this column, while $A_ptr(n)$ holds the total number of entries plus one. The row indices i , $0 \leq i \leq m-1$, and values A_{ij} of the nonzero entries in the j -th column are stored in components $l = A_ptr(j), \dots, A_ptr(j+1)-1$, $0 \leq j \leq n-1$, of the integer array A_row , and real array A_val , respectively. As before, for sparse matrices, this scheme almost always requires less storage than the co-ordinate format.

1.1.9 Symmetric matrix storage formats

Likewise, the symmetric n by n weighted-residual Hessian matrix $H = H(x, y)$ may be presented and stored in a variety of formats. But crucially symmetry is exploited by only storing values from the lower triangular part (i.e., those entries that lie on or below the leading diagonal).

1.1.9.1 Dense storage format

The matrix H is stored as a compact dense matrix by rows, that is, the values of the entries of each row in turn are stored in order within an appropriate real one-dimensional array. Since H is symmetric, only the lower triangular part (that is the part h_{ij} for $0 \leq j \leq i \leq n-1$) need be held. In this case the lower triangle should be stored by rows, that is component $i * i/2 + j$ of the storage array H_val will hold the value h_{ij} (and, by symmetry, h_{ji}) for $0 \leq j \leq i \leq n-1$.

1.1.9.2 Sparse co-ordinate storage format

Only the nonzero entries of the matrices are stored. For the l -th entry, $0 \leq l \leq ne-1$, of H , its row index i , column index j and value h_{ij} , $0 \leq j \leq i \leq n-1$, are stored as the l -th components of the integer arrays H_row and H_col and real array H_val , respectively, while the number of nonzeros is recorded as $H_ne = ne$. Note that only the entries in the lower triangle should be stored.

1.1.9.3 Sparse row-wise storage format

Again only the nonzero entries are stored, but this time they are ordered so that those in row i appear directly before those in row $i+1$. For the i -th row of H the i -th component of the integer array H_ptr holds the position of the first entry in this row, while $H_ptr(n)$ holds the total number of entries plus one. The column indices j , $0 \leq j \leq i$, and values h_{ij} of the entries in the i -th row are stored in components $l = H_ptr(i), \dots, H_ptr(i+1)-1$ of the integer array H_col , and real array H_val , respectively. Note that as before only the entries in the lower triangle should be stored. For sparse matrices, this scheme almost always requires less storage than its predecessor.

1.1.9.4 Diagonal storage format

If H is diagonal (i.e., $H_{ij} = 0$ for all $0 \leq i \neq j \leq n-1$) only the diagonal entries H_{ii} , $0 \leq i \leq n-1$ need be stored, and the first n components of the array H_val may be used for the purpose.

1.1.9.5 Multiples of the identity storage format

If H is a multiple of the identity matrix, (i.e., $H = \alpha I$ where I is the n by n identity matrix and α is a scalar), it suffices to store α as the first component of H_val .

1.1.9.6 The identity matrix format

If H is the identity matrix, no values need be stored.

1.1.9.7 The zero matrix format

The same is true if H is the zero matrix.

Chapter 2

File Index

2.1 File List

Here is a list of all files with brief descriptions:

nls.h	9
---------------------------------	---

Chapter 3

File Documentation

3.1 nls.h File Reference

```
#include <stdbool.h>
#include "galahad_precision.h"
```

Data Structures

- struct [nls_subproblem_control_type](#)
- struct [nls_control_type](#)
- struct [nls_time_type](#)
- struct [nls_subproblem_inform_type](#)
- struct [nls_inform_type](#)

Functions

- void [nls_initialize](#) (void **data, struct [nls_control_type](#) *control, struct [nls_inform_type](#) *inform)
- void [nls_read_specfile](#) (struct [nls_control_type](#) *control, const char specfile[])
- void [nls_import](#) (struct [nls_control_type](#) *control, void **data, int *status, int n, int m, const char J_type[], int J_ne, const int J_row[], const int J_col[], const int J_ptr[], const char H_type[], int H_ne, const int H_row[], const int H_col[], const int H_ptr[], const char P_type[], int P_ne, const int P_row[], const int P_col[], const int P_ptr[], const real_wp_ w[])
- void [nls_reset_control](#) (struct [nls_control_type](#) *control, void **data, int *status,)
- void [nls_solve_with_mat](#) (void **data, void *userdata, int *status, int n, int m, real_wp_ x[], real_wp_ c[], real_wp_ g[], int(*eval_c)(int, int, const real_wp_[], real_wp_[], const void *), int j_ne, int(*eval_j)(int, int, int, const real_wp_[], real_wp_[], const void *), int h_ne, int(*eval_h)(int, int, int, const real_wp_[], const real_wp_[], real_wp_[], const void *), int p_ne, int(*eval_hprods)(int, int, int, const real_wp_[], const real_wp_[], real_wp_[], bool, const void *))
- void [nls_solve_without_mat](#) (void **data, void *userdata, int *status, int n, int m, real_wp_ x[], real_wp_ c[], real_wp_ g[], int(*eval_c)(int, int, const real_wp_[], real_wp_[], const void *), int(*eval_jprod)(int, int, const real_wp_[], const bool, real_wp_[], const real_wp_[], bool, const void *), int(*eval_hprod)(int, int, const real_wp_[], const real_wp_[], real_wp_[], const real_wp_[], bool, const void *), int p_ne, int(*eval_hprods)(int, int, int, const real_wp_[], const real_wp_[], real_wp_[], bool, const void *))
- void [nls_solve_reverse_with_mat](#) (void **data, int *status, int *eval_status, int n, int m, real_wp_ x[], real_wp_ c[], real_wp_ g[], int j_ne, real_wp_ J_val[], const real_wp_ y[], int h_ne, real_wp_ H_val[], real_wp_ v[], int p_ne, real_wp_ P_val[])
- void [nls_solve_reverse_without_mat](#) (void **data, int *status, int *eval_status, int n, int m, real_wp_ x[], real_wp_ c[], real_wp_ g[], bool *transpose, real_wp_ u[], real_wp_ v[], real_wp_ y[], int p_ne, real_wp_ P_val[])
- void [nls_information](#) (void **data, struct [nls_inform_type](#) *inform, int *status)
- void [nls_terminate](#) (void **data, struct [nls_control_type](#) *control, struct [nls_inform_type](#) *inform)

3.1.1 Data Structure Documentation

3.1.1.1 struct nls_subproblem_control_type

control derived type as a C struct

Data Fields

char	alive_file[31]	see alive_unit
int	alive_unit	removal of the file alive_file from unit alive_unit terminates execution
real_wp_	clock_time_limit	the maximum elapsed clock time allowed (-ve means infinite)
real_wp_	cpu_time_limit	the maximum CPU time allowed (-ve means infinite)
bool	deallocate_error_fatal	if .deallocate_error_fatal is true, any array/pointer deallocation error will terminate execution. Otherwise, computation will continue
int	error	error and warning diagnostics occur on stream error
real_wp_	eta_successful	REAL (KIND = wp) :: initial_inner_weight = 0.0001_wp a potential iterate will only be accepted if the actual decrease $f - f(x_{\text{new}})$ is larger than .eta_successful times that predicted by a quadratic model of the decrease. The regularization weight will be decreased if this relative decrease is greater than .eta_very_successful but smaller than .eta_too_successful.
real_wp_	eta_too_successful	see eta_successful
real_wp_	eta_very_successful	see eta_successful
int	hessian_available	is the Hessian matrix of second derivatives available (≥ 2), is access only via matrix-vector products (=1) or is it not available (≤ 0) ?
real_wp_	initial_inner_weight	initial value for the inner regularization weight for tensor GN (-ve $\Rightarrow 0$)
real_wp_	initial_weight	initial value for the regularization weight (-ve $\Rightarrow 1/\ g_0\ $)
int	jacobian_available	is the Jacobian matrix of first derivatives available (≥ 2), is access only via matrix-vector products (=1) or is it not available (≤ 0) ?
real_wp_	large_root	see reduce_gap
bool	magic_step	allow the user to perform a "magic" step to improve the objective
int	maxit	the maximum number of iterations performed
real_wp_	minimum_weight	minimum permitted regularization weight
int	model	the model used. Possible values are <ul style="list-style-type: none"> • 0 dynamic (<i>not yet implemented</i>) • 1 first-order (no Hessian) • 2 barely second-order (identity Hessian) • 3 Gauss-Newton ($J^T J$ Hessian) • 4 second-order (exact Hessian) • 5 Gauss-Newton to Newton transition • 6 tensor Gauss-Newton treated as a least-squares model • 7 tensor Gauss-Newton treated as a general model • 8 tensor Gauss-Newton transition from a least-squares to a general mode
int	non_monotone	non-monotone ≤ 0 monotone strategy used, anything else non-monotone strategy with this history length used

Data Fields

int	norm	<p>the regularization norm used. The norm is defined via $\ v\ ^2 = v^T S v$, and will define the preconditioner used for iterative methods. Possible values for S are</p> <ul style="list-style-type: none"> • -3 user's own regularization norm • -2 S = limited-memory BFGS matrix (with .PSLS_control.lbfgs_vectors history) (<i>not yet implemented</i>) • -1 identity (= Euclidan two-norm) • 0 automatic (<i>not yet implemented</i>) • 1 diagonal, $S = \text{diag}(\max(J^T J \text{ Hessian}, \text{.PSLS_contro.min_diagonal}))$ • 2 diagonal, $S = \text{diag}(\max(\text{Hessian}, \text{.PSLS_contro.min_diagonal}))$ • 3 banded, $S = \text{band}(\text{Hessian})$ with semi-bandwidth .PSLS_control.semi_bandwidth • 4 re-ordered band, $P = \text{band}(\text{order}(A))$ with semi-bandwidth .PSLS_control.semi_bandwidth • 5 full factorization, $S = \text{Hessian}$, Schnabel-Eskow modification • 6 full factorization, $S = \text{Hessian}$, GMPS modification (<i>not yet implemented</i>) • 7 incomplete factorization of Hessian, Lin-More' • 8 incomplete factorization of Hessian, HSL_MI28 • 9 incomplete factorization of Hessian, Munskgaard (<i>not yet implemented</i>) • 10 expanding band of Hessian (<i>not yet implemented</i>)
int	out	general output occurs on stream out
real_wp_	power	the regularization power ($<2 \Rightarrow$ chosen according to the model)
char	prefix[31]	all output lines will be prefixed by .prefix(2:LEN(TRIM(.prefix))-1) where .prefix contains the required string enclosed in quotes, e.g. "string" or 'string'
int	print_gap	the number of iterations between printing
int	print_level	<p>the level of output required.</p> <ul style="list-style-type: none"> • ≤ 0 gives no output, • $= 1$ gives a one-line summary for every iteration, • $= 2$ gives a summary of the inner iteration for each iteration, • ≥ 3 gives increasingly verbose (debugging) output
bool	print_obj	print values of the objective/gradient rather than $\ c\ $ and its gradien

Data Fields

real_wp_	reduce_gap	expert parameters as suggested in Gould, Porcelli and Toint, "Updating t regularization parameter in the adaptive cubic regularization algorithm" RAL-TR-2011-007, Rutherford Appleton Laboratory, England (2011), http://epubs.stfc.ac.uk/bitstream/6181/RAL-TR-2011-007.pdf (these are denoted beta, epsilon_chi and alpha_max in the paper)
bool	renormalize_weight	should the weight be renormalized to account for a change in scaling?
bool	space_critical	if .space_critical true, every effort will be made to use as little space as possible. This may result in longer computation time
int	start_print	any printing will start on this iteration
real_wp_	stop_c_absolute	overall convergence tolerances. The iteration will terminate when $\ c(x)\ _2 \leq \text{MAX}(\text{.stop_c_absolute}, \text{.stop_c_relative} * \ c(x_{\text{initial}})\ _2)$, or when the norm of the gradient, $g = J^T(x)c(x)/\ c(x)\ _2$, of $\ c\ _2$, satisfies $\ g\ _2 \leq \text{MAX}(\text{.stop_g_absolute}, \text{.stop_g_relative} * \ g_{\text{initial}}\ _2)$, or if the step is less than .stop_s
real_wp_	stop_c_relative	see stop_c_absolute
real_wp_	stop_g_absolute	see stop_c_absolute
real_wp_	stop_g_relative	see stop_c_absolute
int	stop_print	any printing will stop on this iteration
real_wp_	stop_s	see stop_c_absolute
bool	subproblem_direct	use a direct (factorization) or (preconditioned) iterative method to find the search direction
real_wp_	switch_to_newton	if the Gauss-Newton to Newton model is specified, switch to Newton as soon as the norm of the gradient g is smaller than switch_to_newton
real_wp_	tiny_gap	see reduce_gap
real_wp_	weight_decrease	REAL (KIND = wp) :: weight_decrease = half.
real_wp_	weight_decrease_min	on very successful iterations, the regularization weight will be reduced by the factor .weight_decrease but no more than .weight_decrease_min while if the iteration is unsuccessful, the weight will be increased by a factor .weight_increase but no more than .weight_increase_max (these are delta_1, delta_2, delta3 and delta_max in Gould, Porcelli and Toint, 2011)
real_wp_	weight_increase	REAL (KIND = wp) :: weight_increase = two.
real_wp_	weight_increase_max	see weight_increase
int	weight_update_strategy	define the weight-update strategy: 1 (basic), 2 (reset to zero when very successful), 3 (imitate TR), 4 (increase lower bound), 5 (GPT)

3.1.1.2 struct nls_control_type

control derived type as a C struct

Examples

[nlst.c](#), and [nlstf.c](#).

Data Fields

char	alive_file[31]	see alive_unit
int	alive_unit	removal of the file alive_file from unit alive_unit terminates execution

Data Fields

real_wp_	clock_time_limit	the maximum elapsed clock time allowed (-ve means infinite)
real_wp_	cpu_time_limit	the maximum CPU time allowed (-ve means infinite)
bool	deallocate_error_fatal	if .deallocate_error_fatal is true, any array/pointer deallocation error will terminate execution. Otherwise, computation will continue
int	error	error and warning diagnostics occur on stream error
real_wp_	eta_successful	REAL (KIND = wp) :: initial_inner_weight = 0.0001_wp a potential iterate will only be accepted if the actual decrease $f - f(x_{\text{new}})$ is larger than .eta_successful times that predicted by a quadratic model of the decrease. The regularization weight will be decreased if this relative decrease is greater than .eta_very_successful but smaller than .eta_too_successful.
real_wp_	eta_too_successful	see eta_successful
real_wp_	eta_very_successful	see eta_successful
bool	f_indexing	use C or Fortran sparse matrix indexing
int	hessian_available	is the Hessian matrix of second derivatives available (≥ 2), is access only via matrix-vector products (=1) or is it not available (≤ 0) ?
real_wp_	initial_inner_weight	initial value for the inner regularization weight for tensor GN (-ve $\Rightarrow 0$)
real_wp_	initial_weight	initial value for the regularization weight (-ve $\Rightarrow 1/\ g_0\ $)
int	jacobian_available	is the Jacobian matrix of first derivatives available (≥ 2), is access only via matrix-vector products (=1) or is it not available (≤ 0) ?
real_wp_	large_root	see reduce_gap
bool	magic_step	allow the user to perform a "magic" step to improve the objective
int	maxit	the maximum number of iterations performed
real_wp_	minimum_weight	minimum permitted regularization weight

Data Fields

	int	model	<p>the model used. Possible values are</p> <ul style="list-style-type: none"> • 0 dynamic (<i>not yet implemented</i>) • 1 first-order (no Hessian) • 2 barely second-order (identity Hessian) • 3 Gauss-Newton ($J^T J$ Hessian) • 4 second-order (exact Hessian) • 5 Gauss-Newton to Newton transition • 6 tensor Gauss-Newton treated as a least-squares model • 7 tensor Gauss-Newton treated as a general model • 8 tensor Gauss-Newton transition from a least-squares to a general mode
	int	non_monotone	<p>non-monotone ≤ 0 monotone strategy used, anything else non-monotone strategy with this history length used</p>

Data Fields

	int	norm	<p>the regularization norm used. The norm is defined via $\ v\ ^2 = v^T S v$, and will define the preconditioner used for iterative methods. Possible values for S are</p> <ul style="list-style-type: none"> • -3 user's own regularization norm • -2 S = limited-memory BFGS matrix (with .PSLS_control.lbfgs_vectors history) (<i>not yet implemented</i>) • -1 identity (= Euclidan two-norm) • 0 automatic (<i>not yet implemented</i>) • 1 diagonal, $S = \text{diag}(\max(J^T J \text{ Hessian}, \text{.PSLS_contro.min_diagonal}))$ • 2 diagonal, $S = \text{diag}(\max(\text{Hessian}, \text{.PSLS_contro.min_diagonal}))$ • 3 banded, $S = \text{band}(\text{Hessian})$ with semi-bandwidth .PSLS_control.semi_bandwidth • 4 re-ordered band, $P = \text{band}(\text{order}(A))$ with semi-bandwidth .PSLS_control.semi_bandwidth • 5 full factorization, $S = \text{Hessian}$, Schnabel-Eskow modification • 6 full factorization, $S = \text{Hessian}$, GMPS modification (<i>not yet implemented</i>) • 7 incomplete factorization of Hessian, Lin-More' • 8 incomplete factorization of Hessian, HSL_MI28 • 9 incomplete factorization of Hessian, Munskgaard (<i>not yet implemented</i>) • 10 expanding band of Hessian (<i>not yet implemented</i>)
	int	out	general output occurs on stream out
	real_wp_	power	the regularization power ($<2 \Rightarrow$ chosen according to the model)
	char	prefix[31]	all output lines will be prefixed by .prefix(2:LEN(TRIM(.prefix))-1) where .prefix contains the required string enclosed in quotes, e.g. "string" or 'string'
	int	print_gap	the number of iterations between printing

Data Fields

int	print_level	the level of output required. <ul style="list-style-type: none"> • ≤ 0 gives no output, • $= 1$ gives a one-line summary for every iteration, • $= 2$ gives a summary of the inner iteration for each iteration, • ≥ 3 gives increasingly verbose (debugging) output
bool	print_obj	print values of the objective/gradient rather than $\ c\ $ and its gradient
real_wp_	reduce_gap	expert parameters as suggested in Gould, Porcelli and Toint, "Updating t regularization parameter in the adaptive cubic regularization algorithm" RAL-TR-2011-007, Rutherford Appleton Laboratory, England (2011), http://epubs.stfc.ac.uk/bitstream/6181/RAL-TR-2011-007.pdf (these are denoted beta, epsilon_chi and alpha_max in the paper)
bool	renormalize_weight	should the weight be renormalized to account for a change in scaling?
bool	space_critical	if .space_critical true, every effort will be made to use as little space as possible. This may result in longer computation time
int	start_print	any printing will start on this iteration
real_wp_	stop_c_absolute	overall convergence tolerances. The iteration will terminate when $\ c(x)\ _2 \leq \text{MAX}(\text{.stop_c_absolute}, \text{.stop_c_relative} * \ c(x_{\text{initial}})\ _2)$ or when the norm of the gradient, $g = J^T(x)c(x)/\ c(x)\ _2$, of $\ c(x)\ _2$ satisfies $\ g\ _2 \leq \text{MAX}(\text{.stop_g_absolute}, \text{.stop_g_relative} * \ g_{\text{initial}}\ _2)$, or if the step is less than .stop_s
real_wp_	stop_c_relative	see stop_c_absolute
real_wp_	stop_g_absolute	see stop_c_absolute
real_wp_	stop_g_relative	see stop_c_absolute
int	stop_print	any printing will stop on this iteration
real_wp_	stop_s	see stop_c_absolute
struct nls_subproblem_control_type	subproblem_control	control parameters for the step-finding subproblem
bool	subproblem_direct	use a direct (factorization) or (preconditioned) iterative method to find the search direction
real_wp_	switch_to_newton	if the Gauss-Newton to Newton model is specified, switch to Newton as soon as the norm of the gradient g is smaller than switch_to_newton
real_wp_	tiny_gap	see reduce_gap

Data Fields

real_wp_	weight_decrease	REAL (KIND = wp) :: weight_decrease = half.
real_wp_	weight_decrease_min	on very successful iterations, the regularization weight will be reduced by the factor .weight_decrease but no more than .weight_decrease_min while if the iteration is unsuccessful, the weight will be increased by a factor .weight_increase but no more than .weight_increase_max (these are delta_1, delta_2, delta3 and delta_max in Gould, Porcelli and Toint, 2011)
real_wp_	weight_increase	REAL (KIND = wp) :: weight_increase = two.
real_wp_	weight_increase_max	see weight_increase
int	weight_update_strategy	define the weight-update strategy: 1 (basic), 2 (reset to zero when very successful), 3 (imitate TR), 4 (increase lower bound), 5 (GPT)

3.1.1.3 struct nls_time_type

time derived type as a C struct

Data Fields

real_sp_	analyse	the CPU time spent analysing the required matrices prior to factorizatio
real_wp_	clock_analyse	the clock time spent analysing the required matrices prior to factorizat
real_wp_	clock_factorize	the clock time spent factorizing the required matrices
real_wp_	clock_preprocess	the clock time spent preprocessing the problem
real_wp_	clock_solve	the clock time spent computing the search direction
real_wp_	clock_total	the total clock time spent in the package
real_sp_	factorize	the CPU time spent factorizing the required matrices
real_sp_	preprocess	the CPU time spent preprocessing the problem
real_sp_	solve	the CPU time spent computing the search direction
real_sp_	total	the total CPU time spent in the package

3.1.1.4 struct nls_subproblem_inform_type

inform derived type as a C struct

Data Fields

int	alloc_status	the status of the last attempted allocation/deallocation
char	bad_alloc[81]	the name of the array for which an allocation/deallocation error occurred
char	bad_eval[13]	the name of the user-supplied evaluation routine for which an error ocur

Data Fields

int	c_eval	the total number of evaluations of the residual function $c(x)$
int	cg_iter	the total number of CG iterations performed
real_wp_	factorization_average	the average number of factorizations per sub-problem solve
int	factorization_integer	the total integer workspace required for the factorization
int	factorization_max	the maximum number of factorizations in a sub-problem solve
int	factorization_real	the total real workspace required for the factorization
int	factorization_status	the return status from the factorization
int	h_eval	the total number of evaluations of the scaled Hessian $H(x,y)$ of $c(x)$
int	iter	the total number of iterations performed
int	j_eval	the total number of evaluations of the Jacobian $J(x)$ of $c(x)$
int	max_entries_factors	the maximum number of entries in the factors
real_wp_	norm_c	the norm of the residual $\ c(x)\ _W$ at the best estimate of the solution x , determined by NLS_solve
real_wp_	norm_g	the norm of the gradient of $\ c(x)\ _W$ of the objective function at the best estimate, x , of the solution determined by NLS_solve
real_wp_	obj	the value of the objective function $\frac{1}{2}\ c(x)\ _W^2$ at the best estimate the solution, x , determined by NLS_solve
int	status	return status. See NLS_solve for details
struct nls_time_type	time	timings (see above)
real_wp_	weight	the final regularization weight used

3.1.1.5 struct nls_inform_type

inform derived type as a C struct

Examples

[nlst.c](#), and [nlstf.c](#).

Data Fields

int	alloc_status	the status of the last attempted allocation/deallocation
char	bad_alloc[81]	the name of the array for which an allocation/deallocation error occurred
char	bad_eval[13]	the name of the user-supplied evaluation routine for which an error occur
int	c_eval	the total number of evaluations of the residual function $c(x)$
int	cg_iter	the total number of CG iterations performed
real_wp_	factorization_average	the average number of factorizations per sub-problem solve
int	factorization_integer	the total integer workspace required for the factorization
int	factorization_max	the maximum number of factorizations in a sub-problem solve
int	factorization_real	the total real workspace required for the factorization

Data Fields

	int	factorization_status	the return status from the factorization
	int	h_eval	the total number of evaluations of the scaled Hessian $H(x,y)$ of $c(x)$
	int	iter	the total number of iterations performed
	int	j_eval	the total number of evaluations of the Jacobian $J(x)$ of $c(x)$
	int	max_entries_factors	the maximum number of entries in the factors
	real_wp_	norm_c	the norm of the residual $\ c(x)\ _W$ at the best estimate of the solution x , determined by NLS_solve
	real_wp_	norm_g	the norm of the gradient of $\ c(x)\ _W$ of the objective function at the best estimate, x , of the solution determined by NLS_solve
	real_wp_	obj	the value of the objective function $\frac{1}{2}\ c(x)\ _W^2$ at the best estimate the solution, x , determined by NLS_solve
	int	status	return status. See NLS_solve for details
	struct nls_subproblem_inform_type	subproblem_inform	inform parameters for subproblem
	struct nls_time_type	time	timings (see above)
	real_wp_	weight	the final regularization weight used

3.1.2 Function Documentation

3.1.2.1 nls_import()

```

void nls_import (
    struct nls\_control\_type * control,
    void ** data,
    int * status,
    int n,
    int m,
    const char J_type[],
    int J_ne,
    const int J_row[],
    const int J_col[],
    const int J_ptr[],
    const char H_type[],
    int H_ne,
    const int H_row[],
    const int H_col[],
    const int H_ptr[],
    const char P_type[],
    int P_ne,
    const int P_row[],
    const int P_col[],
    const int P_ptr[],
    const real_wp_ w[] )

```

Import problem data into internal storage prior to solution.

Parameters

in	<i>control</i>	is a struct whose members provide control paramters for the remaining prcedures (see nls_control_type)
in, out	<i>data</i>	holds private internal data
in, out	<i>status</i>	is a scalar variable of type int, that gives the exit status from the package. Possible values are: <ul style="list-style-type: none"> • 1. The import was succesful, and the package is ready for the solve phase • -1. An allocation error occurred. A message indicating the offending array is written on unit control.error, and the returned allocation status and a string containing the name of the offending array are held in inform.alloc_status and inform.bad_alloc respectively. • -2. A deallocation error occurred. A message indicating the offending array is written on unit control.error and the returned allocation status and a string containing the name of the offending array are held in inform.alloc_status and inform.bad_alloc respectively. • -3. The restrictions $n > 0$, $m > 0$ or requirement that J/H/P_type contains its relevant string 'dense', 'dense_by_columns', 'coordinate', 'sparse_by_rows', 'sparse_by_columns', 'diagonal' or 'absent' has been violated.
in	<i>n</i>	is a scalar variable of type int, that holds the number of variables.
in	<i>m</i>	is a scalar variable of type int, that holds the number of residuals.
in	<i>J_type</i>	is a one-dimensional array of type char that specifies the symmetric storage scheme used for the Jacobian, J . It should be one of 'coordinate', 'sparse_by_rows', 'dense' or 'absent', the latter if access to the Jacobian is via matrix-vector products; lower or upper case variants are allowed.
in	<i>J_ne</i>	is a scalar variable of type int, that holds the number of entries in J in the sparse co-ordinate storage scheme. It need not be set for any of the other schemes.
in	<i>J_row</i>	is a one-dimensional array of size J_ne and type int, that holds the row indices of J in the sparse co-ordinate storage scheme. It need not be set for any of the other schemes, and in this case can be NULL.
in	<i>J_col</i>	is a one-dimensional array of size J_ne and type int, that holds the column indices of J in either the sparse co-ordinate, or the sparse row-wise storage scheme. It need not be set when the dense or diagonal storage schemes are used, and in this case can be NULL.
in	<i>J_ptr</i>	is a one-dimensional array of size n+1 and type int, that holds the starting position of each row of J , as well as the total number of entries plus one, in the sparse row-wise storage scheme. It need not be set when the other schemes are used, and in this case can be NULL.
in	<i>H_type</i>	is a one-dimensional array of type char that specifies the symmetric storage scheme used for the Hessian, H . It should be one of 'coordinate', 'sparse_by_rows', 'dense', 'diagonal' or 'absent', the latter if access to H is via matrix-vector products; lower or upper case variants are allowed.
in	<i>H_ne</i>	is a scalar variable of type int, that holds the number of entries in the lower triangular part of H in the sparse co-ordinate storage scheme. It need not be set for any of the other three schemes.
in	<i>H_row</i>	is a one-dimensional array of size H_ne and type int, that holds the row indices of the lower triangular part of H in the sparse co-ordinate storage scheme. It need not be set for any of the other three schemes, and in this case can be NULL.
in	<i>H_col</i>	is a one-dimensional array of size H_ne and type int, that holds the column indices of the lower triangular part of H in either the sparse co-ordinate, or the sparse row-wise storage scheme. It need not be set when the dense or diagonal storage schemes are used, and in this case can be NULL.

Parameters

in	<i>H_ptr</i>	is a one-dimensional array of size $n+1$ and type <code>int</code> , that holds the starting position of each row of the lower triangular part of H , as well as the total number of entries plus one, in the sparse row-wise storage scheme. It need not be set when the other schemes are used, and in this case can be <code>NULL</code> .
in	<i>P_type</i>	is a one-dimensional array of type <code>char</code> that specifies the symmetric storage scheme used for the residual-Hessians-vector product matrix, P . It should be one of 'coordinate', 'sparse_by_columns', 'dense_by_columns' or 'absent', the latter if access to P is via matrix-vector products; lower or upper case variants are allowed.
in	<i>P_ne</i>	is a scalar variable of type <code>int</code> , that holds the number of entries in P in the sparse co-ordinate storage scheme. It need not be set for any of the other schemes.
in	<i>P_row</i>	is a one-dimensional array of size P_ne and type <code>int</code> , that holds the row indices of P in either the sparse co-ordinate, or the sparse column-wise storage scheme. It need not be set when the dense storage scheme is used, and in this case can be <code>NULL</code> .
in	<i>P_col</i>	is a one-dimensional array of size P_ne and type <code>int</code> , that holds the row indices of P in the sparse co-ordinate storage scheme. It need not be set for any of the other schemes, and in this case can be <code>NULL</code> .
in	<i>P_ptr</i>	is a one-dimensional array of size $m+1$ and type <code>int</code> , that holds the starting position of each row of P , as well as the total number of entries plus one, in the sparse row-wise storage scheme. It need not be set when the other schemes are used, and in this case can be <code>NULL</code> .
in	<i>w</i>	is a one-dimensional array of size m and type <code>double</code> , that holds the values w of the weights on the residuals in the least-squares objective function. It need not be set if the weights are all ones, and in this case can be <code>NULL</code> .

Examples

[nlst.c](#), and [nlstf.c](#).

3.1.2.2 nls_information()

```
void nls_information (
    void ** data,
    struct nls_inform_type * inform,
    int * status )
```

Provides output information

Parameters

in, out	<i>data</i>	holds private internal data
out	<i>inform</i>	is a struct containing output information (see nls_inform_type)
out	<i>status</i>	is a scalar variable of type <code>int</code> , that gives the exit status from the package. Possible values are (currently): <ul style="list-style-type: none"> • 0. The values were recorded successfully

Examples

[nlst.c](#), and [nlstf.c](#).

3.1.2.3 nls_initialize()

```
void nls_initialize (
    void ** data,
    struct nls_control_type * control,
    struct nls_inform_type * inform )
```

Set default control values and initialize private data

Parameters

in, out	<i>data</i>	holds private internal data
out	<i>control</i>	is a struct containing control information (see nls_control_type)
out	<i>inform</i>	is a struct containing output information (see nls_inform_type)

Examples

[nlst.c](#), and [nlstf.c](#).

3.1.2.4 nls_read_specfile()

```
void nls_read_specfile (
    struct nls_control_type * control,
    const char specfile[] )
```

Read the content of a specification file, and assign values associated with given keywords to the corresponding control parameters

Parameters

in, out	<i>control</i>	is a struct containing control information (see nls_control_type)
in	<i>specfile</i>	is a character string containing the name of the specification file

3.1.2.5 nls_reset_control()

```
void nls_reset_control (
    struct nls_control_type * control,
```

```
void ** data,
int * status )
```

Reset control parameters after import if required.

Parameters

in	<i>control</i>	is a struct whose members provide control paramters for the remaining prcedures (see nls_control_type)
in, out	<i>data</i>	holds private internal data
in, out	<i>status</i>	is a scalar variable of type int, that gives the exit status from the package. Possible values are: <ul style="list-style-type: none"> 1. The import was succesful, and the package is ready for the solve phase

3.1.2.6 nls_solve_reverse_with_mat()

```
void nls_solve_reverse_with_mat (
    void ** data,
    int * status,
    int * eval_status,
    int n,
    int m,
    real_wp_ x[],
    real_wp_ c[],
    real_wp_ g[],
    int j_ne,
    real_wp_ J_val[],
    const real_wp_ y[],
    int h_ne,
    real_wp_ H_val[],
    real_wp_ v[],
    int p_ne,
    real_wp_ P_val[] )
```

Find a local minimizer of a given function using a trust-region method.

This call is for the case where $H = \nabla_{xx}f(x)$ is provided specifically, but function/derivative information is only available by returning to the calling procedure

Parameters

in, out	<i>data</i>	holds private internal data
---------	-------------	-----------------------------

Parameters

in, out	status	<p>is a scalar variable of type int, that gives the entry and exit status from the package.</p> <p>On initial entry, status must be set to 1. Possible exit are:</p> <ul style="list-style-type: none"> • 0. The run was succesful • -1. An allocation error occurred. A message indicating the offending array is written on unit control.error, and the returned allocation status and a string containing the name of the offending array are held in inform.alloc_status and inform.bad_alloc respectively. • -2. A deallocation error occurred. A message indicating the offending array is written on unit control.error and the returned allocation status and a string containing the name of the offending array are held in inform.alloc_status and inform.bad_alloc respectively. • -3. The restriction $n > 0$ or requirement that type contains its relevant string 'dense', 'coordinate', 'sparse_by_rows', 'diagonal' or 'absent' has been violated. • -9. The analysis phase of the factorization failed; the return status from the factorization package is given in the component inform.factor_status • -10. The factorization failed; the return status from the factorization package is given in the component inform.factor_status. • -11. The solution of a set of linear equations using factors from the factorization package failed; the return status from the factorization package is given in the component inform.factor_status. • -16. The problem is so ill-conditioned that further progress is impossible. • -17. The step is too small to make further impact. • -18. Too many iterations have been performed. This may happen if control.maxit is too small, but may also be symptomatic of a badly scaled problem. • -19. The CPU time limit has been reached. This may happen if control.cpu_time_limit is too small, but may also be symptomatic of a badly scaled problem. • -82. The user has forced termination of solver by removing the file named control.alive_file from unit unit control.alive_unit. • 2. The user should compute the vector of residuals $c(x)$ at the point x indicated in x and then re-enter the function. The required value should be set in c, and eval_status should be set to 0. If the user is unable to evaluate $c(x)$— for instance, if the function is undefined at x— the user need not set c, but should then set eval_status to a non-zero value. • 3. The user should compute the Jacobian of the vector of residual functions, $\nabla_x c(x)$, at the point x indicated in x and then re-enter the function. The l-th component of the Jacobian stored according to the scheme specified for the remainder of J in the earlier call to nls_import should be set in J_val[l], for $l = 0, \dots, J_ne-1$ and eval_status should be set to 0. If the user is unable to evaluate a component of J — for instance, if a component of the matrix is undefined at x — the user need not set J_val, but should then set eval_status to a non-zero value.
---------	--------	--

Parameters

	<i>status</i>	(continued) <ul style="list-style-type: none"> • 4. The user should compute the matrix $H = \sum_{i=1}^m v_i \nabla_{xx} c_i(x)$ of weighted residual Hessian evaluated at $x = x$ and $v = v$ and then re-enter the function. The l-th component of the matrix stored according to the scheme specified for the remainder of H in the earlier call to <code>nls_import</code> should be set in <code>H_val[l]</code>, for $l = 0, \dots, H_ne-1$ and <code>eval_status</code> should be set to 0. If the user is unable to evaluate a component of H — for instance, if a component of the matrix is undefined at x — the user need not set <code>H_val</code>, but should then set <code>eval_status</code> to a non-zero value. Note that this return will not happen if the Gauss-Newton model is selected. • 7. The user should compute the entries of the matrix P, whose i-th column is the product $\nabla_{xx} c_i(x) v$ between $\nabla_{xx} c_i(x)$, the Hessian of the i-th component of the residual $c(x)$ at $x = x$, and $v = v$ and then re-enter the function. The l-th component of the matrix stored according to the scheme specified for the remainder of P in the earlier call to <code>nls_import</code> should be set in <code>P_val[l]</code>, for $l = 0, \dots, P_ne-1$ and <code>eval_status</code> should be set to 0. If the user is unable to evaluate a component of P — for instance, if a component of the matrix is undefined at x — the user need not set <code>P_val</code>, but should then set <code>eval_status</code> to a non-zero value. Note that this return will not happen if either the Gauss-Newton or Newton models is selected.
in, out	<i>eval_status</i>	is a scalar variable of type int, that is used to indicate if objective function/gradient/Hessian values can be provided (see above)
in	<i>n</i>	is a scalar variable of type int, that holds the number of variables
in	<i>m</i>	is a scalar variable of type int, that holds the number of residuals.
in, out	<i>x</i>	is a one-dimensional array of size n and type double, that holds the values x of the optimization variables. The j -th component of x , $j = 0, \dots, n-1$, contains x_j .
in, out	<i>c</i>	is a one-dimensional array of size m and type double, that holds the residual $c(x)$. The i -th component of c , $j = 0, \dots, m-1$, contains $c_j(x)$. See <code>status = 2</code> , above, for more details.
in, out	<i>g</i>	is a one-dimensional array of size n and type double, that holds the gradient $g = \nabla_x f(x)$ of the objective function. The j -th component of g , $j = 0, \dots, n-1$, contains g_j .
in	<i>j_ne</i>	is a scalar variable of type int, that holds the number of entries in the Jacobian matrix J .
in	<i>J_val</i>	is a one-dimensional array of size <code>j_ne</code> and type double, that holds the values of the entries of the Jacobian matrix J in any of the available storage schemes. See <code>status = 3</code> , above, for more details.
in, out	<i>y</i>	is a one-dimensional array of size m and type double, that is used for reverse communication. See <code>status = 4</code> above for more details.
in	<i>h_ne</i>	is a scalar variable of type int, that holds the number of entries in the lower triangular part of the Hessian matrix H .
in	<i>H_val</i>	is a one-dimensional array of size <code>h_ne</code> and type double, that holds the values of the entries of the lower triangular part of the Hessian matrix H in any of the available storage schemes. See <code>status = 4</code> , above, for more details.
in, out	<i>v</i>	is a one-dimensional array of size n and type double, that is used for reverse communication. See <code>status = 7</code> , above, for more details.
in	<i>p_ne</i>	is a scalar variable of type int, that holds the number of entries in the residual-Hessians-vector product matrix, P .
in	<i>P_val</i>	is a one-dimensional array of size <code>p_ne</code> and type double, that holds the values of the entries of the residual-Hessians-vector product matrix, P . See <code>status = 7</code> , above, for more details.

Examples

[nlst.c](#), and [nlstf.c](#).

3.1.2.7 nls_solve_reverse_without_mat()

```
void nls_solve_reverse_without_mat (
    void ** data,
    int * status,
    int * eval_status,
    int n,
    int m,
    real_wp_ x[],
    real_wp_ c[],
    real_wp_ g[],
    bool * transpose,
    real_wp_ u[],
    real_wp_ v[],
    real_wp_ y[],
    int p_ne,
    real_wp_ P_val[] )
```

Find a local minimizer of a given function using a trust-region method.

This call is for the case where access to $H = \nabla_{xx}f(x)$ is provided by Hessian-vector products, but function/derivative information is only available by returning to the calling procedure.

Parameters

in, out	<i>data</i>	holds private internal data
in, out	<i>status</i>	<p>is a scalar variable of type int, that gives the entry and exit status from the package.</p> <p>On initial entry, status must be set to 1. Possible exit are:</p> <ul style="list-style-type: none"> • 0. The run was succesful • -1. An allocation error occurred. A message indicating the offending array is written on unit control.error, and the returned allocation status and a string containing the name of the offending array are held in inform.alloc_status and inform.bad_alloc respectively. • -2. A deallocation error occurred. A message indicating the offending array is written on unit control.error and the returned allocation status and a string containing the name of the offending array are held in inform.alloc_status and inform.bad_alloc respectively. • -3. The restriction $n > 0$ or requirement that type contains its relevant string 'dense', 'coordinate', 'sparse_by_rows', 'diagonal' or 'absent' has been violated. • -9. The analysis phase of the factorization failed; the return status from the factorization package is given in the component inform.factor_status

Parameters

	<i>status</i>	<p>(continued)</p> <ul style="list-style-type: none"> • -10. The factorization failed; the return status from the factorization package is given in the component <code>inform.factor_status</code>. • -11. The solution of a set of linear equations using factors from the factorization package failed; the return status from the factorization package is given in the component <code>inform.factor_status</code>. • -16. The problem is so ill-conditioned that further progress is impossible. • -17. The step is too small to make further impact. • -18. Too many iterations have been performed. This may happen if <code>control.maxit</code> is too small, but may also be symptomatic of a badly scaled problem. • -19. The CPU time limit has been reached. This may happen if <code>control.cpu_time_limit</code> is too small, but may also be symptomatic of a badly scaled problem. • -82. The user has forced termination of solver by removing the file named <code>control.alive_file</code> from unit <code>unit control.alive_unit</code>. • 2. The user should compute the vector of residuals $c(x)$ at the point x indicated in <code>x</code> and then re-enter the function. The required value should be set in <code>c</code>, and <code>eval_status</code> should be set to 0. If the user is unable to evaluate $c(x)$— for instance, if the function is undefined at x— the user need not set <code>c</code>, but should then set <code>eval_status</code> to a non-zero value. • 5. The user should compute the sum $u + \nabla_x c(x)v$ (if <code>tranpose</code> is false) or $u + (\nabla_x c(x))^T v$ (if <code>tranpose</code> is true) between the product of the Jacobian $\nabla_x c(x)$ or its tranpose with the vector <code>v= v</code> and the vector <code>u = \$ u</code>, and then re-enter the function. The result should be set in <code>u</code>, and <code>eval_status</code> should be set to 0. If the user is unable to evaluate the sum — for instance, if the Jacobian is undefined at x — the user need not set <code>u</code>, but should then set <code>eval_status</code> to a non-zero value. • 6. The user should compute the sum $u + \sum_{i=1}^m y_i \nabla_{xx} c_i(x)v$ between the product of the weighted residual Hessian $H = \sum_{i=1}^m y_i \nabla_{xx} c_i(x)$ evaluated at $x= x$ and $y= y$ with the vector <code>v= v</code> and the the vector <code>u = \$ u</code>, and then re-enter the function. The result should be set in <code>u</code>, and <code>eval_status</code> should be set to 0. If the user is unable to evaluate the sum — for instance, if the weifghted residual Hessian is undefined at x — the user need not set <code>u</code>, but should then set <code>eval_status</code> to a non-zero value. • 7. The user should compute the entries of the matrix P, whose i-th column is the product $\nabla_{xx} c_i(x)v$ between $\nabla_{xx} c_i(x)$, the Hessian of the i-th component of the residual $c(x)$ at $x= x$, and <code>v= v</code> and then re-enter the function. The i-th component of the matrix stored according to the scheme specified for the remainder of P in the earlier call to <code>nls_import</code> should be set in <code>P_val[l]</code>, for $l = 0, \dots, P_ne-1$ and <code>eval_status</code> should be set to 0. If the user is unable to evaluate a component of P — for instance, if a component of the matrix is undefined at x — the user need not set <code>P_val</code>, but should then set <code>eval_status</code> to a non-zero value. Note that this return will not happen if either the Gauss-Newton or Newton models is selected.
<i>in, out</i>	<i>eval_status</i>	is a scalar variable of type <code>int</code> , that is used to indicate if objective function/gradient/Hessian values can be provided (see above)

Parameters

in	n	is a scalar variable of type int, that holds the number of variables
in	m	is a scalar variable of type int, that holds the number of residuals.
in, out	x	is a one-dimensional array of size n and type double, that holds the values x of the optimization variables. The j -th component of x , $j = 0, \dots, n-1$, contains x_j .
in, out	c	is a one-dimensional array of size m and type double, that holds the residual $c(x)$. The i -th component of c , $j = 0, \dots, m-1$, contains $c_j(x)$. See status = 2, above, for more details.
in, out	g	is a one-dimensional array of size n and type double, that holds the gradient $g = \nabla_x f(x)$ of the objective function. The j -th component of g , $j = 0, \dots, n-1$, contains g_j .
out	<i>transpose</i>	is a scalar variable of type bool, that indicates whether the product with Jacobian or its transpose should be obtained when status=5.
in, out	u	is a one-dimensional array of size $\max(n,m)$ and type double, that is used for reverse communication. See status = 5,6 above for more details.
in, out	v	is a one-dimensional array of size $\max(n,m)$ and type double, that is used for reverse communication. See status = 5,6,7 above for more details.
in, out	y	is a one-dimensional array of size m and type double, that is used for reverse communication. See status = 6 above for more details.
in	p_ne	is a scalar variable of type int, that holds the number of entries in the residual-Hessians-vector product matrix, P .
in	P_val	is a one-dimensional array of size P_ne and type double, that holds the values of the entries of the residual-Hessians-vector product matrix, P . See status = 7, above, for more details.

Examples

[nlst.c](#), and [nlstf.c](#).

3.1.2.8 nls_solve_with_mat()

```

void nls_solve_with_mat (
    void ** data,
    void * userdata,
    int * status,
    int n,
    int m,
    real_wp_ x[],
    real_wp_ c[],
    real_wp_ g[],
    int(*) (int, int, const real_wp_[], real_wp_[], const void *) eval_c,
    int j_ne,
    int(*) (int, int, int, const real_wp_[], real_wp_[], const void *) eval_j,
    int h_ne,
    int(*) (int, int, int, const real_wp_[], const real_wp_[], real_wp_[], const void
*) eval_h,
    int p_ne,
    int(*) (int, int, int, const real_wp_[], const real_wp_[], real_wp_[], bool, const
void *) eval_hprods )

```


Find a local minimizer of a given function using a trust-region method.

This call is for the case where $H = \nabla_{xx}f(x)$ is provided specifically, and all function/derivative information is available by function calls.

Parameters

in, out	<i>data</i>	holds private internal data
in	<i>userdata</i>	is a structure that allows data to be passed into the function and derivative evaluation programs.
in, out	<i>status</i>	<p>is a scalar variable of type int, that gives the entry and exit status from the package. On initial entry, status must be set to 1. Possible exit are:</p> <ul style="list-style-type: none"> • 0. The run was succesful • -1. An allocation error occurred. A message indicating the offending array is written on unit control.error, and the returned allocation status and a string containing the name of the offending array are held in inform.alloc_status and inform.bad_alloc respectively. • -2. A deallocation error occurred. A message indicating the offending array is written on unit control.error and the returned allocation status and a string containing the name of the offending array are held in inform.alloc_status and inform.bad_alloc respectively. • -3. The restriction $n > 0$ or requirement that type contains its relevant string 'dense', 'coordinate', 'sparse_by_rows', 'diagonal' or 'absent' has been violated. • -9. The analysis phase of the factorization failed; the return status from the factorization package is given in the component inform.factor_status • -10. The factorization failed; the return status from the factorization package is given in the component inform.factor_status. • -11. The solution of a set of linear equations using factors from the factorization package failed; the return status from the factorization package is given in the component inform.factor_status. • -16. The problem is so ill-conditioned that further progress is impossible. • -17. The step is too small to make further impact. • -18. Too many iterations have been performed. This may happen if control.maxit is too small, but may also be symptomatic of a badly scaled problem. • -19. The CPU time limit has been reached. This may happen if control.cpu_time_limit is too small, but may also be symptomatic of a badly scaled problem. • -82. The user has forced termination of solver by removing the file named control.alive_file from unit unit control.alive_unit.
in	<i>n</i>	is a scalar variable of type int, that holds the number of variables.
in	<i>m</i>	is a scalar variable of type int, that holds the number of residuals.
in, out	<i>x</i>	is a one-dimensional array of size n and type double, that holds the values x of the optimization variables. The j-th component of x , $j = 0, \dots, n-1$, contains x_j .
out	<i>c</i>	is a one-dimensional array of size m and type double, that holds the residual $c(x)$. The i-th component of c , $j = 0, \dots, n-1$, contains $c_j(x)$.
out	<i>g</i>	is a one-dimensional array of size n and type double, that holds the gradient $g = \nabla_x f(x)$ of the objective function. The j-th component of g , $j = 0, \dots, n-1$, contains g_j .

Parameters

	<i>eval_c</i>	is a user-supplied function that must have the following signature: <pre>int eval_c(int n, const double x[], double c[], const void *userdata)</pre> The components of the residual function $c(x)$ evaluated at $x=x$ must be assigned to <i>c</i> , and the function return value set to 0. If the evaluation is impossible at <i>x</i> , return should be set to a nonzero value. Data may be passed into <i>eval_c</i> via the structure <i>userdata</i> .
in	<i>j_ne</i>	is a scalar variable of type int, that holds the number of entries in the Jacobian matrix <i>J</i> .
	<i>eval_j</i>	is a user-supplied function that must have the following signature: <pre>int eval_j(int n, int m, int jne, const double x[], double j[], const void *userdata)</pre> The components of the Jacobian $J = \nabla_x c(x)$ of the residuals must be assigned to <i>j</i> in the same order as presented to <i>nls_import</i> , and the function return value set to 0. If the evaluation is impossible at <i>x</i> , return should be set to a nonzero value. Data may be passed into <i>eval_j</i> via the structure <i>userdata</i> .
in	<i>h_ne</i>	is a scalar variable of type int, that holds the number of entries in the lower triangular part of the Hessian matrix <i>H</i> if it is used.
	<i>eval_h</i>	is a user-supplied function that must have the following signature: <pre>int eval_h(int n, int m, int hne, const double x[], const double y[], double h[], const void *userdata)</pre> The nonzeros of the matrix $H = \sum_{i=1}^m y_i \nabla_{xx} c_i(x)$ of the weighted residual Hessian evaluated at $x=x$ and $y=y$ must be assigned to <i>h</i> in the same order as presented to <i>nls_import</i> , and the function return value set to 0. If the evaluation is impossible at <i>x</i> , return should be set to a nonzero value. Data may be passed into <i>eval_h</i> via the structure <i>userdata</i> .
in	<i>p_ne</i>	is a scalar variable of type int, that holds the number of entries in the residual-Hessians-vector product matrix <i>P</i> if it is used.
	<i>eval_hprods</i>	is an optional user-supplied function that may be NULL. If non-NULL, it must have the following signature: <pre>int eval_hprods(int n, int m, int pne, const double x[], const double v[], double p[], bool got_h, const void *userdata)</pre> The entries of the matrix <i>P</i> , whose <i>i</i> -th column is the product $\nabla_{xx} c_i(x)v$ between $\nabla_{xx} c_i(x)$, the Hessian of the <i>i</i> -th component of the residual $c(x)$ at $x=x$, and $v=v$ must be returned in <i>p</i> and the function return value set to 0. If the evaluation is impossible at <i>x</i> , return should be set to a nonzero value. Data may be passed into <i>eval_hprods</i> via the structure <i>userdata</i> .

Examples

[nlst.c](#), and [nlsf.c](#).

3.1.2.9 nls_solve_without_mat()

```
void nls_solve_without_mat (
    void ** data,
    void * userdata,
    int * status,
    int n,
    int m,
    real_wp_ x[],
    real_wp_ c[],
```

```

    real_wp_ g[],
    int(*) (int, int, const real_wp[], real_wp[], const void *) eval_c,
    int(*) (int, int, const real_wp[], const bool, real_wp[], const real_wp[],
bool, const void *) eval_jprod,
    int(*) (int, int, const real_wp[], const real_wp[], real_wp[], const real_wp[],
[], bool, const void *) eval_hprod,
    int p_ne,
    int(*) (int, int, int, const real_wp[], const real_wp[], real_wp[], bool, const
void *) eval_hprods )

```

Find a local minimizer of a given function using a trust-region method.

This call is for the case where access to $H = \nabla_{xx}f(x)$ is provided by Hessian-vector products, and all function/derivative information is available by function calls.

Parameters

in, out	<i>data</i>	holds private internal data
in	<i>userdata</i>	is a structure that allows data to be passed into the function and derivative evaluation programs.

Parameters

in, out	status	<p>is a scalar variable of type int, that gives the entry and exit status from the package.</p> <p>On initial entry, status must be set to 1.</p> <p>Possible exit are:</p> <ul style="list-style-type: none"> • 0. The run was succesful • -1. An allocation error occurred. A message indicating the offending array is written on unit control.error, and the returned allocation status and a string containing the name of the offending array are held in inform.alloc_status and inform.bad_alloc respectively. • -2. A deallocation error occurred. A message indicating the offending array is written on unit control.error and the returned allocation status and a string containing the name of the offending array are held in inform.alloc_status and inform.bad_alloc respectively. • -3. The restriction $n > 0$ or requirement that type contains its relevant string 'dense', 'coordinate', 'sparse_by_rows', 'diagonal' or 'absent' has been violated. • -9. The analysis phase of the factorization failed; the return status from the factorization package is given in the component inform.factor_status • -10. The factorization failed; the return status from the factorization package is given in the component inform.factor_status. • -11. The solution of a set of linear equations using factors from the factorization package failed; the return status from the factorization package is given in the component inform.factor_status. • -16. The problem is so ill-conditioned that further progress is impossible. • -17. The step is too small to make further impact. • -18. Too many iterations have been performed. This may happen if control.maxit is too small, but may also be symptomatic of a badly scaled problem. • -19. The CPU time limit has been reached. This may happen if control.cpu_time_limit is too small, but may also be symptomatic of a badly scaled problem. • -82. The user has forced termination of solver by removing the file named control.alive_file from unit unit control.alive_unit.
in	n	is a scalar variable of type int, that holds the number of variables
in	m	is a scalar variable of type int, that holds the number of residuals.
in, out	x	is a one-dimensional array of size n and type double, that holds the values x of the optimization variables. The j-th component of x, $j = 0, \dots, n-1$, contains x_j .
out	c	is a one-dimensional array of size m and type double, that holds the residual $c(x)$. The i-th component of c, $j = 0, \dots, n-1$, contains $c_j(x)$.
out	g	is a one-dimensional array of size n and type double, that holds the gradient $g = \nabla_x f(x)$ of the objective function. The j-th component of g, $j = 0, \dots, n-1$, contains g_j .

Parameters

	<i>eval_c</i>	is a user-supplied function that must have the following signature: <pre>int eval_c(int n, const double x[], double c[], const void *userdata)</pre> The components of the residual function $c(x)$ evaluated at $x=x$ must be assigned to c , and the function return value set to 0. If the evaluation is impossible at x , return should be set to a nonzero value. Data may be passed into <i>eval_c</i> via the structure <i>userdata</i> .
	<i>eval_jprod</i>	is a user-supplied function that must have the following signature: <pre>int eval_jprod(int n, int m, const double x[], bool transpose, double u[], const double v[], bool got_j, const void *userdata)</pre> The sum $u + \nabla_x c(x)v$ (if <i>transpose</i> is false) or The sum $u + (\nabla_x c(x))^T v$ (if <i>transpose</i> is true) between the product of the Jacobian $\nabla_x c(x)$ or its transpose with the vector $v=v$ and the vector u must be returned in u , and the function return value set to 0. If the evaluation is impossible at x , return should be set to a nonzero value. Data may be passed into <i>eval_jprod</i> via the structure <i>userdata</i> .
	<i>eval_hprod</i>	is a user-supplied function that must have the following signature: <pre>int eval_hprod(int n, int m, const double x[], const double y[], double u[], const double v[], bool got_h, const void *userdata)</pre> The sum $u + \sum_{i=1}^m y_i \nabla_{xx} c_i(x)v$ of the product of the weighted residual Hessian $H = \sum_{i=1}^m y_i \nabla_{xx} c_i(x)$ evaluated at $x=x$ and $y=y$ with the vector $v=v$ and the vector u must be returned in u , and the function return value set to 0. If the evaluation is impossible at x , return should be set to a nonzero value. The Hessians have already been evaluated or used at x if <i>got_h</i> is true. Data may be passed into <i>eval_hprod</i> via the structure <i>userdata</i> .
in	<i>p_ne</i>	is a scalar variable of type <i>int</i> , that holds the number of entries in the residual-Hessians-vector product matrix P if it is used.
	<i>eval_hprods</i>	is an optional user-supplied function that may be NULL. If non-NULL, it must have the following signature: <pre>int eval_hprods(int n, int m, int p_ne, const double x[], const double v[], double pval[], bool got_h, const void *userdata)</pre> The entries of the matrix P , whose i -th column is the product $\nabla_{xx} c_i(x)v$ between $\nabla_{xx} c_i(x)$, the Hessian of the i -th component of the residual $c(x)$ at $x=x$, and $v=v$ must be returned in <i>pval</i> and the function return value set to 0. If the evaluation is impossible at x , return should be set to a nonzero value. Data may be passed into <i>eval_hprods</i> via the structure <i>userdata</i> .

Examples

[nlst.c](#), and [nlstf.c](#).

3.1.2.10 nls_terminate()

```
void nls_terminate (
    void ** data,
    struct nls_control_type * control,
    struct nls_inform_type * inform )
```

Deallocate all internal private storage

Parameters

<code>in, out</code>	<i>data</i>	holds private internal data
<code>out</code>	<i>control</i>	is a struct containing control information (see nls_control_type)
<code>out</code>	<i>inform</i>	is a struct containing output information (see nls_inform_type)

Examples

[nlst.c](#), and [nlstf.c](#).

Chapter 4

Example Documentation

4.1 nlst.c

This is an example of how to use the package both when the matrices (Jacobian, Hessian, residual-Hessians-vector product) are directly available and when their product with vectors may be found. Both function call evaluations and returns to the calling program to find the required values are illustrated. A variety of supported Hessian storage formats are shown.

Notice that C-style indexing is used, and that this is flagged by setting `control.f_indexing` to `false`. In addition, see how parameters may be passed into the evaluation functions via `userdata`.

```
/* nlst.c */
/* Full test for the NLS C interface using C sparse matrix indexing */
/* Jari Fowkes & Nick Gould, STFC-Rutherford Appleton Laboratory, 2021 */
#include <stdio.h>
#include <math.h>
#include "nls.h"
#define max(a,b) \
({ __typeof__ (a) _a = (a); \
  __typeof__ (b) _b = (b); \
  _a > _b ? _a : _b; })
// Custom userdata struct
struct userdata_type {
    double p;
};
// Function prototypes
int res( int n, int m, const double x[], double c[], const void * );
int jac( int n, int m, int jne, const double x[], double jval[], const void * );
int hess( int n, int m, int hne, const double x[], const double y[],
         double hval[], const void * );
int jacprod( int n, int m, const double x[], const bool transpose, double u[],
            const double v[], bool got_j, const void * );
int hessprod( int n, int m, const double x[], const double y[], double u[],
            const double v[], bool got_h, const void * );
int rhessprods( int n, int m, int pne, const double x[], const double v[],
              double pval[], bool got_h, const void * );
int scale( int n, int m, const double x[], double u[],
          const double v[], const void * );
int jac_dense( int n, int m, int jne, const double x[], double jval[],
             const void * );
int hess_dense( int n, int m, int hne, const double x[], const double y[],
             double hval[], const void * );
int rhessprods_dense( int n, int m, int pne, const double x[],
                    const double v[], double pval[], bool got_h,
                    const void * );

int main(void) {
    // Derived types
    void *data;
    struct nls_control_type control;
    struct nls_inform_type inform;
    // Set user data
    struct userdata_type userdata;
```

```

userdata.p = 1.0;
// Set problem data
int n = 2; // # variables
int m = 3; // # residuals
int j_ne = 5; // Jacobian elements
int h_ne = 2; // Hesssian elements
int p_ne = 2; // residual-Hessians-vector products elements
int J_row[] = {0, 1, 1, 2, 2}; // Jacobian J
int J_col[] = {0, 0, 1, 0, 1}; //
int J_ptr[] = {0, 1, 3, 5}; // row pointers
int H_row[] = {0, 1}; // Hessian H
int H_col[] = {0, 1}; // NB lower triangle
int H_ptr[] = {0, 1, 2}; // row pointers
int P_row[] = {0, 1}; // residual-Hessians-vector product matrix
int P_ptr[] = {0, 1, 2, 2}; // column pointers
// Set storage
double g[n]; // gradient
double c[m]; // residual
double y[m]; // multipliers
double W[] = {1.0, 1.0, 1.0}; // weights
char st;
int status;
printf(" C sparse matrix indexing\n\n");
printf(" tests options for all-in-one storage format\n\n");
for( int d=1; d <= 5; d++){
// for( int d=5; d <= 5; d++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );
    // Set user-defined control options
    control.f_indexing = false; // C sparse matrix indexing
    //control.print_level = 1;
    control.jacobian_available = 2;
    control.hessian_available = 2;
    control.model = 6;
    double x[] = {1.5,1.5}; // starting point
    double W[] = {1.0, 1.0, 1.0}; // weights
    switch(d){
        case 1: // sparse co-ordinate storage
            st = 'C';
            nls_import( &control, &data, &status, n, m,
                "coordinate", j_ne, J_row, J_col, NULL,
                "coordinate", h_ne, H_row, H_col, NULL,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
            nls_solve_with_mat( &data, &userdata, &status,
                n, m, x, c, g, res, j_ne, jac,
                h_ne, hess, p_ne, rhessprods );

            break;
        case 2: // sparse by rows
            st = 'R';
            nls_import( &control, &data, &status, n, m,
                "sparse_by_rows", j_ne, NULL, J_col, J_ptr,
                "sparse_by_rows", h_ne, NULL, H_col, H_ptr,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
            nls_solve_with_mat( &data, &userdata, &status,
                n, m, x, c, g, res, j_ne, jac,
                h_ne, hess, p_ne, rhessprods );

            break;
        case 3: // dense
            st = 'D';
            nls_import( &control, &data, &status, n, m,
                "dense", j_ne, NULL, NULL, NULL,
                "dense", h_ne, NULL, NULL, NULL,
                "dense", p_ne, NULL, NULL, NULL, W );
            nls_solve_with_mat( &data, &userdata, &status,
                n, m, x, c, g, res, j_ne, jac_dense,
                h_ne, hess_dense, p_ne, rhessprods_dense );

            break;
        case 4: // diagonal
            st = 'I';
            nls_import( &control, &data, &status, n, m,
                "sparse_by_rows", j_ne, NULL, J_col, J_ptr,
                "diagonal", h_ne, NULL, NULL, NULL,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
            nls_solve_with_mat( &data, &userdata, &status,
                n, m, x, c, g, res, j_ne, jac,
                h_ne, hess, p_ne, rhessprods );

            break;
        case 5: // access by products
            st = 'P';
            nls_import( &control, &data, &status, n, m,
                "absent", j_ne, NULL, NULL, NULL,
                "absent", h_ne, NULL, NULL, NULL,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
            nls_solve_without_mat( &data, &userdata, &status,
                n, m, x, c, g, res, jacprod,
                hessprod, p_ne, rhessprods );

            break;
    }
}

```

```

    }
    nls_information( &data, &inform, &status );
    if(inform.status == 0){
        printf("%c:%i iterations. Optimal objective value = %5.2f"
            " status = %li\n",
            st, inform.iter, inform.obj, inform.status);
    }else{
        printf("%c: NLS_solve exit status = %li\n", st, inform.status);
    }
    // Delete internal workspace
    nls_terminate( &data, &control, &inform );
}
printf("\n tests reverse-communication options\n\n");
// reverse-communication input/output
int eval_status;
double u[max(m,n)], v[max(m,n)];
double J_val[j_ne], J_dense[m*n];
double H_val[h_ne], H_dense[n*(n+1)/2], H_diag[n];
double P_val[p_ne], P_dense[m*n];
bool transpose;
bool got_j = false;
bool got_h = false;
for( int d=1; d <= 5; d++){
// for( int d=1; d <= 4; d++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );
    // Set user-defined control options
    control.f_indexing = false; // C sparse matrix indexing
    //control.print_level = 1;
    control.jacobian_available = 2;
    control.hessian_available = 2;
    control.model = 6;
    double x[] = {1.5,1.5}; // starting point
    double W[] = {1.0, 1.0, 1.0}; // weights
    switch(d){
        case 1: // sparse co-ordinate storage
            st = 'C';
            nls_import( &control, &data, &status, n, m,
                "coordinate", j_ne, J_row, J_col, NULL,
                "coordinate", h_ne, H_row, H_col, NULL,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
            while(true){ // reverse-communication loop
                nls_solve_reverse_with_mat( &data, &status, &eval_status,
                    n, m, x, c, g, j_ne, J_val, y,
                    h_ne, H_val, v, p_ne, P_val );
                if(status == 0){ // successful termination
                    break;
                }else if(status < 0){ // error exit
                    break;
                }else if(status == 2){ // evaluate c
                    eval_status = res( n, m, x, c, &userdata );
                }else if(status == 3){ // evaluate J
                    eval_status = jac( n, m, j_ne, x, J_val, &userdata );
                }else if(status == 4){ // evaluate H
                    eval_status = hess( n, m, h_ne, x, y, H_val, &userdata );
                }else if(status == 7){ // evaluate P
                    eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                        got_h, &userdata );
                }else{
                    printf(" the value %li of status should not occur\n",
                        status);
                    break;
                }
            }
            break;
        case 2: // sparse by rows
            st = 'R';
            nls_import( &control, &data, &status, n, m,
                "sparse_by_rows", j_ne, NULL, J_col, J_ptr,
                "sparse_by_rows", h_ne, NULL, H_col, H_ptr,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
            while(true){ // reverse-communication loop
                nls_solve_reverse_with_mat( &data, &status, &eval_status,
                    n, m, x, c, g, j_ne, J_val, y,
                    h_ne, H_val, v, p_ne, P_val );
                if(status == 0){ // successful termination
                    break;
                }else if(status < 0){ // error exit
                    break;
                }else if(status == 2){ // evaluate c
                    eval_status = res( n, m, x, c, &userdata );
                }else if(status == 3){ // evaluate J
                    eval_status = jac( n, m, j_ne, x, J_val, &userdata );
                }else if(status == 4){ // evaluate H
                    eval_status = hess( n, m, h_ne, x, y, H_val, &userdata );
                }else if(status == 7){ // evaluate P
                    eval_status = rhessprods( n, m, p_ne, x, v, P_val,

```

```

                                got_h, &userdata );
    }else{
        printf(" the value %li of status should not occur\n",
            status);
        break;
    }
}
break;
case 3: // dense
    st = 'D';
    nls_import( &control, &data, &status, n, m,
        "dense", j_ne, NULL, NULL, NULL,
        "dense", h_ne, NULL, NULL, NULL,
        "dense", p_ne, NULL, NULL, NULL, W );
    while(true){ // reverse-communication loop
        nls_solve_reverse_with_mat( &data, &status, &eval_status,
            n, m, x, c, g, j_ne, J_dense, y,
            h_ne, H_dense, v, p_ne, P_dense );
        if(status == 0){ // successful termination
            break;
        }else if(status < 0){ // error exit
            break;
        }else if(status == 2){ // evaluate c
            eval_status = res( n, m, x, c, &userdata );
        }else if(status == 3){ // evaluate J
            eval_status = jac_dense( n, m, j_ne, x, J_dense,
                &userdata );
        }else if(status == 4){ // evaluate H
            eval_status = hess_dense( n, m, h_ne, x, y, H_dense,
                &userdata );
        }else if(status == 7){ // evaluate P
            eval_status = rhessprods_dense( n, m, p_ne, x, v, P_dense,
                got_h, &userdata );
        }else{
            printf(" the value %li of status should not occur\n",
                status);
            break;
        }
    }
    break;
case 4: // diagonal
    st = 'I';
    nls_import( &control, &data, &status, n, m,
        "sparse_by_rows", j_ne, NULL, J_col, J_ptr,
        "diagonal", h_ne, NULL, NULL, NULL,
        "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
    while(true){ // reverse-communication loop
        nls_solve_reverse_with_mat( &data, &status, &eval_status,
            n, m, x, c, g, j_ne, J_val, y,
            h_ne, H_diag, v, p_ne, P_val );
        if(status == 0){ // successful termination
            break;
        }else if(status < 0){ // error exit
            break;
        }else if(status == 2){ // evaluate c
            eval_status = res( n, m, x, c, &userdata );
        }else if(status == 3){ // evaluate J
            eval_status = jac( n, m, j_ne, x, J_val, &userdata );
        }else if(status == 4){ // evaluate H
            eval_status = hess( n, m, h_ne, x, y, H_diag, &userdata );
        }else if(status == 7){ // evaluate P
            eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                got_h, &userdata );
        }else{
            printf(" the value %li of status should not occur\n",
                status);
            break;
        }
    }
    break;
case 5: // access by products
    st = 'P';
    // control.print_level = 1;
    nls_import( &control, &data, &status, n, m,
        "absent", j_ne, NULL, NULL, NULL,
        "absent", h_ne, NULL, NULL, NULL,
        "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
    while(true){ // reverse-communication loop
        nls_solve_reverse_without_mat( &data, &status, &eval_status,
            n, m, x, c, g, &transpose,
            u, v, y, p_ne, P_val );
        if(status == 0){ // successful termination
            break;
        }else if(status < 0){ // error exit
            break;
        }else if(status == 2){ // evaluate c
            eval_status = res( n, m, x, c, &userdata );

```

```

    }else if(status == 5){ // evaluate u + J v or u + J'v
        eval_status = jacprod( n, m, x, transpose, u, v, got_j,
                               &userdata );
    }else if(status == 6){ // evaluate u + H v
        eval_status = hessprod( n, m, x, y, u, v, got_h,
                                &userdata );
    }else if(status == 7){ // evaluate P
        eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                                   got_h, &userdata );
    }else{
        printf(" the value %li of status should not occur\n",
               status);
        break;
    }
}
break;
}
nls_information( &data, &inform, &status );
if(inform.status == 0){
    printf("%c:%6i iterations. Optimal objective value = %5.2f"
           " status = %li\n",
           st, inform.iter, inform.obj, inform.status);
}else{
    printf("%c: NLS_solve exit status = %li\n", st, inform.status);
}
// Delete internal workspace
nls_terminate( &data, &control, &inform );
}
printf("\n basic tests of models used, direct access\n\n");
for( int model=3; model <= 8; model++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );
    // Set user-defined control options
    control.f_indexing = false; // C sparse matrix indexing
    //control.print_level = 1;
    control.jacobian_available = 2;
    control.hessian_available = 2;
    control.model = model;
    double x[] = {1.5,1.5}; // starting point
    double W[] = {1.0, 1.0, 1.0}; // weights
    nls_import( &control, &data, &status, n, m,
                "sparse_by_rows", j_ne, NULL, J_col, J_ptr,
                "sparse_by_rows", h_ne, NULL, H_col, H_ptr,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
    nls_solve_with_mat( &data, &userdata, &status,
                        n, m, x, c, g, res, j_ne, jac,
                        h_ne, hess, p_ne, rhessprods );
    nls_information( &data, &inform, &status );
    if(inform.status == 0){
        printf(" %li:%6i iterations. Optimal objective value = %5.2f"
               " status = %li\n",
               model, inform.iter, inform.obj, inform.status);
    }else{
        printf(" %i: NLS_solve exit status = %li\n", model, inform.status);
    }
    // Delete internal workspace
    nls_terminate( &data, &control, &inform );
}
printf("\n basic tests of models used, access by products\n\n");
for( int model=3; model <= 8; model++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );
    // Set user-defined control options
    control.f_indexing = false; // C sparse matrix indexing
    //control.print_level = 1;
    control.jacobian_available = 2;
    control.hessian_available = 2;
    control.model = model;
    double x[] = {1.5,1.5}; // starting point
    double W[] = {1.0, 1.0, 1.0}; // weights
    nls_import( &control, &data, &status, n, m,
                "absent", j_ne, NULL, NULL, NULL,
                "absent", h_ne, NULL, NULL, NULL,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
    nls_solve_without_mat( &data, &userdata, &status,
                           n, m, x, c, g, res, jacprod,
                           hessprod, p_ne, rhessprods );
    nls_information( &data, &inform, &status );
    if(inform.status == 0){
        printf("P%li:%6i iterations. Optimal objective value = %5.2f"
               " status = %li\n",
               model, inform.iter, inform.obj, inform.status);
    }else{
        printf("P%i: NLS_solve exit status = %li\n", model, inform.status);
    }
    // Delete internal workspace
    nls_terminate( &data, &control, &inform );
}

```

```

}
printf("\n basic tests of models used, reverse access\n\n");
for( int model=3; model <= 8; model++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );
    // Set user-defined control options
    control.f_indexing = false; // C sparse matrix indexing
    //control.print_level = 1;
    control.jacobian_available = 2;
    control.hessian_available = 2;
    control.model = model;
    double x[] = {1.5,1.5}; // starting point
    double W[] = {1.0, 1.0, 1.0}; // weights
    nls_import( &control, &data, &status, n, m,
        "sparse_by_rows", j_ne, NULL, J_col, J_ptr,
        "sparse_by_rows", h_ne, NULL, H_col, H_ptr,
        "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
    while(true){ // reverse-communication loop
        nls_solve_reverse_with_mat( &data, &status, &eval_status,
            n, m, x, c, g, j_ne, J_val, y,
            h_ne, H_val, v, p_ne, P_val );
        if(status == 0){ // successful termination
            break;
        }else if(status < 0){ // error exit
            break;
        }else if(status == 2){ // evaluate c
            eval_status = res( n, m, x, c, &userdata );
        }else if(status == 3){ // evaluate J
            eval_status = jac( n, m, j_ne, x, J_val, &userdata );
        }else if(status == 4){ // evaluate H
            eval_status = hess( n, m, h_ne, x, y, H_val, &userdata );
        }else if(status == 7){ // evaluate P
            eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                got_h, &userdata );
        }else{
            printf(" the value %li of status should not occur\n",
                status);
            break;
        }
    }
    nls_information( &data, &inform, &status );
    if(inform.status == 0){
        printf("P%li:%6i iterations. Optimal objective value = %5.2f"
            " status = %li\n",
            model, inform.iter, inform.obj, inform.status);
    }else{
        printf(" %i: NLS_solve exit status = %li\n", model, inform.status);
    }
    // Delete internal workspace
    nls_terminate( &data, &control, &inform );
}
printf("\n basic tests of models used, reverse access by products\n\n");
for( int model=3; model <= 8; model++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );
    // Set user-defined control options
    control.f_indexing = false; // C sparse matrix indexing
    //control.print_level = 1;
    control.jacobian_available = 2;
    control.hessian_available = 2;
    control.model = model;
    double x[] = {1.5,1.5}; // starting point
    double W[] = {1.0, 1.0, 1.0}; // weights
    nls_import( &control, &data, &status, n, m,
        "absent", j_ne, NULL, NULL, NULL,
        "absent", h_ne, NULL, NULL, NULL,
        "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
    while(true){ // reverse-communication loop
        nls_solve_reverse_without_mat( &data, &status, &eval_status,
            n, m, x, c, g, &transpose,
            u, v, y, p_ne, P_val );
        if(status == 0){ // successful termination
            break;
        }else if(status < 0){ // error exit
            break;
        }else if(status == 2){ // evaluate c
            eval_status = res( n, m, x, c, &userdata );
        }else if(status == 5){ // evaluate u + J'v or u + J'v
            eval_status = jacprod( n, m, x, transpose, u, v, got_j,
                &userdata );
        }else if(status == 6){ // evaluate u + H v
            eval_status = hessprod( n, m, x, y, u, v, got_h,
                &userdata );
        }else if(status == 7){ // evaluate P
            eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                got_h, &userdata );
        }else{

```

```

        printf(" the value %li of status should not occur\n",
               status);
        break;
    }
}
nls_information( &data, &inform, &status );
if(inform.status == 0){
    printf("P%i:%6i iterations. Optimal objective value = %5.2f"
           " status = %li\n",
           model, inform.iter, inform.obj, inform.status);
}else{
    printf("P%i: NLS_solve exit status = %li\n", model, inform.status);
}
// Delete internal workspace
nls_terminate( &data, &control, &inform );
}

// compute the residuals
int res( int n, int m, const double x[], double c[], const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    double p = myuserdata->p;
    c[0] = pow(x[0],2.0) + p;
    c[1] = x[0] + pow(x[1],2.0);
    c[2] = x[0] - x[1];
    return 0;
}

// compute the Jacobian
int jac( int n, int m, int jne, const double x[], double jval[],
        const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    jval[0] = 2.0 * x[0];
    jval[1] = 1.0;
    jval[2] = 2.0 * x[1];
    jval[3] = 1.0;
    jval[4] = - 1.0;
    return 0;
}

// compute the Hessian
int hess( int n, int m, int hne, const double x[], const double y[],
         double hval[], const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    hval[0] = 2.0 * y[0];
    hval[1] = 2.0 * y[1];
    return 0;
}

// compute Jacobian-vector products
int jacprod( int n, int m, const double x[], const bool transpose, double u[],
            const double v[], bool got_j, const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    if (transpose) {
        u[0] = u[0] + 2.0 * x[0] * v[0] + v[1] + v[2];
        u[1] = u[1] + 2.0 * x[1] * v[1] - v[2];
    }else{
        u[0] = u[0] + 2.0 * x[0] * v[0];
        u[1] = u[1] + v[0] + 2.0 * x[1] * v[1];
        u[2] = u[2] + v[0] - v[1];
    }
    return 0;
}

// compute Hessian-vector products
int hessprod( int n, int m, const double x[], const double y[], double u[],
            const double v[], bool got_h, const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    u[0] = u[0] + 2.0 * y[0] * v[0];
    u[1] = u[1] + 2.0 * y[1] * v[1];
    return 0;
}

// compute residual-Hessians-vector products
int rhessprods( int n, int m, int pne, const double x[], const double v[],
               double pval[], bool got_h, const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    pval[0] = 2.0 * v[0];
    pval[1] = 2.0 * v[1];
    return 0;
}

// scale v
int scale( int n, int m, const double x[], double u[],
          const double v[], const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    u[0] = v[0];
    u[1] = v[1];
    return 0;
}

// compute the dense Jacobian
int jac_dense( int n, int m, int jne, const double x[], double jval[],
              const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;

```

```

    jval[0] = 2.0 * x[0];
    jval[1] = 0.0;
    jval[2] = 1.0;
    jval[3] = 2.0 * x[1];
    jval[4] = 1.0;
    jval[5] = - 1.0;
    return 0;
}
// compute the dense Hessian
int hess_dense( int n, int m, int hne, const double x[], const double y[],
               double hval[], const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    hval[0] = 2.0 * y[0];
    hval[1] = 0.0;
    hval[2] = 2.0 * y[1];
    return 0;
}
// compute dense residual-Hessians-vector products
int rhessprods_dense( int n, int m, int pne, const double x[],
                     const double v[], double pval[], bool got_h,
                     const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    pval[0] = 2.0 * v[0];
    pval[1] = 0.0;
    pval[2] = 0.0;
    pval[3] = 2.0 * v[1];
    pval[4] = 0.0;
    pval[5] = 0.0;
    return 0;
}

```

4.2 nlstf.c

This is the same example, but now fortran-style indexing is used.

```

/* nlstf.c */
/* Full test for the NLS interface using Fortran sparse matrix indexing */
/* Jari Fowkes & Nick Gould, STFC-Rutherford Appleton Laboratory, 2021 */
#include <stdio.h>
#include <math.h>
#include "nls.h"
#define max(a,b) \
    ( ( __typeof__ (a) )_a = (a); \
      __typeof__ (b) )_b = (b); \
      _a > _b ? _a : _b; )
// Custom userdata struct
struct userdata_type {
    double p;
};
// Function prototypes
int res( int n, int m, const double x[], double c[], const void * );
int jac( int n, int m, int jne, const double x[], double jval[], const void * );
int hess( int n, int m, int hne, const double x[], const double y[],
         double hval[], const void * );
int jacprod( int n, int m, const double x[], const bool transpose, double u[],
            const double v[], bool got_j, const void * );
int hessprod( int n, int m, const double x[], const double y[], double u[],
            const double v[], bool got_h, const void * );
int rhessprods( int n, int m, int pne, const double x[], const double v[],
              double pval[], bool got_h, const void * );
int scale( int n, int m, const double x[], double u[],
          const double v[], const void * );
int jac_dense( int n, int m, int jne, const double x[], double jval[],
              const void * );
int hess_dense( int n, int m, int hne, const double x[], const double y[],
              double hval[], const void * );
int rhessprods_dense( int n, int m, int pne, const double x[],
                    const double v[], double pval[], bool got_h,
                    const void * );

int main(void) {
    // Derived types
    void *data;
    struct nls_control_type control;
    struct nls_inform_type inform;
    // Set user data
    struct userdata_type userdata;
    userdata.p = 1.0;
    // Set problem data
    int n = 2; // # variables
    int m = 3; // # residuals
    int j_ne = 5; // Jacobian elements
}

```



```

int h_ne = 2; // Hesssian elements
int p_ne = 2; // residual-Hessians-vector products elements
int J_row[] = {1, 2, 2, 3, 3}; // Jacobian J
int J_col[] = {1, 1, 2, 1, 2}; //
int J_ptr[] = {1, 2, 4, 6}; // row pointers
int H_row[] = {1, 2}; // Hessian H
int H_col[] = {1, 2}; // NB lower triangle
int H_ptr[] = {1, 2, 3}; // row pointers
int P_row[] = {1, 2}; // residual-Hessians-vector product matrix
int P_ptr[] = {1, 2, 3, 3}; // column pointers
// Set storage
double g[n]; // gradient
double c[m]; // residual
double y[m]; // multipliers
double W[] = {1.0, 1.0, 1.0}; // weights
char st;
int status;
printf(" Fortran sparse matrix indexing\n\n");
printf(" tests options for all-in-one storage format\n\n");
for( int d=1; d <= 5; d++){
// for( int d=5; d <= 5; d++){
// Initialize NLS
nls_initialize( &data, &control, &inform );
// Set user-defined control options
control.f_indexing = true; // Fortran sparse matrix indexing
//control.print_level = 1;
control.jacobian_available = 2;
control.hessian_available = 2;
control.model = 6;
double x[] = {1.5,1.5}; // starting point
double W[] = {1.0, 1.0, 1.0}; // weights
switch(d){
case 1: // sparse co-ordinate storage
st = 'C';
nls_import( &control, &data, &status, n, m,
"coordinate", j_ne, J_row, J_col, NULL,
"coordinate", h_ne, H_row, H_col, NULL,
"sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
nls_solve_with_mat( &data, &userdata, &status,
n, m, x, c, g, res, j_ne, jac,
h_ne, hess, p_ne, rhessprods );
break;
case 2: // sparse by rows
st = 'R';
nls_import( &control, &data, &status, n, m,
"sparse_by_rows", j_ne, NULL, J_col, J_ptr,
"sparse_by_rows", h_ne, NULL, H_col, H_ptr,
"sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
nls_solve_with_mat( &data, &userdata, &status,
n, m, x, c, g, res, j_ne, jac,
h_ne, hess, p_ne, rhessprods );
break;
case 3: // dense
st = 'D';
nls_import( &control, &data, &status, n, m,
"dense", j_ne, NULL, NULL, NULL,
"dense", h_ne, NULL, NULL, NULL,
"dense", p_ne, NULL, NULL, NULL, W );
nls_solve_with_mat( &data, &userdata, &status,
n, m, x, c, g, res, j_ne, jac_dense,
h_ne, hess_dense, p_ne, rhessprods_dense );
break;
case 4: // diagonal
st = 'I';
nls_import( &control, &data, &status, n, m,
"sparse_by_rows", j_ne, NULL, J_col, J_ptr,
"diagonal", h_ne, NULL, NULL, NULL,
"sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
nls_solve_with_mat( &data, &userdata, &status,
n, m, x, c, g, res, j_ne, jac,
h_ne, hess, p_ne, rhessprods );
break;
case 5: // access by products
st = 'P';
nls_import( &control, &data, &status, n, m,
"absent", j_ne, NULL, NULL, NULL,
"absent", h_ne, NULL, NULL, NULL,
"sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
nls_solve_without_mat( &data, &userdata, &status,
n, m, x, c, g, res, jacprod,
hessprod, p_ne, rhessprods );
break;
}
nls_information( &data, &inform, &status );
if(inform.status == 0){
printf("%c:%6i iterations. Optimal objective value = %5.2f"
" status = %li\n",

```

```

        st, inform.iter, inform.obj, inform.status);
    }else{
        printf("%c: NLS_solve exit status = %li\n", st, inform.status);
    }
    // Delete internal workspace
    nls_terminate( &data, &control, &inform );
}
printf("\n tests reverse-communication options\n\n");
// reverse-communication input/output
int eval_status;
double u[max(m,n)], v[max(m,n)];
double J_val[j_ne], J_dense[m*n];
double H_val[h_ne], H_dense[n*(n+1)/2], H_diag[n];
double P_val[p_ne], P_dense[m*n];
bool transpose;
bool got_j = false;
bool got_h = false;
for( int d=1; d <= 5; d++){
// for( int d=1; d <= 4; d++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );
    // Set user-defined control options
    control.f_indexing = true; // Fortran sparse matrix indexing
    //control.print_level = 1;
    control.jacobian_available = 2;
    control.hessian_available = 2;
    control.model = 6;
    double x[] = {1.5,1.5}; // starting point
    double W[] = {1.0, 1.0, 1.0}; // weights
    switch(d){
        case 1: // sparse co-ordinate storage
            st = 'C';
            nls_import( &control, &data, &status, n, m,
                "coordinate", j_ne, J_row, J_col, NULL,
                "coordinate", h_ne, H_row, H_col, NULL,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
            while(true){ // reverse-communication loop
                nls_solve_reverse_with_mat( &data, &status, &eval_status,
                    n, m, x, c, g, j_ne, J_val, y,
                    h_ne, H_val, v, p_ne, P_val );
                if(status == 0){ // successful termination
                    break;
                }else if(status < 0){ // error exit
                    break;
                }else if(status == 2){ // evaluate c
                    eval_status = res( n, m, x, c, &userdata );
                }else if(status == 3){ // evaluate J
                    eval_status = jac( n, m, j_ne, x, J_val, &userdata );
                }else if(status == 4){ // evaluate H
                    eval_status = hess( n, m, h_ne, x, y, H_val, &userdata );
                }else if(status == 7){ // evaluate P
                    eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                        got_h, &userdata );
                }else{
                    printf(" the value %li of status should not occur\n",
                        status);
                    break;
                }
            }
            break;
        case 2: // sparse by rows
            st = 'R';
            nls_import( &control, &data, &status, n, m,
                "sparse_by_rows", j_ne, NULL, J_col, J_ptr,
                "sparse_by_rows", h_ne, NULL, H_col, H_ptr,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
            while(true){ // reverse-communication loop
                nls_solve_reverse_with_mat( &data, &status, &eval_status,
                    n, m, x, c, g, j_ne, J_val, y,
                    h_ne, H_val, v, p_ne, P_val );
                if(status == 0){ // successful termination
                    break;
                }else if(status < 0){ // error exit
                    break;
                }else if(status == 2){ // evaluate c
                    eval_status = res( n, m, x, c, &userdata );
                }else if(status == 3){ // evaluate J
                    eval_status = jac( n, m, j_ne, x, J_val, &userdata );
                }else if(status == 4){ // evaluate H
                    eval_status = hess( n, m, h_ne, x, y, H_val, &userdata );
                }else if(status == 7){ // evaluate P
                    eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                        got_h, &userdata );
                }else{
                    printf(" the value %li of status should not occur\n",
                        status);
                    break;
                }
            }

```

```

    }
}
break;
case 3: // dense
st = 'D';
nls_import( &control, &data, &status, n, m,
            "dense", j_ne, NULL, NULL, NULL,
            "dense", h_ne, NULL, NULL, NULL,
            "dense", p_ne, NULL, NULL, NULL, W );
while(true){ // reverse-communication loop
    nls_solve_reverse_with_mat( &data, &status, &eval_status,
                               n, m, x, c, g, j_ne, J_dense, y,
                               h_ne, H_dense, v, p_ne, P_dense );
    if(status == 0){ // successful termination
        break;
    }else if(status < 0){ // error exit
        break;
    }else if(status == 2){ // evaluate c
        eval_status = res( n, m, x, c, &userdata );
    }else if(status == 3){ // evaluate J
        eval_status = jac_dense( n, m, j_ne, x, J_dense,
                                &userdata );
    }else if(status == 4){ // evaluate H
        eval_status = hess_dense( n, m, h_ne, x, y, H_dense,
                                &userdata );
    }else if(status == 7){ // evaluate P
        eval_status = rhessprods_dense( n, m, p_ne, x, v, P_dense,
                                       got_h, &userdata );
    }else{
        printf(" the value %li of status should not occur\n",
               status);
        break;
    }
}
break;
case 4: // diagonal
st = 'I';
nls_import( &control, &data, &status, n, m,
            "sparse_by_rows", j_ne, NULL, J_col, J_ptr,
            "diagonal", h_ne, NULL, NULL, NULL,
            "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
while(true){ // reverse-communication loop
    nls_solve_reverse_with_mat( &data, &status, &eval_status,
                               n, m, x, c, g, j_ne, J_val, y,
                               h_ne, H_diag, v, p_ne, P_val );
    if(status == 0){ // successful termination
        break;
    }else if(status < 0){ // error exit
        break;
    }else if(status == 2){ // evaluate c
        eval_status = res( n, m, x, c, &userdata );
    }else if(status == 3){ // evaluate J
        eval_status = jac( n, m, j_ne, x, J_val, &userdata );
    }else if(status == 4){ // evaluate H
        eval_status = hess( n, m, h_ne, x, y, H_diag, &userdata );
    }else if(status == 7){ // evaluate P
        eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                                  got_h, &userdata );
    }else{
        printf(" the value %li of status should not occur\n",
               status);
        break;
    }
}
break;
case 5: // access by products
st = 'P';
control.print_level = 1;
nls_import( &control, &data, &status, n, m,
            "absent", j_ne, NULL, NULL, NULL,
            "absent", h_ne, NULL, NULL, NULL,
            "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
while(true){ // reverse-communication loop
    nls_solve_reverse_without_mat( &data, &status, &eval_status,
                                   n, m, x, c, g, &transpose,
                                   u, v, y, p_ne, P_val );
    if(status == 0){ // successful termination
        break;
    }else if(status < 0){ // error exit
        break;
    }else if(status == 2){ // evaluate c
        eval_status = res( n, m, x, c, &userdata );
    }else if(status == 5){ // evaluate u + J v or u + J'v
        eval_status = jacprod( n, m, x, transpose, u, v, got_j,
                               &userdata );
    }else if(status == 6){ // evaluate u + H v
        eval_status = hessprod( n, m, x, y, u, v, got_h,

```

```

                                &userdata );
    }else if(status == 7){ // evaluate P
        eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                                got_h, &userdata );
    }else{
        printf(" the value %li of status should not occur\n",
            status);
        break;
    }
}
break;
}
nls_information( &data, &inform, &status );
if(inform.status == 0){
    printf("%c:%6i iterations. Optimal objective value = %5.2f"
        " status = %li\n",
        st, inform.iter, inform.obj, inform.status);
}else{
    printf("%c: NLS_solve exit status = %li\n", st, inform.status);
}
// Delete internal workspace
nls_terminate( &data, &control, &inform );
}
printf("\n basic tests of models used, direct access\n\n");
for( int model=3; model <= 8; model++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );
    // Set user-defined control options
    control.f_indexing = true; // Fortran sparse matrix indexing
    //control.print_level = 1;
    control.jacobian_available = 2;
    control.hessian_available = 2;
    control.model = model;
    double x[] = {1.5,1.5}; // starting point
    double W[] = {1.0, 1.0, 1.0}; // weights
    nls_import( &control, &data, &status, n, m,
        "sparse_by_rows", j_ne, NULL, J_col, J_ptr,
        "sparse_by_rows", h_ne, NULL, H_col, H_ptr,
        "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
    nls_solve_with_mat( &data, &userdata, &status,
        n, m, x, c, g, res, j_ne, jac,
        h_ne, hess, p_ne, rhessprods );
    nls_information( &data, &inform, &status );
    if(inform.status == 0){
        printf(" %li:%6i iterations. Optimal objective value = %5.2f"
            " status = %li\n",
            model, inform.iter, inform.obj, inform.status);
    }else{
        printf(" %i: NLS_solve exit status = %li\n", model, inform.status);
    }
}
// Delete internal workspace
nls_terminate( &data, &control, &inform );
}
printf("\n basic tests of models used, access by products\n\n");
for( int model=3; model <= 8; model++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );
    // Set user-defined control options
    control.f_indexing = true; // Fortran sparse matrix indexing
    //control.print_level = 1;
    control.jacobian_available = 2;
    control.hessian_available = 2;
    control.model = model;
    double x[] = {1.5,1.5}; // starting point
    double W[] = {1.0, 1.0, 1.0}; // weights
    nls_import( &control, &data, &status, n, m,
        "absent", j_ne, NULL, NULL, NULL,
        "absent", h_ne, NULL, NULL, NULL,
        "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
    nls_solve_without_mat( &data, &userdata, &status,
        n, m, x, c, g, res, jacprod,
        hessprod, p_ne, rhessprods );
    nls_information( &data, &inform, &status );
    if(inform.status == 0){
        printf("P%i:%6i iterations. Optimal objective value = %5.2f"
            " status = %li\n",
            model, inform.iter, inform.obj, inform.status);
    }else{
        printf("P%i: NLS_solve exit status = %li\n", model, inform.status);
    }
}
// Delete internal workspace
nls_terminate( &data, &control, &inform );
}
printf("\n basic tests of models used, reverse access\n\n");
for( int model=3; model <= 8; model++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );

```

```

// Set user-defined control options
control.f_indexing = true; // Fortran sparse matrix indexing
//control.print_level = 1;
control.jacobian_available = 2;
control.hessian_available = 2;
control.model = model;
double x[] = {1.5,1.5}; // starting point
double W[] = {1.0, 1.0, 1.0}; // weights
nls_import( &control, &data, &status, n, m,
            "sparse_by_rows", j_ne, NULL, J_col, J_ptr,
            "sparse_by_rows", h_ne, NULL, H_col, H_ptr,
            "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
while(true){ // reverse-communication loop
    nls_solve_reverse_with_mat( &data, &status, &eval_status,
                               n, m, x, c, g, j_ne, J_val, y,
                               h_ne, H_val, v, p_ne, P_val );

    if(status == 0){ // successful termination
        break;
    }else if(status < 0){ // error exit
        break;
    }else if(status == 2){ // evaluate c
        eval_status = res( n, m, x, c, &userdata );
    }else if(status == 3){ // evaluate J
        eval_status = jac( n, m, j_ne, x, J_val, &userdata );
    }else if(status == 4){ // evaluate H
        eval_status = hess( n, m, h_ne, x, y, H_val, &userdata );
    }else if(status == 7){ // evaluate P
        eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                                   got_h, &userdata );
    }else{
        printf(" the value %li of status should not occur\n",
               status);
        break;
    }
}
nls_information( &data, &inform, &status );
if(inform.status == 0){
    printf("P%i:%i iterations. Optimal objective value = %5.2f"
           " status = %li\n",
           model, inform.iter, inform.obj, inform.status);
}else{
    printf(" %i: NLS_solve exit status = %li\n", model, inform.status);
}
// Delete internal workspace
nls_terminate( &data, &control, &inform );
}
printf("\n basic tests of models used, reverse access by products\n\n");
for( int model=3; model <= 8; model++){
    // Initialize NLS
    nls_initialize( &data, &control, &inform );
    // Set user-defined control options
    control.f_indexing = true; // Fortran sparse matrix indexing
    //control.print_level = 1;
    control.jacobian_available = 2;
    control.hessian_available = 2;
    control.model = model;
    double x[] = {1.5,1.5}; // starting point
    double W[] = {1.0, 1.0, 1.0}; // weights
    nls_import( &control, &data, &status, n, m,
                "absent", j_ne, NULL, NULL, NULL,
                "absent", h_ne, NULL, NULL, NULL,
                "sparse_by_columns", p_ne, P_row, NULL, P_ptr, W );
    while(true){ // reverse-communication loop
        nls_solve_reverse_without_mat( &data, &status, &eval_status,
                                       n, m, x, c, g, &transpose,
                                       u, v, y, p_ne, P_val );

        if(status == 0){ // successful termination
            break;
        }else if(status < 0){ // error exit
            break;
        }else if(status == 2){ // evaluate c
            eval_status = res( n, m, x, c, &userdata );
        }else if(status == 5){ // evaluate u + J v or u + J'v
            eval_status = jacprod( n, m, x, transpose, u, v, got_j,
                                   &userdata );
        }else if(status == 6){ // evaluate u + H v
            eval_status = hessprod( n, m, x, y, u, v, got_h,
                                    &userdata );
        }else if(status == 7){ // evaluate P
            eval_status = rhessprods( n, m, p_ne, x, v, P_val,
                                       got_h, &userdata );
        }else{
            printf(" the value %li of status should not occur\n",
                   status);
            break;
        }
    }
}
}

```

```

    nls_information( &data, &inform, &status );
    if(inform.status == 0){
        printf("P%i:%i iterations. Optimal objective value = %5.2f"
            " status = %i\n",
            model, inform.iter, inform.obj, inform.status);
    }else{
        printf("P%i: NLS_solve exit status = %i\n", model, inform.status);
    }
    // Delete internal workspace
    nls_terminate( &data, &control, &inform );
}

// compute the residuals
int res( int n, int m, const double x[], double c[], const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    double p = myuserdata->p;
    c[0] = pow(x[0],2.0) + p;
    c[1] = x[0] + pow(x[1],2.0);
    c[2] = x[0] - x[1];
    return 0;
}

// compute the Jacobian
int jac( int n, int m, int jne, const double x[], double jval[],
    const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    jval[0] = 2.0 * x[0];
    jval[1] = 1.0;
    jval[2] = 2.0 * x[1];
    jval[3] = 1.0;
    jval[4] = - 1.0;
    return 0;
}

// compute the Hessian
int hess( int n, int m, int hne, const double x[], const double y[],
    double hval[], const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    hval[0] = 2.0 * y[0];
    hval[1] = 2.0 * y[1];
    return 0;
}

// compute Jacobian-vector products
int jacprod( int n, int m, const double x[], const bool transpose, double u[],
    const double v[], bool got_j, const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    if (transpose) {
        u[0] = u[0] + 2.0 * x[0] * v[0] + v[1] + v[2];
        u[1] = u[1] + 2.0 * x[1] * v[1] - v[2];
    }else{
        u[0] = u[0] + 2.0 * x[0] * v[0];
        u[1] = u[1] + v[0] + 2.0 * x[1] * v[1];
        u[2] = u[2] + v[0] - v[1];
    }
    return 0;
}

// compute Hessian-vector products
int hessprod( int n, int m, const double x[], const double y[], double u[],
    const double v[], bool got_h, const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    u[0] = u[0] + 2.0 * y[0] * v[0];
    u[1] = u[1] + 2.0 * y[1] * v[1];
    return 0;
}

// compute residual-Hessians-vector products
int rhessprods( int n, int m, int pne, const double x[], const double v[],
    double pval[], bool got_h, const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    pval[0] = 2.0 * v[0];
    pval[1] = 2.0 * v[1];
    return 0;
}

// scale v
int scale( int n, int m, const double x[], double u[],
    const double v[], const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    u[0] = v[0];
    u[1] = v[1];
    return 0;
}

// compute the dense Jacobian
int jac_dense( int n, int m, int jne, const double x[], double jval[],
    const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    jval[0] = 2.0 * x[0];
    jval[1] = 0.0;
    jval[2] = 1.0;
    jval[3] = 2.0 * x[1];
    jval[4] = 1.0;
}

```

```
    jval[5] = - 1.0;
    return 0;
}
// compute the dense Hessian
int hess_dense( int n, int m, int hne, const double x[], const double y[],
               double hval[], const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    hval[0] = 2.0 * y[0];
    hval[1] = 0.0;
    hval[2] = 2.0 * y[1];
    return 0;
}
// compute dense residual-Hessians-vector products
int rhessprods_dense( int n, int m, int pne, const double x[],
                    const double v[], double pval[], bool got_h,
                    const void *userdata ){
    struct userdata_type *myuserdata = ( struct userdata_type * ) userdata;
    pval[0] = 2.0 * v[0];
    pval[1] = 0.0;
    pval[2] = 0.0;
    pval[3] = 2.0 * v[1];
    pval[4] = 0.0;
    pval[5] = 0.0;
    return 0;
}
```


Index

- nls.h, [9](#)
 - nls_import, [19](#)
 - nls_information, [21](#)
 - nls_initialize, [22](#)
 - nls_read_specfile, [22](#)
 - nls_reset_control, [22](#)
 - nls_solve_reverse_with_mat, [23](#)
 - nls_solve_reverse_without_mat, [26](#)
 - nls_solve_with_mat, [28](#)
 - nls_solve_without_mat, [31](#)
 - nls_terminate, [34](#)
- nls_control_type, [12](#)
- nls_import
 - nls.h, [19](#)
- nls_inform_type, [18](#)
- nls_information
 - nls.h, [21](#)
- nls_initialize
 - nls.h, [22](#)
- nls_read_specfile
 - nls.h, [22](#)
- nls_reset_control
 - nls.h, [22](#)
- nls_solve_reverse_with_mat
 - nls.h, [23](#)
- nls_solve_reverse_without_mat
 - nls.h, [26](#)
- nls_solve_with_mat
 - nls.h, [28](#)
- nls_solve_without_mat
 - nls.h, [31](#)
- nls_subproblem_control_type, [10](#)
- nls_subproblem_inform_type, [17](#)
- nls_terminate
 - nls.h, [34](#)
- nls_time_type, [17](#)