GALAHAD SLLS

1 SUMMARY

This package uses a preconditioned, projected-gradient method to solve the **simplex-constrained regularized linear least-squares problem**

$$\text{minimize } q(\mathbf{x}) = \tfrac{1}{2}\|\mathbf{A}\mathbf{x} - \mathbf{b}\|_W^2 + \tfrac{1}{2}\sigma\|\mathbf{x}\|_2^2,$$

where the variables $\mathbf{x}$ are required to lie within the **regular simplex**

$$\mathbf{e}^T\mathbf{x} = 1 \text{ and } \mathbf{x} \geq 0,$$

where the m by n real matrix $\mathbf{A}$, the vector $\mathbf{b}$ and the non-negative weights $\mathbf{w}$ and σ are given, where $\mathbf{e}$ is the vector of ones, and where the Euclidean and weighted-Euclidean norms are given by $\|\mathbf{v}\|_2^2 = \mathbf{v}^T\mathbf{v}$ and $\|\mathbf{v}\|_W^2 = \mathbf{v}^T\mathbf{W}\mathbf{v}$, respectively, with $\mathbf{W} = \text{diag}(\mathbf{w})$. Full advantage is taken of any zero coefficients of the Jacobian matrix $\mathbf{A}$ of the **residuals** $\mathbf{c}(\mathbf{x}) = \mathbf{A}\mathbf{x} - \mathbf{b}$; the matrix need not be provided as there are options to obtain matrix-vector products involving $\mathbf{A}$ and its transpose either by reverse communication or from a user-provided subroutine.

ATTRIBUTES — Versions: `GALAHAD_SLLS_single`, `GALAHAD_SLLS_double`. **Uses:** `GALAHAD_CPU_time`, `GALAHAD-_SYMBOLS`, `GALAHAD_SPACE`, `GALAHAD_STRING`, `GALAHAD_SORT`, `GALAHAD_NORMS`, `GALAHAD_CONVERT`, `GALAHAD_SBLS`, `GALAHAD_QPT`, `GALAHAD_QPD`, `GALAHAD_USERDATA`. `GALAHAD_SPECFILE`. **Date:** July 2022. **Origin:** N. I. M. Gould, STFC-Rutherford Appleton Laboratory. **Language:** Fortran 95 + TR 15581 or Fortran 2003.

2 HOW TO USE THE PACKAGE

Access to the package requires a `USE` statement such as

Single precision version

        USE GALAHAD_SLLS_single

Double precision version

        USE GALAHAD_SLLS_double

If it is required to use both modules at the same time, the derived types `SMT_type`, `QPT_problem_type`, `NLPT_userdata-_type`, `SLLS_time_type`, `SLLS_control_type`, `SLLS_inform_type` and `SLLS_data_type` (Section 2.2) and the sub-routines `SLLS_initialize`, `SLLS_solve`, `SLLS_terminate`, (Section 2.3) and `SLLS_read_specfile` (Section 2.7) must be renamed on one of the `USE` statements.

2.1 Matrix storage formats

When it is explicitly available, the coefficient matrix $\mathbf{A}$ may be stored in a variety of input formats.

2.1.1 Dense row-wise storage format

The matrix $\mathbf{A}$ is stored as a compact dense matrix by rows, that is, the values of the entries of each row in turn are stored in order within an appropriate real one-dimensional array. Component $n * (i-1) + j$ of the storage array `A%val` will hold the value $\mathbf{A}_{i,j}$ for $i = 1, \ldots, m$, $j = 1, \ldots, n$.

2.1.2 Dense column-wise storage format

The matrix $\mathbf{A}$ is stored as a compact dense matrix by columns, that is, the values of the entries of each column in turn are stored in order within an appropriate real one-dimensional array. Component $m * (j - 1) + i$ of the storage array A%val will hold the value $\mathbf{A}_{i,j}$ for $i = 1, \ldots, m$, $j = 1, \ldots, n$.

2.1.3 Sparse coordinate storage format

Only the nonzero entries of the matrices are stored. For the l-th entry of $\mathbf{A}$, its row index i, column index j and value $\mathbf{A}_{ij}$ are stored in the l-th components of the integer arrays A%row, A%col and real array A%val. The order is unimportant, but the total number of entries A%ne is required.

2.1.4 Sparse row-wise storage format

Again only the nonzero entries are stored, but this time they are ordered so that those in row i appear directly before those in row $i + 1$. For the i-th row of $\mathbf{A}$, the i-th component of the integer array A%ptr holds the position of the first entry in this row, while A%ptr $(m + 1)$ holds the total number of entries plus one. The column indices j and values $\mathbf{A}_{ij}$ of the entries in the i-th row are stored in components $l = $ A%ptr$(i), \ldots,$A%ptr $(i + 1) - 1$ of the integer array A%col, and real array A%val, respectively.

2.1.5 Sparse column-wise storage format

Yet again only the nonzero entries are stored, but this time they are ordered so that those in column j appear directly before those in column $j + 1$. For the j-th column of $\mathbf{A}$, the j-th component of the integer array A%ptr holds the position of the first entry in this column, while A%ptr $(n + 1)$ holds the total number of entries plus one. The row indices i and values $\mathbf{A}_{ij}$ of the entries in the j-th column are stored in components $l = $ A%ptr$(j), \ldots,$A%ptr $(j + 1) - 1$ of the integer array A%row, and real array A%val, respectively.

For sparse matrices, the row- and column-wise storage schemes almost always requires less storage than their predecessor.

2.2 The derived data types

Ten derived data types are accessible from the package.

2.2.1 The derived data type for holding matrices

The derived data type SMT_TYPE is used to hold the matrix $\mathbf{A}$. The components of SMT_TYPE used here are:

m   is a scalar component of type default INTEGER, that holds the number of rows in the matrix.

n   is a scalar component of type default INTEGER, that holds the number of columns in the matrix.

ne   is a scalar variable of type default INTEGER, that holds the number of matrix entries.

type   is a rank-one allocatable array of type default CHARACTER, that is used to indicate the matrix storage scheme used. Its precise length and content depends on the type of matrix to be stored (see §2.2.2).

val   is a rank-one allocatable array of type default REAL (double precision in GALAHAD_SLLS_double) and dimension at least ne, that holds the values of the entries. Any duplicated entries that appear in the sparse coordinate or row-wise schemes will be summed.

---

row   is a rank-one allocatable array of type default `INTEGER`, and dimension at least `ne`, that may hold the row indices of the entries. (see §2.1.3 and §2.1.5).

col   is a rank-one allocatable array of type default `INTEGER`, and dimension at least `ne`, that may the column indices of the entries (see §2.1.3–2.1.4).

ptr   is a rank-one allocatable array of type default `INTEGER`, and dimension at least `m + 1`, that may hold the pointers to the first entry in each row (see §2.1.4). or dimension at least `n + 1`, that may hold the pointers to the first entry in each column (see §2.1.5).

2.2.2 The derived data type for holding the problem

The derived data type `QPT_problem_type` is used to hold the problem. The components of `QPT_problem_type` are:

n   is a scalar variable of type default `INTEGER`, that holds the number of optimization variables, n.

m   is a scalar variable of type default `INTEGER`, that holds the number of residuals, m.

A   is scalar variable of type `SMT_TYPE` that holds the Jacobian matrix **A** of the residuals (if it is available). The following components are used here:

   A%type is an allocatable array of rank one and type default `CHARACTER`, that is used to indicate the storage scheme used. If the dense row-wise storage scheme (see Section 2.1.1) is used, the first twelve components of A%type must contain the string `DENSE_BY_ROWS`, while if the column-wise scheme (see Section 2.1.2) is used, the fifteen components of A%type must contain the string `DENSE_BY_COLUMNS`. For the sparse coordinate scheme (see Section 2.1.3), the first ten components of A%type must contain the string `COORDINATE`, for the sparse row-wise storage scheme (see Section 2.1.4), the first fourteen components of A%type must contain the string `SPARSE_BY_ROWS`. and for the sparse column-wise storage scheme (see Section 2.1.5), the first seventeen components of A%type must contain the string `SPARSE_BY_COLUMNS`.

   For convenience, the procedure `SMT_put` may be used to allocate sufficient space and insert the required keyword into A%type. For example, if nlp is of derived type `SLLS_problem_type` and involves a Jacobian we wish to store using the coordinate scheme, we may simply

   ```
   CALL SMT_put( nlp%A%type, 'COORDINATE' )
   ```

   See the documentation for the GALAHAD package `SMT` for further details on the use of `SMT_put`.

   A%type should **not** be allocated if **A** is unavailable (and access provided to **A** provided by other means, see later sections), and in this case the remaining components of A need not be set.

   A%ne is a scalar variable of type default `INTEGER`, that holds the number of entries in **A** in the sparse coordinate storage scheme (see Section 2.1.3). It need not be set for any of the other three schemes.

   A%val is a rank-one allocatable array of type default `REAL` (double precision in `GALAHAD_SLLS_double`), that holds the values of the entries of the Jacobian matrix **A** in any of the storage schemes discussed in Section 2.1.

   A%row is a rank-one allocatable array of type default `INTEGER`, that holds the row indices of **A** in the sparse coordinate storage scheme (see Section 2.1.3). It need not be allocated for any of the other three schemes.

   A%col is a rank-one allocatable array variable of type default `INTEGER`, that holds the column indices of **A** in either the sparse coordinate (see Section 2.1.3), or the sparse row-wise (see Section 2.1.4) storage scheme. It need not be allocated when the dense or column-wise storage schemes are used.

   A%row is a rank-one allocatable array variable of type default `INTEGER`, that holds the row indices of **A** in either the sparse coordinate (see Section 2.1.3), or the sparse column-wise (see Section 2.1.5) storage scheme. It need not be allocated when the dense or row-wise storage schemes are used.

A%ptr is a rank-one allocatable array and type default INTEGER, that must be of dimension m+1 and hold the starting position of each row of **A**, as well as the total number of entries plus one, in the sparse row-wise storage scheme (see Section 2.1.4). If the sparse column-wise storage scheme (see Section 2.1.5) is used, it must instead be of dimension n+1 and hold the starting position of each column of **A**, as well as the total number of entries plus one. It need not be allocated when the other schemes are used.

B   is a rank-one allocatable array of dimension m and type default REAL (double precision in GALAHAD_SLLS_-double), that holds the constant term **b** in the residuals. The i-th component of B, $i = 1, \dots, m$, contains $\mathbf{b}_i$.

C   is a rank-one allocatable array of dimension m and type default REAL (double precision in GALAHAD_SLLS_-double), that holds the residuals, $\mathbf{c}(\mathbf{x}) = \mathbf{A}\mathbf{x} - \mathbf{b}$, at the point **x**. The i-th component of C, $i = 1, \dots, m$, contains $\mathbf{c}_i(\mathbf{x})$.

G   is a rank-one allocatable array of dimension n and type default REAL (double precision in GALAHAD_SLLS_-double), that holds the gradient of the objective $\mathbf{g}(\mathbf{x}) = \mathbf{A}^T\mathbf{c}(\mathbf{x})$, at the point **x**. The j-th component of G, $j = 1, \dots, n$, contains $\mathbf{g}_j(\mathbf{x})$.

X   is a rank-one allocatable array of dimension n and type default REAL (double precision in GALAHAD_SLLS_-double), that holds the values **x** of the optimization variables. The j-th component of X, $j = 1, \dots, n$, contains x_j.

Z   is a rank-one allocatable array of dimension n and type default REAL (double precision in GALAHAD_SLLS_-double), that holds the values **z** of estimates of the dual variables corresponding to the non-negativity constraints (see Section 4). The j-th component of Z, $j = 1, \dots, n$, contains z_j.

2.2.3 The derived data type for holding control parameters

The derived data type SLLS_control_type is used to hold controlling data. Default values may be obtained by calling SLLS_initialize (see Section 2.3.1), while components may also be changed by calling SLLS_read_specfile (see Section 2.7.1). The components of SLLS_control_type are:

error is a scalar variable of type default INTEGER, that holds the stream number for error messages. Printing of error messages in SLLS_solve and SLLS_terminate is suppressed if error $\leq$ 0. The default is error = 6.

out   is a scalar variable of type default INTEGER, that holds the stream number for informational messages. Printing of informational messages in SLLS_solve is suppressed if out < 0. The default is out = 6.

print_level is a scalar variable of type default INTEGER, that is used to control the amount of informational output which is required. No informational output will occur if print_level $\leq$ 0. If print_level = 1, a single line of output will be produced for each iteration of the process. If print_level $\geq$ 2, this output will be increased to provide significant detail of each iteration. The default is print_level = 0.

start_print is a scalar variable of type default INTEGER, that specifies the first iteration for which printing will be permitted in GALAHAD_SLLS_solve. If start_print is negative, printing will be permitted from the outset. The default is start_print = -1.

stop_print is a scalar variable of type default INTEGER, that specifies the last iteration for which printing will be permitted in GALAHAD_SLLS_solve. If stop_print is negative, printing will be permitted once it has been started by start_print. The default is stop_print = -1.

print_gap is a scalar variable of type default INTEGER. Once printing has been started, output will occur once every print_gap iterations. If print_gap is no larger than 1, printing will be permitted on every iteration. The default is print_gap = 1.

maxit is a scalar variable of type default INTEGER, that holds the maximum number of iterations which will be allowed in GALAHAD_SLLS_solve. The default is maxit = 1000.

cold_start is a scalar variable of type default INTEGER, that should be set to 0 if a warm start is required (with variables assigned according to X_stat, see below), and to any other value if the values given in prob%X suffice. The default is cold_start = 1.

preconditioner is a scalar variable of type default INTEGER, that specifies the type of preconditioner (scaling) used when computing the search direction. Currently this can be 0 (no preconditioner), 1 (a diagonal preconditioner that normalises the rows of **A**) or 2 (a preconditioner supplied by the user either via a subroutine call of eval_PREC, see §2.4.4, or via reverse communication, see §2.5), although other values may be introduced in future. The default is preconditioner = 1.

change_max is a scalar variable of type default INTEGER, that specifies the maximum number of per-iteration changes in the working set permitted when allowing subspace solution rather than steepest descent (see §4). The default is change_max = 2.

cg_maxit is a scalar variable of type default INTEGER, that holds the maximum number of conjugate-gradient iterations which will be allowed per main iteration in GALAHAD_SLLS_solve. The default is cg_maxit = 1000, and any negative value will be interpreted as $n + 1$.

arcsearch_max_steps is a scalar variable of type default INTEGER, that specifies the maximum number of steps allowed in an individual piecewise arc search. The default is arcsearch_max_steps = 1000, and a negative value removes the limit.

weight is a scalar variable of type default REAL (double precision in GALAHAD_SLLS_double), that is used to specify the weight σ that controls regularization of the objective function. Any value smaller than 0 will be regarded as zero. The default is weight = 0.0.

stop_d is a scalar variable of type default REAL (double precision in GALAHAD_SLLS_double), that holds the required accuracy for the dual infeasibility (see Section 4). The default is stop_d = $u^{1/3}$, where u is EPSILON(1.0) (EPSILON(1.0D0) in GALAHAD_SLLS_double).

stop_cg_relative and stop_cg_absolute are scalar variables of type default REAL (double precision in GALAHAD_SLLS_double), that hold the relative and absolute convergence tolerances for the conjugate-gradient iteration that occurs in the face of currently-active constraints when constructing the search direction. _stop_cg_relative = 0.01 and stop_cg_absolute = $\sqrt{u}$, where u is EPSILON(1.0) (EPSILON(1.0D0) in GALAHAD_SLLS_double).

alpha_max is a scalar variable of type default REAL (double precision in GALAHAD_SLLS_double), that specifies The default is alpha_max = 10^{20}.

alpha_initial is a scalar variable of type default REAL (double precision in GALAHAD_SLLS_double), that specifies the initial arc length during the inexact piecewise arc search. The default is alpha_initial = 1.0.

alpha_reduction is a scalar variable of type default REAL (double precision in GALAHAD_SLLS_double), that specifies the arc length reduction factor for the inexact piecewise arc search. The default is alpha_reduction = 0.5.

arcsearch_acceptance_tol is a scalar variable of type default REAL (double precision in GALAHAD_SLLS-_double), that specifies the required relative reduction during the inexact arc search The default is arcsearch_acceptance_tol = 0.01.

stabilisation_weight is a scalar variable of type default REAL (double precision in GALAHAD_SLLS_double), that specifies the weight added to the search-direction subproblem to stabilise the computation. The default is stabilisation_weight = 10^{-12}.

---

cpu_time_limit is a scalar variable of type default REAL (double precision in GALAHAD_SLLS_double), that is used to specify the maximum permitted CPU time. Any negative value indicates no limit will be imposed. The default is cpu_time_limit = - 1.0.

direct_subproblem_solve is a scalar variable of type default LOGICAL, that should be set .TRUE. if the search over a promising subspace is carried out using matrix factorization, and .FALSE. if the conjugate-gradient least-squares (CGLS) method is to be preferred. The former is generally more effective so long as the cost of factorization isn't exorbitant. The default is direct_subproblem_solve = .TRUE., but the package will override this if the Jacobian is not explicitly available.

exact_arc_search is a scalar variable of type default LOGICAL, that should be set .TRUE. if the exact minimizer along the search arc is required, and .FALSE. if an approximation found by backtracking or advancing suffices. The default is exact_arc_search = .TRUE..

space_critical is a scalar variable of type default LOGICAL, that must be set .TRUE. if space is critical when allocating arrays and .FALSE. otherwise. The package may run faster if space_critical is .FALSE. but at the possible expense of a larger storage requirement. The default is space_critical = .FALSE..

deallocate_error_fatal is a scalar variable of type default LOGICAL, that must be set .TRUE. if the user wishes to terminate execution if a deallocation fails, and .FALSE. if an attempt to continue will be made. The default is deallocate_error_fatal = .FALSE..

prefix is a scalar variable of type default CHARACTER and length 30, that may be used to provide a user-selected character string to preface every line of printed output. Specifically, each line of output will be prefaced by the string prefix(2:LEN(TRIM(prefix))-1), thus ignoring the first and last non-null components of the supplied string. If the user does not want to preface lines by such a string, they may use the default prefix = "".

SBLS_control is a scalar variable of type SBLS_control_type whose components are used to control the factorization and/or preconditioner used, performed by the package GALAHAD_SBLS. See the documentation for GALAHAD_SBLS for further details.

2.2.4 The derived data type for holding timing information

The derived data type SLLS_time_type is used to hold elapsed CPU times for the various parts of the calculation. The components of SLLS_time_type are:

total is a scalar variable of type default REAL, that gives the total time spent in the package.

analyse is a scalar variable of type default REAL, that gives the time spent analysing the required matrices prior to factorization.

factorize is a scalar variable of type default REAL, that gives the time spent factorizing the required matrices.

solve is a scalar variable of type default REAL, that gives the time spent computing the search direction.

2.2.5 The derived data type for holding informational parameters

The derived data type SLLS_inform_type is used to hold parameters that give information about the progress and needs of the algorithm. The components of SLLS_inform_type are:

status is a scalar variable of type default INTEGER, that gives the exit status of the algorithm. See Section 2.6 for details.

alloc_status is a scalar variable of type default INTEGER, that gives the status of the last attempted array allocation or deallocation. This will be 0 if status = 0.

bad_alloc is a scalar variable of type default CHARACTER and length 80, that gives the name of the last internal array for which there were allocation or deallocation errors. This will be the null string if status = 0.

factorization_status is a scalar variable of type default INTEGER, that gives the return status from the matrix factorization.

iter is a scalar variable of type default INTEGER, that gives the number of iterations performed.

obj is a scalar variable of type default REAL (double precision in GALAHAD_SLLS_double), that holds the value of the objective function at the best estimate of the solution found.

time is a scalar variable of type SLLS_time_type whose components are used to hold elapsed CPU times for the various parts of the calculation (see Section 2.2.4).

SBLS_inform is a scalar variable of type SBLS_inform_type whose components provide information about the progress and needs of the factorization/preconditioner performed by the package GALAHAD_SBLS. See the documentation for GALAHAD_SBLS for further details.

2.2.6 The derived data type for holding problem data

The derived data type SLLS_data_type is used to hold all the data for a particular problem, or sequences of problems with the same structure, between calls of SLLS procedures. This data should be preserved, untouched, from the initial call to SLLS_initialize to the final call to SLLS_terminate.

2.2.7 The derived data type for holding user data

The derived data type NLPT_userdata_type is available to allow the user to pass data to and from user-supplied matrix-vector product and preconditioning subroutines (see Section 2.4). Components of variables of type NLPT_userdata-_type may be allocated as necessary. The following components are available:

integer is a rank-one allocatable array of type default INTEGER.

real is a rank-one allocatable array of type default REAL (double precision in GALAHAD_SLLS_double)

complex is a rank-one allocatable array of type default COMPLEX (double precision complex in GALAHAD_SLLS_-double).

character is a rank-one allocatable array of type default CHARACTER.

logical is a rank-one allocatable array of type default LOGICAL.

integer_pointer is a rank-one pointer array of type default INTEGER.

real_pointer is a rank-one pointer array of type default REAL (double precision in GALAHAD_SLLS_double)

complex_pointer is a rank-one pointer array of type default COMPLEX (double precision complex in GALAHAD_SLLS_-double).

character_pointer is a rank-one pointer array of type default CHARACTER.

logical_pointer is a rank-one pointer array of type default LOGICAL.

2.2.8 The derived data type for holding reverse-communication data

The derived data type `SLLS_reverse_type` is used to hold data needed for reverse communication when this is required. The components of `SLLS_reverse_type` are:

`nz_v_start` is a scalar variable of type default `INTEGER`, that may be used to hold the starting position in `NZ_v` (see below) of the list of indices of nonzero components of **v**.

`nz_v_end` is a scalar variable of type default `INTEGER`, that may be used to hold the finishing position in `NZ_v` (see below) of the list of indices of nonzero components of **v**.

`NZ_v` is a rank-one allocatable array of dimension n and type default `INTEGER`, that may be used to hold the indices of the nonzero components of **v**. If used, components `NZ_v(nz_v_start:nz_v_end)` of V (see below) will be nonzero.

`V` is a rank-one allocatable array of dimension n and type default `REAL` (double precision in `GALAHAD_SLLS_double`), that is used to hold the components of the output vector **v**.

`P` is a rank-one allocatable array of dimension m and type default `REAL` (double precision in `GALAHAD_SLLS_double`), that is used to record the components of the resulting vector **Av**.

`nz_p_end` is a scalar variable of type default `INTEGER`, that is used to record the finishing position in `NZ_p` (see below) of the list of indices of nonzero components of **Av** if required.

`NZ_p` is a rank-one allocatable array of dimension n and type default `INTEGER`, that is used to record the list of indices of nonzero components of **Av** if required. Components `NZ_p(1:nz_prod_end)` of P should then be nonzero.

2.3 Argument lists and calling sequences

There are three procedures for user calls (see Section 2.7 for further features):

1. The subroutine `SLLS_initialize` is used to set default values, and initialize private data, before solving one or more problems with the same sparsity and bound structure.

2. The subroutine `SLLS_solve` is called to solve the problem.

3. The subroutine `SLLS_terminate` is provided to allow the user to automatically deallocate array components of the private data, allocated by `SLLS_solve`, at the end of the solution process.

We use square brackets `[ ]` to indicate `OPTIONAL` arguments.

2.3.1 The initialization subroutine

Default values are provided as follows:

```
CALL SLLS_initialize( data, control )
```

`data` is a scalar `INTENT(INOUT)` argument of type `SLLS_data_type` (see Section 2.2.6). It is used to hold data about the problem being solved.

`control` is a scalar `INTENT(OUT)` argument of type `SLLS_control_type` (see Section 2.2.3). On exit, `control` contains default values for the components as described in Section 2.2.3. These values should only be changed after calling `SLLS_initialize`.

2.3.2 The bound-constrained linear least-squares subroutine

The bound-constrained linear least-squares solution algorithm is called as follows:

```
CALL SLLS_solve( prob, X_stat, data, control, inform, userdata[, W, reverse, &
                 eval_APROD, eval_ASPROD, eval_AFPROD, eval_PREC] )
```

prob   is a scalar `INTENT(INOUT)` argument of type `QPT_problem_type` (see Section 2.2.2). It is used to hold data about the problem being solved. The user must allocate and set values for the array components, and set values for the component `prob%B`. Additionally, the user can provide **A** by allocating the relevant array components and setting values for `prob%A` using whichever of the matrix formats described in Section 2.1 is appropriate for the user's application; if the effect of **A** and its transpose are only available to form products via reverse communication (see `reverse` below) or with a set of user-supplied subroutines (see `eval_APROD eval_ASPROD` and `eval_AFPROD` below), `prob%A` is not needed.

    The components `prob%X` must be set to initial estimates of the primal variables, **x**. Inappropriate initial values will be altered, so the user should not be overly concerned if suitable values are not apparent, and may be content with merely setting `prob%X=0.0`. The components `prob%C`, `prob%G` and `prob%Z` need not be set or allocated on entry.

    On exit, the components `prob%X` and `prob%Z` will contain the best estimates of the primal variables **x**, and dual variables for the bound constraints **z**, respectively. The components `prob%C` and `prob%G` will contain the residuals $\mathbf{c}(\mathbf{x})$ and gradients $\mathbf{g}(\mathbf{x})$ at **x**, respectively. **Restrictions:** `prob%n` > 0, `prob%m` > 0 and (if **A** is provided) `prob%A%ne` ≥ 0. `prob%A_type` $\in \{$'DENSE_BY_ROWS', 'DENSE_BY_COLUMNS', 'COORDINATE', 'SPARSE_BY_ROWS', SPARSE_BY_COLUMNS $\}$.

X_stat is a rank-one `INTENT(INOUT)` array argument of dimension `prob%n` and type default `INTEGER`, that indicates which of the simple bound constraints are in the current working set. Possible values for `X_stat(j)`, j= 1,..., `prob%n`, and their meanings are

    `<0` the *j*-th simple bound constraint is in the working set, on its lower bound,

    `>0` the *j*-th simple bound constraint is in the working set, on its upper bound, and

    `0` the *j*-th simple bound constraint is not in the working set.

    Suitable values must be supplied if `control%slls_control%cold_start = 0` on entry, but need not be provided for other input values of `control%cold_start`. Inappropriate values will be ignored. On exit, `X_stat` will contain values appropriate for the ultimate working set.

data   is a scalar `INTENT(INOUT)` argument of type `SLLS_data_type` (see Section 2.2.6). It is used to hold data about the problem being solved. It must not have been altered **by the user** since the last call to `SLLS_initialize`.

control is a scalar `INTENT(IN)` argument of type `SLLS_control_type` (see Section 2.2.3). Default values may be assigned by calling `SLLS_initialize` prior to the first call to `SLLS_solve`.

inform is a scalar `INTENT(INOUT)` argument of type `SLLS_inform_type` (see Section 2.2.5). On initial entry, the component `status` must be set to 1, while other components need not be set. A successful call to `SLLS_solve` is indicated when the component `status` has the value 0. For other return values of `status`, see Sections 2.5 and 2.6.

userdata is a scalar `INTENT(INOUT)` argument of type `NLPT_userdata_type` whose components may be used to communicate user-supplied data (see Section 2.2.7) to and from the `OPTIONAL` subroutines `eval_APROD` and `eval_ASPROD` (see below).

W   is an `OPTIONAL` rank-one `INTENT(INOUT)` array argument of dimension `prob%n` and type default `REAL` (double precision in `GALAHAD_SLLS_double`) that if present contains the weights **w**. If W is absent, weights of one will be used instead.

---

reverse is an OPTIONAL scalar INTENT(INOUT) argument of type SLLS_reverse_type (see Section 2.2.8). It is used to communicate reverse-communication data between the subroutine and calling program. If reverse is PRESENT and eval%APROD or eval%ASPROD (see below) are absent, the user should monitor inform%status on exit (see Section 2.5).

eval_APROD is an OPTIONAL user-supplied subroutine whose purpose is to evaluate the value of the product $\mathbf{p} + \mathbf{Av}$ or $\mathbf{p} + \mathbf{A}^T\mathbf{v}$ involving the Jacobian (and its transpose) and given vectors $\mathbf{v}$ and $\mathbf{p}$. See Section 2.4.1 for details. If eval_APROD is present, it must be declared EXTERNAL in the calling program. If eval_APROD is absent, GALAHAD_SLLS_solve will use reverse communication (see Section 2.5) to obtain Jacobian-vector products if reverse is PRESENT or otherwise require that the user has provided all relevant components of prob%A.

eval_ASPROD is an OPTIONAL user-supplied subroutine whose purpose is to evaluate the value of the product $\mathbf{Av}$ involving the Jacobian and a given *sparse* vector $\mathbf{v}$. See Section 2.4.2 for details. If eval_ASPROD is present, it must be declared EXTERNAL in the calling program. If eval_ASPROD is absent, GALAHAD_SLLS_solve will use reverse communication (see Section 2.5) to obtain Jacobian-sparse-vector products if reverse is PRESENT or otherwise require that the user has provided all relevant components of prob%A.

eval_AFPROD is an OPTIONAL user-supplied subroutine whose purpose is to evaluate the value of the product $\mathbf{Av}$ or $\mathbf{A}^T v$ involving the Jacobian (and its transpose) and a given vector $\mathbf{v}$ in which either only some components of $\mathbf{v}$ or the resulting product are set/required. See Section 2.4.3 for details. If eval_AFPROD is present, it must be declared EXTERNAL in the calling program. If eval_AFPROD is absent, GALAHAD_SLLS_solve will use reverse communication (see Section 2.5) to obtain Jacobian-sparse-vector products if reverse is PRESENT or otherwise require that the user has provided all relevant components of prob%A.

eval_PREC is an OPTIONAL user-supplied subroutine whose purpose is to evaluate the value of the product $\mathbf{P}^{-1}\mathbf{v}$ involving a symmetric, positive definite preconditioner $\mathbf{P}$ and a given vector $\mathbf{v}$. See Section 2.4.4 for details. If eval_PREC is present, it must be declared EXTERNAL in the calling program. If eval_PREC is absent, GALAHAD_SLLS_solve will use reverse communication (see Section 2.5) to obtain preconditioning products so long as reverse is PRESENT; if reverse is not PRESENT, no preconditioning will be performed.

2.3.3 The termination subroutine

All previously allocated arrays are deallocated as follows:

```
CALL SLLS_terminate( data, control, inform )
```

data is a scalar INTENT(INOUT) argument of type SLLS_data_type exactly as for SLLS_solve, which must not have been altered **by the user** since the last call to SLLS_initialize. On exit, array components will have been deallocated.

control is a scalar INTENT(IN) argument of type SLLS_control_type exactly as for SLLS_solve.

inform is a scalar INTENT(OUT) argument of type SLLS_inform_type exactly as for SLLS_solve. Only the component status will be set on exit, and a successful call to SLLS_terminate is indicated when this component status has the value 0. For other return values of status, see Section 2.6.

2.4 Matrix-vector operations

2.4.1 Jacobian-vector products via internal evaluation

If the argument eval_APROD is present when calling GALAHAD_SLLS_solve, the user is expected to provide a subroutine of that name to evaluate the sum $\mathbf{p} + \mathbf{Av}$ or $\mathbf{p} + \mathbf{A}^T\mathbf{v}$ involving the product of the residual Jacobian $\mathbf{A}$ or its transpose $\mathbf{A}^T$ and a given vector $\mathbf{v}$. The routine must be specified as

```
SUBROUTINE eval_APROD( status, userdata, transpose, V, P )
```

All use is subject to the conditions of a BSD-3-Clause License.
See http://galahad.rl.ac.uk/galahad-www/cou.html **for full details.**

whose arguments are as follows:

status is a scalar INTENT(OUT) argument of type default INTEGER, that should be set to 0 if the routine has been able to evaluate the sum $\mathbf{p} + \mathbf{A}\mathbf{v}$ or $\mathbf{p} + \mathbf{A}^T\mathbf{v}$ and to a non-zero value if the evaluation has not been possible.

userdata is a scalar INTENT(INOUT) argument of type NLPT_userdata_type whose components may be used to communicate user-supplied data to and from the subroutines eval_APROD (see Section 2.2.7).

transpose is a scalar INTENT(IN) array argument of type default that will be set .TRUE. if the product involves the transpose of the Jacobian $\mathbf{A}^T$ and .FALSE. if the product involves the Jacobian $\mathbf{A}$ itself.

V   is a rank-one INTENT(IN) array argument of type default REAL (double precision in GALAHAD_SLLS_double) whose components contain the vector $\mathbf{v}$.

P   is a rank-one INTENT(INOUT) array argument of type default REAL (double precision in GALAHAD_SLLS_double) whose components on input contain the vector $\mathbf{p}$ and on output the sum $\mathbf{p} + \mathbf{A}\mathbf{v}$ when %transpose is .FALSE. or $\mathbf{p} + \mathbf{A}^T\mathbf{v}$ when %transpose is .TRUE..

2.4.2 Jacobian-sparse-vector products via internal evaluation

If the argument eval_ASPROD is present when calling GALAHAD_SLLS_solve, the user is expected to provide a subroutine of that name to evaluate the product of the Jacobian $\mathbf{A}$ with a given vector $\mathbf{v}$. The routine must be specified as

```
SUBROUTINE eval_ASPROD( status, userdata, V, P[, NZ_in, nz_in_start, nz_in_end,    &
                        NZ_out, nz_out_end] )
```

whose arguments are as follows:

status is a scalar INTENT(OUT) argument of type default INTEGER, that should be set to 0 if the routine has been able to evaluate the product $\mathbf{A}\mathbf{v}$ and to a non-zero value if the evaluation has not been possible.

userdata is a scalar INTENT(INOUT) argument of type NLPT_userdata_type whose components may be used to communicate user-supplied data to and from the subroutine (see Section 2.2.7).

V   is a rank-one INTENT(IN) array argument of type default REAL (double precision in GALAHAD_SLLS_double) whose components contain the vector $\mathbf{v}$. If components nz_in_start, nz_in_end and NZ_in (see below) are PRESENT, only components NZ_in(nz_in_start:nz_in_end) of V will be nonzero and the remaining components of V should be ignored. Otherwise, all components of V should be presumed to be nonzero.

P   is a rank-one INTENT(OUT) array argument of type default REAL (double precision in GALAHAD_SLLS_double) whose components on output contain the required components of $\mathbf{A}\mathbf{v}$. If components nz_out_end and NZ_out (see below) are PRESENT, only the nonzero components NZ_out(1:nz_out_end) of P need be assigned. Otherwise, all components of P must be set.

nz_in_start is an OPTIONAL scalar variable of type default INTEGER, that, if PRESENT, holds the starting position in NZ_in of the list of indices of nonzero components of $\mathbf{v}$.

nz_in_end is an OPTIONAL scalar variable of type default INTEGER, that, if PRESENT, holds the finishing position in NZ_in of the list of indices of nonzero components of $\mathbf{v}$.

NZ_in is an OPTIONAL rank-one array of dimension n and type default INTEGER, that, if PRESENT, holds the indices of the nonzero components of $\mathbf{v}$. If any of nz_in_start, nz_in_end and NZ_in are absent, all components of V are assumed to be nonzero.

---

nz_out_end is an OPTIONAL scalar variable of type default INTEGER, that, if PRESENT, must be set to record the number of non-zeros in **Av**.

NZ_out is an OPTIONAL rank-one array of dimension *m* and type default INTEGER, that, if PRESENT, must be set to record the list of indices of nonzero components of **Av**. If either of nz_out_end and NZ_out are absent, all components of P should be set even if they are zero.

2.4.3   Jacobian-vector sub-products via internal evaluation

If the argument eval_AFPROD is present when calling GALAHAD_SLLS_solve, the user is expected to provide a subroutine of that name to evaluate the product of the Jacobian **A**, or its transpose, with a given vector **v**. Here, either only a subset of the components of the vector **v** are nonzero, or only a subset of the components product $\mathbf{A}^T\mathbf{v}$ are to required. The routine must be specified as

```
SUBROUTINE eval_AFPROD( status, userdata, transpose, V, P, FREE, n_free )
```

whose arguments are as follows:

status is a scalar INTENT(OUT) argument of type default INTEGER, that should be set to 0 if the routine has been able to evaluate the sum **Av** or $\mathbf{A}^T\mathbf{v}$ and to a non-zero value if the evaluation has not been possible.

userdata is a scalar INTENT(INOUT) argument of type NLPT_userdata_type whose components may be used to communicate user-supplied data to and from the subroutines eval_APROD (see Section 2.2.7).

transpose is a scalar INTENT(IN) array argument of type default that will be set .TRUE. if the product involves the transpose of the Jacobian $\mathbf{A}^T$ and .FALSE. if the product involves the Jacobian **A** itself.

V   is a rank-one INTENT(IN) array argument of type default REAL (double precision in GALAHAD_SLLS_double) whose components contain the vector **v**. If transpose is .FALSE. only those components whose indices FREE(:n_free) (see below) will by set, and the remainder will be presumed to be zero.

P   is a rank-one INTENT(OUT) array argument of type default REAL (double precision in GALAHAD_SLLS_double) whose components on output must be set to the product **Av** when %transpose is .FALSE.. If %transpose is .TRUE., components with indices FREE(:n_free) (see below) should be set to the corresponding components of the product $\mathbf{A}^T\mathbf{v}$, and the remaining components ignored.

FREE is a rank-one INTENT(IN) array argument of type default INTEGER that flags the input components of **v** that are set (when transpose is .FALSE.) or output components of $\mathbf{A}^T\mathbf{v}$ that are required (when transpose is .TRUE.). Specifically, only indices FREE(:n_free) of the relevant vector is set or required, and the remainder should be treated as zero (if transpose is .FALSE.) or ignored (if transpose is .TRUE.).

n_free is a scalar INTENT(IN) argument of type default INTEGER, that specifies the number of components of FREE that need be considered.

2.4.4   Application of a preconditioner via internal evaluation

If the argument eval_PREC is present when calling GALAHAD_SLLS_solve, the user is expected to provide a subroutine of that name to evaluate the product $\mathbf{P}^{-1}\mathbf{v}$ involving a symmetric, positive definite preconditioner **P** and a given *n*-vector **v**. Ideally the *n* by *n* matrix **P** should be chosen so that the of the product is inexpensive to compute, but also so that **P** is a good approximation of $\mathbf{A}^T\mathbf{A}$ in the sense that the eigenvalues of $\mathbf{P}^{-1}\mathbf{A}^T\mathbf{A}$ are clustered around a small number of values (preferably all around 1.0). The routine must be specified as

```
SUBROUTINE eval_PREC( status, userdata, V, P )
```

whose arguments are as follows:

`status` is a scalar `INTENT(OUT)` argument of type default `INTEGER`, that should be set to 0 if the routine has been able to evaluate the product $\mathbf{P}^{-1}\mathbf{v}$, and to a non-zero value if the evaluation has not been possible.

`userdata` is a scalar `INTENT(INOUT)` argument of type `NLPT_userdata_type` whose components may be used to communicate user-supplied data to and from the subroutines `eval_APROD` (see Section 2.2.7).

`V` is a rank-one `INTENT(IN)` array argument of type default `REAL` (double precision in `GALAHAD_SLLS_double`) whose components contain the vector $\mathbf{v}$.

`P` is a rank-one `INTENT(INOUT)` array argument of type default `REAL` (double precision in `GALAHAD_SLLS_double`) whose components on output contain the vector $\mathbf{p} = \mathbf{P}^{-1}\mathbf{v}$.

2.5 Reverse Communication Information

A positive value of `inform%status` on exit from `SLLS_solve` indicates that `GALAHAD_SLLS_solve` is seeking further information—this will happen if the user has chosen to evaluate matrix-vector products by reverse communication. The user should compute the required information and re-enter `GALAHAD_SLLS_solve` with `inform%status` and all other arguments (except those specifically mentioned below) unchanged.

Possible values of `inform%status` and the information required are

2. The user should compute the product $\mathbf{A}\mathbf{v}$ involving the product of the residual Jacobian, $\mathbf{A}$, with a given vector $\mathbf{v}$. The vector $\mathbf{v}$ is given in `reverse%V`, and the product $\mathbf{A}\mathbf{v}$ should be written to `reverse%P`, and `reverse%eval_status` should be set to 0. If the user is unable to evaluate the product, the user need not set `reverse%P`, but should then set `reverse%eval_status` to a non-zero value.

3. The user should compute the product $\mathbf{A}^T\mathbf{v}$ involving the product of the transpose of the residual Jacobian, $\mathbf{A}$, with a given vector $\mathbf{v}$. The vector $\mathbf{v}$ is given in `reverse%V`, and the product $\mathbf{A}^T\mathbf{v}$ should be written to `reverse%P`, and `reverse%eval_status` should be set to 0. If the user is unable to evaluate the product, the user need not set `reverse%P`, but should then set `reverse%eval_status` to a non-zero value.

4. The user should compute the matrix-vector product $\mathbf{A}\mathbf{v}$ using the vector $\mathbf{v}$ whose nonzero components are stored in positions `reverse%NZ_in(reverse%nz_in_start:reverse%nz_in_end)` of `reverse%V`. The remaining components of `reverse%V` should be ignored. The component `reverse%eval_status` should set to 0 unless the user is unable to evaluate the product, in which case `reverse%eval_status` should be set to a non-zero value, and the remaining components left unaltered.

5. The user should compute the nonzero components of the matrix-vector product $\mathbf{A}\mathbf{v}$ using the vector $\mathbf{v}$ whose nonzero components are stored in positions `reverse%NZ_in(reverse%nz_in_start:reverse%nz_in_end)` of `reverse%V`. The remaining components of `reverse%V` should be ignored. The nonzero components must occupy positions `reverse%NZ_out(1:reverse%nz_out_end)` of `reverse%P`, and the components `reverse%NZ_out` and `reverse%nz_out_end` must be set. The component `reverse%eval_status` should set to 0 unless the user is unable to evaluate the product, in which case `reverse%eval_status` should be set to a non-zero value, and the remaining components left unaltered.

6. The user should compute the matrix-vector product $\mathbf{A}^T\mathbf{v}$ using the vector $\mathbf{v}$ given in `reverse%V`. The components $(\mathbf{A}^T\mathbf{v})_j$ of the product $\mathbf{A}^T\mathbf{v}$, for `j = reverse%NZ_in(reverse%nz_in_start:reverse%nz_in_end)`, should be written to `reverse%P(j)`. The component `reverse%eval_status` should set to 0 unless the user is unable to evaluate the product, in which case `reverse%eval_status` should be set to a non-zero value, and the remaining components left unaltered.

7. The user should compute the product $\mathbf{P}^{-1}\mathbf{v}$ involving a symmetric, positive definite preconditioner $\mathbf{P}$ and a given vector $\mathbf{v}$ The vector $\mathbf{v}$ is given in `reverse%V`, and the product $\mathbf{P}^{-1}\mathbf{v}$ should be written to `reverse%P`, and `reverse%eval_status` should be set to 0. If the user is unable to evaluate the product, the user need not set `reverse%P`, but should then set `reverse%eval_status` to a non-zero value. This value of `inform%status` can only occur if the user has set `control%preconditioner = 2`.

2.6 Warning and error messages

A negative value of `inform%status` on exit from `SLLS_solve` or `SLLS_terminate` indicates that an error has occurred. No further calls should be made until the error has been corrected. Possible values are:

−1. An allocation error occurred. A message indicating the offending array is written on unit `control%error`, and the returned allocation status and a string containing the name of the offending array are held in `inform%alloc_status` and `inform%bad_alloc` respectively.

−2. A deallocation error occurred. A message indicating the offending array is written on unit `control%error` and the returned allocation status and a string containing the name of the offending array are held in `inform%alloc_status` and `inform%bad_alloc` respectively.

−3. One of the restrictions `prob%n > 0`, `prob%m > 0` or the requirement that `prob%A_type` contain its relevant string `'DENSE_BY_ROWS'`, `'DENSE_BY_COLUMNS'`, `'COORDINATE'`, `'SPARSE_BY_ROWS'` or `SPARSE_BY_COLUMNS`, when **A** is available, has been violated.

−4. The bound constraints are inconsistent.

−9. The analysis phase of the factorization failed; the return status from the factorization package is given in the component `inform%factor_status`.

−10. The factorization failed; the return status from the factorization package is given in the component `inform%factor_status`.

−18. Too many iterations have been performed. This may happen if `control%maxit` is too small, but may also be symptomatic of a badly scaled problem.

−19. The CPU time limit has been reached. This may happen if `control%cpu_time_limit` is too small, but may also be symptomatic of a badly scaled problem.

2.7 Further features

In this section, we describe an alternative means of setting control parameters, that is components of the variable `control` of type `SLLS_control_type` (see Section 2.2.3), by reading an appropriate data specification file using the subroutine `SLLS_read_specfile`. This facility is useful as it allows a user to change `SLLS` control parameters without editing and recompiling programs that call `SLLS`.

A specification file, or specfile, is a data file containing a number of "specification commands". Each command occurs on a separate line, and comprises a "keyword", which is a string (in a close-to-natural language) used to identify a control parameter, and an (optional) "value", which defines the value to be assigned to the given control parameter. All keywords and values are case insensitive, keywords may be preceded by one or more blanks but values must not contain blanks, and each value must be separated from its keyword by at least one blank. Values must not contain more than 30 characters, and each line of the specfile is limited to 80 characters, including the blanks separating keyword and value.

The portion of the specification file used by `SLLS_read_specfile` must start with a "BEGIN SLLS" command and end with an "END" command. The syntax of the specfile is thus defined as follows:

All use is subject to the conditions of a BSD-3-Clause License.
See `http://galahad.rl.ac.uk/galahad-www/cou.html` **for full details.**

```
( .. lines ignored by SLLS_read_specfile .. )
  BEGIN SLLS
     keyword    value
     .......    .....
     keyword    value
  END
( .. lines ignored by SLLS_read_specfile .. )
```

where keyword and value are two strings separated by (at least) one blank. The "`BEGIN SLLS`" and "`END`" delimiter command lines may contain additional (trailing) strings so long as such strings are separated by one or more blanks, so that lines such as

```
  BEGIN SLLS SPECIFICATION
```

and

```
  END SLLS SPECIFICATION
```

are acceptable. Furthermore, between the "`BEGIN SLLS`" and "`END`" delimiters, specification commands may occur in any order. Blank lines and lines whose first non-blank character is ! or * are ignored. The content of a line after a ! or * character is also ignored (as is the ! or * character itself). This provides an easy manner to "comment out" some specification commands, or to comment specific values of certain control parameters.

The value of a control parameters may be of three different types, namely integer, logical or real. Integer and real values may be expressed in any relevant Fortran integer and floating-point formats (respectively). Permitted values for logical parameters are "`ON`", "`TRUE`", "`.TRUE.`", "`T`", "`YES`", "`Y`", or "`OFF`", "`NO`", "`N`", "`FALSE`", "`.FALSE.`" and "`F`". Empty values are also allowed for logical control parameters, and are interpreted as "`TRUE`".

The specification file must be open for input when `SLLS_read_specfile` is called, and the associated device number passed to the routine in device (see below). Note that the corresponding file is `REWIND`ed, which makes it possible to combine the specifications for more than one program/routine. For the same reason, the file is not closed by `SLLS_read_specfile`.

2.7.1 To read control parameters from a specification file

Control parameters may be read from a file as follows:

```
    CALL SLLS_read_specfile( control, device )
```

`control` is a scalar `INTENT(INOUT)` argument of type `SLLS_control_type` (see Section 2.2.3). Default values should have already been set, perhaps by calling `SLLS_initialize`. On exit, individual components of `control` may have been changed according to the commands found in the specfile. Specfile commands and the component (see Section 2.2.3) of `control` that each affects are given in Table 2.1.

`device` is a scalar `INTENT(IN)` argument of type default `INTEGER`, that must be set to the unit number on which the specfile has been opened. If `device` is not open, `control` will not be altered and execution will continue, but an error message will be printed on unit `control%error`.

2.8 Information printed

If `control%print_level` is positive, information about the progress of the algorithm will be printed on unit `control%out`. If `control%print_level = 1`, a single line of output will be produced for each iteration of the process. This will give the total number of CG iterations performed (if any), the value of the objective function and the norm of the projected gradient, the stepsize taken, the number of free variables (i.e., those that have a positive value), the change in

All use is subject to the conditions of a BSD-3-Clause License.
See `http://galahad.rl.ac.uk/galahad-www/cou.html` **for full details.**

command	component of `control`	value type
`error-printout-device`	`%error`	integer
`printout-device`	`%out`	integer
`print-level`	`%print_level`	integer
`start-print`	`%start_print`	integer
`stop-print`	`%stop_print`	integer
`iterations-between-printing`	`%print_gap`	integer
`maximum-number-of-iterations`	`%maxit`	integer
`cold-start`	`%cold_start`	integer
`preconditioner`	`%preconditioner`	integer
`max-change-to-working-set-for-subspace-solution`	`%change_max`	integer
`maximum-number-of-cg-iterations-per-iteration`	`%cg_maxit`	integer
`maximum-number-of-arcsearch-steps`	`%arcsearch_max_steps`	integer
`primal-accuracy-required`	`%stop_p`	real
`dual-accuracy-required`	`%stop_d`	real
`complementary-slackness-accuracy-required`	`%stop_c`	real
`cg-relative-accuracy-required`	`%stop_cg_relative`	real
`cg-absolute-accuracy-required`	`%stop_cg_absolute`	real
`maximum-arcsearch-stepsize`	`%alpha_max`	real
`initial-arcsearch-stepsize`	`%alpha_initial`	real
`arcsearch-reduction-factor`	`%alpha_reduction`	real
`arcsearch-acceptance-tolerance`	`%arcsearch_acceptance_tol`	real
`regularization-weight`	`%regularization_weight`	real
`maximum-cpu-time-limit`	`%cpu_time_limit`	real
`direct-subproblem-solve`	`%direct_subproblem_solve`	logical
`exact-arc-search-used`	`%exact_arc_search`	logical
`space-critical`	`%space_critical`	logical
`deallocate-error-fatal`	`%deallocate_error_fatal`	logical
`output-line-prefix`	`%prefix`	character

Table 2.1: Specfile commands and associated components of `control`.

the number of free variables since the last iteration, and the elapsed CPU time in seconds. If `control%print_level` ≥ 2 this output will be increased to provide significant detail of each iteration. This extra output includes residuals of the linear systems solved, and, for larger values of `control%print_level`, values of the primal and dual variables and Lagrange multiplier.

3 GENERAL INFORMATION

Use of common: None.

Workspace: Provided automatically by the module.

Other routines called directly: None.

Other modules used directly: `SLLS_solve` calls the GALAHAD packages `GALAHAD_CPU_time`, `GALAHAD_SYMBOLS`, `GALAHAD_SPACE`, `GALAHAD_STRING`, `GALAHAD_SORT`, `GALAHAD_NORMS`, `GALAHAD_CONVERT`, `GALAHAD_SBLS`, `GALAHAD_QPT`, `GALAHAD_QPD`, `GALAHAD_USERDATA` and `GALAHAD_SPECFILE`.

Input/output: Output is under control of the arguments `control%error`, `control%out` and `control%print_level`.

Restrictions: `prob%n > 0`, `prob%m > 0`, `prob%A_type` ∈ {`'DENSE_BY_ROWS'`,`'DENSE_BY_COLUMNS'`,`'COORDINATE'`, `'SPARSE_BY_ROWS'`, `SPARSE_BY_COLUMNS` } (if **A** is explicit).

Portability: ISO Fortran 95 + TR 15581 or Fortran 2003. The package is thread-safe.

4   METHOD

The required solution **x** necessarily satisfies the primal optimality conditions

$$\mathbf{e}^T \mathbf{x} = 1 \text{ and } \mathbf{x} \geq 0, \tag{4.1}$$

the dual optimality conditions

$$\mathbf{A}^T \mathbf{W}(\mathbf{A}\mathbf{x} - \mathbf{b}) + \sigma x = \lambda \mathbf{e} + \mathbf{z}, \tag{4.2}$$

for some scalar Lagrange multiplier λ and dual variables

$$\mathbf{z} \geq 0, \tag{4.3}$$

and the complementary slackness conditions

$$\mathbf{x}^T \mathbf{z} = 0, \tag{4.4}$$

where the vector inequalities hold componentwise. Projected-gradient methods iterate towards a point that satisfies these conditions by ultimately aiming to satisfy (4.2), while ensuring that (4.1), and (4.3) and (4.4) are satisfied at each stage. Appropriate norms of the amounts by which (4.1), (4.2) and (4.4) fail to be satisfied are known as the primal and dual infeasibility, and the violation of complementary slackness, respectively.

The method is iterative. Each iteration proceeds in two stages. Firstly, a search direction **s** from the current estimate of the solution **x** is computed. This may be in a scaled steepest-descent direction, or, if the working set of variables on bounds has not changed dramatically, in a direction that provides an approximate minimizer of the objective over a subspace comprising the currently free-variables. The latter is computed either using an appropriate sparse factorization by the package `GALAHAD_SBLS`, or by the conjugate-gradient least-squares (CGLS) method; tt may be necessary to regularize the subproblem very slightly to avoid a ill-posedness. Thereafter, a piecewise linesearch (arc search) is carried out along the arc $\mathbf{x}(\alpha) = P(\mathbf{x} + \alpha\mathbf{s})$ for $\alpha > 0$, where the projection operator $P(\mathbf{v})$ gives the nearest point to **v** within the regular simplex; thus this arc bends the search direction into the feasible region. The arc search is performed either exactly, by passing through a set of increasing breakpoints at which it changes direction, or inexactly, by evaluating a sequence of different α on the arc. All computation is designed to exploit sparsity in **A**.

References:

Full details are provided in

N. I. M. Gould (2022). Linear least-squares over the unit simplex. In preparation.

5   EXAMPLE OF USE

Suppose we wish to minimize

$$\frac{1}{2} \left\| \begin{pmatrix} x_1 \\ x_1 + x_2 - 2 \\ x_3 - 1 \\ x_3 - 2 \end{pmatrix} \right\|_2^2$$

within the simplex

$$x_1 + x_2 + x_3 = 1, \ x_1, x_2, x_3 \geq 0.$$

Then, on writing the data for this problem as

$$\mathbf{A} = \begin{pmatrix} 1 & & \\ 1 & 1 & \\ & & 1 \\ & & 1 \end{pmatrix}, \ \mathbf{b} = \begin{pmatrix} 0 \\ 2 \\ 1 \\ 2 \end{pmatrix}, \ \mathbf{x}^l = \begin{pmatrix} -1 \\ -\infty \\ 0 \end{pmatrix} \text{ and } \mathbf{x}^u = \begin{pmatrix} \infty \\ 1 \\ 2 \end{pmatrix}$$

in sparse coordinate format, we may use the following code:

```
! THIS VERSION: GALAHAD 4.1 - 2022-05-26 AT 07:15 GMT
   PROGRAM GALAHAD_SLLS_EXAMPLE
   USE GALAHAD_SLLS_double          ! double precision version
   IMPLICIT NONE
   INTEGER, PARAMETER :: wp = KIND( 1.0D+0 ) ! set precision
   TYPE ( QPT_problem_type ) :: p
   TYPE ( SLLS_data_type ) :: data
   TYPE ( SLLS_control_type ) :: control
   TYPE ( SLLS_inform_type ) :: inform
   TYPE ( GALAHAD_userdata_type ) :: userdata
   INTEGER, ALLOCATABLE, DIMENSION( : ) :: X_stat
   INTEGER :: s
   INTEGER, PARAMETER :: n = 3, m = 4, a_ne = 5
! start problem data
   ALLOCATE( p%B( m ), p%X( n ), X_stat( n ) )
   p%n = n ; p%m = m                          ! dimensions
   p%B = (/ 0.0_wp, 2.0_wp, 1.0_wp, 2.0_wp /) ! right-hand side
   p%X = 0.0_wp ! start from zero
!  sparse co-ordinate storage format
   CALL SMT_put( p%A%type, 'COORDINATE', s )     ! Co-ordinate  storage for A
   ALLOCATE( p%A%val( a_ne ), p%A%row( a_ne ), p%A%col( A_ne ) )
   p%A%m = m ; p%A%n = n
   p%A%val = (/ 1.0_wp, 1.0_wp, 1.0_wp, 1.0_wp, 1.0_wp /) ! Jacobian A
   p%A%row = (/ 1, 2, 2, 3, 4 /)                     !
   p%A%col = (/ 1, 1, 2, 3, 3 /) ; p%A%ne = a_ne
! problem data complete
   CALL SLLS_initialize( data, control, inform ) ! Initialize control parameters
!  control%print_level = 1                       ! print one line/iteration
   control%exact_arc_search = .FALSE.
!  control%CONVERT_control%print_level = 3
   inform%status = 1
   CALL SLLS_solve( p, X_stat, data, control, inform, userdata )
   IF ( inform%status == 0 ) THEN              !  Successful return
     WRITE( 6, "( /, ' SLLS: ', I0, ' iterations ', /,                  &
   &     ' Optimal objective value =',                                  &
   &        ES12.4, /, ' Optimal solution = ', ( 5ES12.4 ) )" )         &
     inform%iter, inform%obj, p%X
   ELSE                                        ! Error returns
     WRITE( 6, "( /, ' SLLS_solve exit status = ', I0 ) " ) inform%status
     WRITE( 6, * ) inform%alloc_status, inform%bad_alloc
   END IF
   CALL SLLS_terminate( data, control, inform )  !  delete workspace
   DEALLOCATE( p%B, p%X, p%Z, X_stat )
   DEALLOCATE( p%A%val, p%A%row, p%A%col, p%A%type )
   END PROGRAM GALAHAD_SLLS_EXAMPLE
```

This produces the following output:

All use is subject to the conditions of a BSD-3-Clause License.
See `http://galahad.rl.ac.uk/galahad-www/cou.html` **for full details.**

```
 SLLS: 5 iterations
 Optimal objective value =  2.3333E+00
 Optimal solution =   3.3346E-13  3.3333E-01  6.6667E-01
```

The same problem may be solved holding the data in a sparse row-wise storage format by replacing the lines

```
!  sparse coordinate storage format
...
! problem data complete
```

by

```
! sparse row-wise storage format
   CALL SMT_put( p%A%type, 'SPARSE_BY_ROWS', s )  ! Specify sparse-by-rows
   ALLOCATE( p%A%val( a_ne ), p%A%col( a_ne ), p%A%ptr( m + 1 ) )
   p%A%val = (/ 1.0_wp, 1.0_wp, 1.0_wp, 1.0_wp, 1.0_wp /) ! Jacobian A
   p%A%col = (/ 1, 1, 2, 3, 3 /)                          ! Column indices
   p%A%ptr = (/ 1, 2, 4, 5, 6 /)                          ! Set row pointers
! problem data complete
```

a sparse column-wise storage format by replacing the lines

```
! sparse column-wise storage format
   CALL SMT_put( p%A%type, 'SPARSE_BY_COLUMNS', s )  ! Specify sparse-by-columns
   ALLOCATE( p%A%val( a_ne ), p%A%row( a_ne ), p%A%ptr( n + 1 ) )
   p%A%val = (/ 1.0_wp, 1.0_wp, 1.0_wp, 1.0_wp, 1.0_wp /) ! Jacobian A
   p%A%row = (/ 1, 2, 2, 3, 4 /)                          ! Row indices
   p%A%ptr = (/ 1, 3, 4, 6 /)                             ! Set column pointers
! problem data complete
```

a dense-by-rows storage format with the replacement lines

```
! dense storage format
   CALL SMT_put( p%A%type, 'DENSE_BY_ROWS', s )  ! Specify dense-by-rows
   ALLOCATE( p%A%val( m * n  )
   p%A%val = (/ 1.0_wp, 0.0_wp, 0.0_wp, 1.0_wp, 1.0_wp, 0.0_wp, 0.0_wp, 0.0_wp, &
               1.0_wp, 0.0_wp, 01.0_wp, 1.0_wp /)
! problem data complete
```

or a dense-by-columns storage format using the replacement

```
! dense storage format
   CALL SMT_put( p%A%type, 'DENSE_BY_COLUMS', s )  ! Specify dense-by-columns
   ALLOCATE( p%A%val( m * n  )
   p%A%val = (/ 1.0_wp, 1.0_wp, 0.0_wp, 0.0_wp, 0.0_wp, 1.0_wp, 0.0_wp, 0.0_wp, &
               0.0_wp, 0.0_wp, 1.0_wp, 1.0_wp /)
! problem data complete
```

respectively.

The same problem may be solved using reverse communication with the following code:

```
! THIS VERSION: GALAHAD 4.1 - 2022-07-06 AT 11:30 GMT.
   PROGRAM GALAHAD_SLLS_SECOND_EXAMPLE ! reverse commmunication interface
   USE GALAHAD_SLLS_double            ! double precision version
   IMPLICIT NONE
   INTEGER, PARAMETER :: wp = KIND( 1.0D+0 ) ! set precision
```

```
      TYPE ( QPT_problem_type ) :: p
      TYPE ( SLLS_data_type ) :: data
      TYPE ( SLLS_control_type ) :: control
      TYPE ( SLLS_inform_type ) :: inform
      TYPE ( SLLS_reverse_type ) :: reverse
      TYPE ( GALAHAD_userdata_type ) :: userdata
      INTEGER, ALLOCATABLE, DIMENSION( : ) :: X_stat
      INTEGER :: i, j, k, l, nflag
      REAL ( KIND = wp ) :: val
      INTEGER, PARAMETER :: n = 3, m = 4, a_ne = 5
      INTEGER, ALLOCATABLE, DIMENSION( : ) :: A_row, A_ptr, FLAG
      REAL ( KIND = wp ), ALLOCATABLE, DIMENSION( : ) :: A_val
! start problem data
      ALLOCATE( p%B( m ), p%X( n ), X_stat( n ) )
      p%n = n ; p%m = m                      ! dimensions
      p%B = (/ 0.0_wp, 2.0_wp, 1.0_wp, 2.0_wp /) ! right-hand side
      p%X = 0.0_wp ! start from zero
!  sparse column storage format
      ALLOCATE( A_val( a_ne ), A_row( a_ne ), A_ptr( n + 1 ) )
      A_val = (/ 1.0_wp, 1.0_wp, 1.0_wp, 1.0_wp, 1.0_wp /) ! Jacobian A by columns
      A_row = (/ 1, 2, 2, 3, 4 /)                ! row indices
      A_ptr = (/ 1, 3, 4, 6 /)                   ! pointers to column starts
! problem data complete
      CALL SLLS_initialize( data, control, inform ) ! Initialize control parameters
      control%print_level = 1                    ! print one line/iteration
      control%exact_arc_search = .FALSE.
      ALLOCATE( FLAG( n ) )
      nflag = 0 ; FLAG = 0  ! Flag if index already used in current (nflag) product
      inform%status = 1
10 CONTINUE ! Solve problem - reverse commmunication loop
        CALL SLLS_solve( p, X_stat, data, control, inform, userdata,        &
                       reverse = reverse )
      SELECT CASE ( inform%status )
      CASE ( 0 ) !  successful return
        WRITE( 6, "( /, ' SLLS: ', I0, ' iterations  ', /,                  &
       &    ' Optimal objective value =',                                   &
       &       ES12.4, /, ' Optimal solution = ', ( 5ES12.4 ) )" )          &
          inform%iter, inform%obj, p%X
      CASE ( 2 ) ! compute A * v
        reverse%P( : m ) = 0.0_wp
        DO j = 1, n
          val = reverse%V( j )
          DO k = A_ptr( j ), A_ptr( j + 1 ) - 1
            i = A_row( k )
            reverse%P( i ) = reverse%P( i ) + A_val( k ) * val
          END DO
        END DO
        GO TO 10
      CASE ( 3 ) ! compute A^T * v
        reverse%P( : n ) = 0.0_wp
        DO j = 1, n
          val = 0.0_wp
          DO k = A_ptr( j ), A_ptr( j + 1 ) - 1
            val = val + A_val( k ) * reverse%V( A_row( k ) )
          END DO
```

```
        reverse%P( j ) = val
      END DO
      GO TO 10
    CASE ( 4 ) ! compute A * sparse v
      reverse%P( : m ) = 0.0_wp
      DO l = reverse%nz_in_start, reverse%nz_in_end
        j = reverse%NZ_in( l )
        val = reverse%V( j )
        DO k = A_ptr( j ), A_ptr( j + 1 ) - 1
          i = A_row( k )
          reverse%P( i ) = reverse%P( i ) + A_val( k ) * val
        END DO
      END DO
      GO TO 10
    CASE ( 5 ) ! compute sparse( A * sparse v )
      nflag = nflag + 1
      reverse%nz_out_end = 0
      DO l = reverse%nz_in_start, reverse%nz_in_end
        j = reverse%NZ_in( l )
        val = reverse%V( j )
        DO k = A_ptr( j ), A_ptr( j + 1 ) - 1
          i = A_row( k )
          IF ( FLAG( i ) < nflag ) THEN
            FLAG( i ) = nflag
            reverse%P( i ) = A_val( k ) * val
            reverse%nz_out_end = reverse%nz_out_end + 1
            reverse%NZ_out( reverse%nz_out_end ) = i
          ELSE
            reverse%P( i ) = reverse%P( i ) + A_val( k ) * val
          END IF
        END DO
      END DO
      GO TO 10
    CASE ( 6 ) ! compute sparse( A^T * v )
      reverse%P( : n ) = 0.0_wp
      DO l = reverse%nz_in_start, reverse%nz_in_end
        j = reverse%NZ_in( l )
        val = 0.0_wp
        DO k = A_ptr( j ), A_ptr( j + 1 ) - 1
          val = val + A_val( k ) * reverse%V( A_row( k ) )
        END DO
        reverse%P( j ) = val
      END DO
      GO TO 10
    CASE DEFAULT ! error returns
      WRITE( 6, "( /, ' SLLS_solve exit status = ', I0 ) " ) inform%status
    END SELECT
  CALL SLLS_terminate( data, control, inform )  !  delete workspace
  DEALLOCATE( p%B, p%X, p%Z, X_stat, FLAG )
  DEALLOCATE( A_val, A_row, A_ptr )
  END PROGRAM GALAHAD_SLLS_SECOND_EXAMPLE
```

This produces the following output:

```
      S=steepest descent, F=factorization used
```

```
#its  #cg          f          proj gr     step   #free change   time
   0   -  2.55555555555556E+00 6.667E-01     -      -      -     0.0
   1   S  2.50000000000000E+00 6.667E-01 6.667E-01   1      -     0.0
   2   S  2.37500000000000E+00 1.667E-01 2.500E-01   2      1     0.0
   3   0  2.37500000000000E+00 1.667E-01 0.000E+00   2      0     0.0
   4   S  2.36718750000000E+00 2.500E-01 2.500E-01   3      1     0.0
   5   2  2.33333333333333E+00 1.332E-15 1.875E-01   2      1     0.0

 SLLS: 5 iterations
 Optimal objective value =  2.3333E+00
 Optimal solution =   4.7184E-16  3.3333E-01  6.6667E-01
```

The same problem may also be solved by user-provided matrix-vector products as follows:

```
! THIS VERSION: GALAHAD 4.1 - 2022-07-06 AT 16:20 GMT.
   PROGRAM GALAHAD_SLLS_THIRD_EXAMPLE ! subroutine evaluation interface
   USE GALAHAD_SLLS_double            ! double precision version
   IMPLICIT NONE
   INTEGER, PARAMETER :: wp = KIND( 1.0D+0 ) ! set precision
   TYPE ( QPT_problem_type ) :: p
   TYPE ( SLLS_data_type ) :: data
   TYPE ( SLLS_control_type ) :: control
   TYPE ( SLLS_inform_type ) :: inform
   TYPE ( GALAHAD_userdata_type ) :: userdata
   INTEGER, ALLOCATABLE, DIMENSION( : ) :: X_stat
   INTEGER, PARAMETER :: n = 3, m = 4, a_ne = 5
! partition userdata%integer so that it holds
!   m n nflag  flag      a_ptr          a_row
! |1|2|  3  |4 to n+3 |n+4 to 2n+4|2n+5 to 2n+4+a_ne|
! partition userdata%real so that it holds
!     a_val
! |1 to a_ne|
   INTEGER, PARAMETER :: mn = MAX( m, n )
   INTEGER, PARAMETER :: nflag = 3, st_flag = 3, st_ptr = st_flag + mn
   INTEGER, PARAMETER :: st_row = st_ptr + n + 1, st_val = 0
   INTEGER, PARAMETER :: len_integer = st_row + a_ne + 1, len_real = a_ne
   EXTERNAL :: APROD, ASPROD, AFPROD
! start problem data
   ALLOCATE( p%B( m ), p%X( n ), X_stat( n ) )
   p%n = n ; p%m = m                          ! dimensions
   p%B = (/ 0.0_wp, 2.0_wp, 1.0_wp, 2.0_wp /) ! right-hand side
   p%X = 0.0_wp ! start from zero
!  sparse co-ordinate storage format
   ALLOCATE( userdata%integer( len_integer ), userdata%real( len_real ) )
   userdata%integer( 1 ) = m                  ! load Jacobian data into userdata
   userdata%integer( 2 ) = n
   userdata%integer( st_ptr + 1 : st_ptr + n + 1 ) = (/ 1, 3, 4, 6 /)
   userdata%integer( st_row + 1 : st_row + a_ne ) = (/ 1, 2, 2, 3, 4 /)
   userdata%real( st_val + 1 : st_val + a_ne )                        &
     = (/ 1.0_wp, 1.0_wp, 1.0_wp, 1.0_wp, 1.0_wp /)
! problem data complete
   CALL SLLS_initialize( data, control, inform ) ! Initialize control parameters
   control%print_level = 1                     ! print one line/iteration
   control%exact_arc_search = .FALSE.
!  load workspace into userdata
   userdata%integer( nflag ) = 0
```

```
   userdata%integer( st_flag + 1 : st_flag + mn ) = 0
   inform%status = 1
   CALL SLLS_solve( p, X_stat, data, control, inform, userdata,            &
                    eval_APROD = APROD, eval_ASPROD = ASPROD,              &
                    eval_AFPROD = AFPROD )
   IF ( inform%status == 0 ) THEN              ! Successful return
     WRITE( 6, "( /, ' SLLS: ', I0, ' iterations  ', /,                    &
    &     ' Optimal objective value =',                                    &
    &        ES12.4, /, ' Optimal solution = ', ( 5ES12.4 ) )" )           &
     inform%iter, inform%obj, p%X
   ELSE                                        ! Error returns
     WRITE( 6, "( /, ' SLLS_solve exit status = ', I0 ) " ) inform%status
   END IF
   CALL SLLS_terminate( data, control, inform )  !  delete workspace
   DEALLOCATE( p%B, p%X, p%Z, X_stat )
   DEALLOCATE( userdata%integer, userdata%real )
   END PROGRAM GALAHAD_SLLS_THIRD_EXAMPLE

   SUBROUTINE APROD( status, userdata, transpose, V, P )
   USE GALAHAD_USERDATA_double
   INTEGER, PARAMETER :: wp = KIND( 1.0D+0 )
   INTEGER, INTENT( OUT ) :: status
   TYPE ( GALAHAD_userdata_type ), INTENT( INOUT ) :: userdata
   LOGICAL, INTENT( IN ) :: transpose
   REAL ( KIND = wp ), DIMENSION( : ), INTENT( IN ) :: V
   REAL ( KIND = wp ), DIMENSION( : ), INTENT( INOUT ) :: P
   INTEGER :: i, j, k
   REAL ( KIND = wp ) :: val
! recover problem data from userdata
   INTEGER :: m, n, nflag, st_flag, st_ptr, st_row, st_val
   m = userdata%integer( 1 )
   n = userdata%integer( 2 )
   nflag = 3
   st_flag = 3
   st_ptr = st_flag + MAX( m, n )
   st_row = st_ptr + n + 1
   st_val = 0
   IF ( transpose ) THEN
     DO j = 1, n
       DO k = userdata%integer( st_ptr + j ),                             &
              userdata%integer( st_ptr + j + 1 ) - 1
         P( j ) = P( j ) + userdata%real( st_val + k ) *                  &
                  V( userdata%integer( st_row + k ) )
       END DO
     END DO
   ELSE
     DO j = 1, n
       val = V( j )
       DO k = userdata%integer( st_ptr + j ),                             &
              userdata%integer( st_ptr + j + 1 ) - 1
         i = userdata%integer( st_row + k )
         P( i ) = P( i ) + userdata%real( st_val + k ) * val
       END DO
     END DO
   END IF
```

```
        status = 0
        RETURN
        END SUBROUTINE APROD

        SUBROUTINE ASPROD( status, userdata, V, P, NZ_in, nz_in_start, nz_in_end,   &
                           NZ_out, nz_out_end )
        USE GALAHAD_USERDATA_double
        INTEGER, PARAMETER :: wp = KIND( 1.0D+0 )
        INTEGER, INTENT( OUT ) :: status
        TYPE ( GALAHAD_userdata_type ), INTENT( INOUT ) :: userdata
        REAL ( KIND = wp ), DIMENSION( : ), INTENT( IN ) :: V
        REAL ( KIND = wp ), DIMENSION( : ), INTENT( OUT ) :: P
        INTEGER, OPTIONAL, INTENT( IN ) :: nz_in_start, nz_in_end
        INTEGER, OPTIONAL, INTENT( INOUT ) :: nz_out_end
        INTEGER, DIMENSION( : ), OPTIONAL, INTENT( IN ) :: NZ_in
        INTEGER, DIMENSION( : ), OPTIONAL, INTENT( INOUT ) :: NZ_out
        INTEGER :: i, j, k, l
        REAL ( KIND = wp ) :: val
!  recover problem data from userdata
        INTEGER :: m, n, nflag, st_flag, st_ptr, st_row, st_val
        IF ( PRESENT( NZ_in ) ) THEN
          IF ( .NOT. ( PRESENT( nz_in_start ) .AND. PRESENT( nz_in_end ) ) ) THEN
              status = - 1 ; RETURN
          END IF
        END IF
        m = userdata%integer( 1 )
        n = userdata%integer( 2 )
        nflag = 3
        st_flag = 3
        st_ptr = st_flag + MAX( m, n )
        st_row = st_ptr + n + 1
        st_val = 0
        IF ( PRESENT( NZ_in ) ) THEN
          IF ( PRESENT( NZ_out ) ) THEN
            IF ( .NOT. PRESENT( nz_out_end ) ) THEN
              status = - 1 ; RETURN
            END IF
            userdata%integer( nflag ) = userdata%integer( nflag ) + 1
            nz_out_end = 0
            DO l = nz_in_start, nz_in_end
              j = NZ_in( l )
              val = V( j )
              DO k = userdata%integer( st_ptr + j ),                           &
                     userdata%integer( st_ptr + j + 1 ) - 1
                i = userdata%integer( st_row + k )
                IF ( userdata%integer( st_flag + i ) <                         &
                     userdata%integer( nflag )  ) THEN
                  userdata%integer( st_flag + i ) = userdata%integer( nflag )
                  P( i ) = userdata%real( st_val + k ) * val
                  nz_out_end = nz_out_end + 1
                  NZ_out( nz_out_end ) = i
                ELSE
                  P( i ) = P( i ) + userdata%real( st_val + k ) * val
                END IF
              END DO
```

All use is subject to the conditions of a BSD-3-Clause License.
See `http://galahad.rl.ac.uk/galahad-www/cou.html` **for full details.**

```
        END DO
      ELSE
        P( : m ) = 0.0_wp
        DO l = nz_in_start, nz_in_end
          j = NZ_in( l )
          val = V( j )
          DO k = userdata%integer( st_ptr + j ),                      &
                 userdata%integer( st_ptr + j + 1 ) - 1
            i = userdata%integer( st_row + k )
            P( i ) = P( i ) + userdata%real( st_val + k ) * val
          END DO
        END DO
      END IF
  ELSE
    IF ( PRESENT( NZ_out ) ) THEN
      IF ( .NOT. PRESENT( nz_out_end ) ) THEN
        status = - 1 ; RETURN
      END IF
      userdata%integer( nflag ) = userdata%integer( nflag ) + 1
      nz_out_end = 0
      DO j = 1, n
        val = V( j )
        DO k = userdata%integer( st_ptr + j ),                      &
               userdata%integer( st_ptr + j + 1 ) - 1
          i = userdata%integer( st_row + k )
          IF ( userdata%integer( st_flag + i ) <                    &
               userdata%integer( nflag )  ) THEN
            userdata%integer( st_flag + i ) = userdata%integer( nflag )
            P( i ) = userdata%real( st_val + k ) * val
            nz_out_end = nz_out_end + 1
            NZ_out( nz_out_end ) = i
          ELSE
            P( i ) = P( i ) + userdata%real( st_val + k ) * val
          END IF
        END DO
      END DO
    ELSE
      P( : m ) = 0.0_wp
      DO j = 1, n
        val = V( j )
        DO k = userdata%integer( st_ptr + j ),                      &
               userdata%integer( st_ptr + j + 1 ) - 1
          i = userdata%integer( st_row + k )
          P( i ) = P( i ) + userdata%real( st_val + k ) * val
        END DO
      END DO
    END IF
  END IF
END IF
status = 0
RETURN
END SUBROUTINE ASPROD

SUBROUTINE AFPROD( status, userdata, transpose, V, P, FREE, n_free )
USE GALAHAD_USERDATA_double
INTEGER, PARAMETER :: wp = KIND( 1.0D+0 )
```

```
      INTEGER, INTENT( OUT ) :: status
      TYPE ( GALAHAD_userdata_type ), INTENT( INOUT ) :: userdata
      LOGICAL, INTENT( IN ) :: transpose
      INTEGER, INTENT( IN ) :: n_free
      INTEGER, INTENT( IN ), DIMENSION( : ) :: FREE
      REAL ( KIND = wp ), DIMENSION( : ), INTENT( IN ) :: V
      REAL ( KIND = wp ), DIMENSION( : ), INTENT( OUT ) :: P
      INTEGER :: i, j, k, l
      REAL ( KIND = wp ) :: val
! recover problem data from userdata
      INTEGER :: m, n, nflag, st_flag, st_ptr, st_row, st_val
      m = userdata%integer( 1 )
      n = userdata%integer( 2 )
      nflag = 3
      st_flag = 3
      st_ptr = st_flag + MAX( m, n )
      st_row = st_ptr + n + 1
      st_val = 0
      IF ( transpose ) THEN
        DO l = 1, n_free
          j = FREE( l )
          val = 0.0_wp
          DO k = userdata%integer( st_ptr + j ),                   &
                 userdata%integer( st_ptr + j + 1 ) - 1
            val = val + userdata%real( st_val + k ) *              &
                  V( userdata%integer( st_row + k ) )
          END DO
          P( j ) = val
        END DO
      ELSE
        P( : m ) = 0.0_wp
        DO l = 1, n_free
          j = FREE( l )
          val = V( j )
          DO k = userdata%integer( st_ptr + j ),                   &
                 userdata%integer( st_ptr + j + 1 ) - 1
            i = userdata%integer( st_row + k )
            P( i ) = P( i ) + userdata%real( st_val + k ) * val
          END DO
        END DO
      END IF
      status = 0
      RETURN
      END SUBROUTINE AFPROD
```

This produces the same output. Now notice how the matrix **A** is passed to the matrix-vector product evaluation routines via the `integer` and `real` components of the derived type `userdata`.