GALAHAD   TRS

1   SUMMARY

Given real n by n symmetric matrices $\mathbf{H}$ and $\mathbf{M}$ (with $\mathbf{M}$ diagonally dominant), another real m by n matrix $\mathbf{A}$, a real n vector $\mathbf{c}$ and scalars $\Delta > 0$ and f, this package finds a **global minimizer of the quadratic objective function** $\frac{1}{2}\mathbf{x}^T\mathbf{H}\mathbf{x} + \mathbf{c}^T\mathbf{x} + f$, **where the vector x is required to satisfy the constraint** $\|\mathbf{x}\|_{\mathbf{M}} \leq \Delta$ **and possibly** $\mathbf{A}\mathbf{x} = \mathbf{0}$, and where the $\mathbf{M}$-norm of $\mathbf{x}$ is $\|\mathbf{x}\|_{\mathbf{M}} = \sqrt{\mathbf{x}^T\mathbf{M}\mathbf{x}}$. This problem commonly occurs as a trust-region subproblem in nonlinear optimization calculations. The package may also be used to solve the related problem in which $\mathbf{x}$ is instead required to satisfy the **equality constraint** $\|\mathbf{x}\|_{\mathbf{M}} = \Delta$. The matrix $\mathbf{M}$ need not be provided in the commonly-occurring ℓ_2-trust-region case for which $\mathbf{M} = \mathbf{I}$, the n by n identity matrix.

Factorization of matrices of the form $\mathbf{H} + \lambda\mathbf{M}$—or

$$\begin{pmatrix} \mathbf{H} + \lambda\mathbf{M} & \mathbf{A}^T \\ \mathbf{A} & \mathbf{0} \end{pmatrix} \tag{1.1}$$

in cases where $\mathbf{A}\mathbf{x} = \mathbf{0}$ is imposed—for a succession of scalars λ will be required, so this package is most suited for the case where such a factorization may be found efficiently. If this is not the case, the package `GALAHAD_GLTR` may be preferred.

ATTRIBUTES — Versions: `GALAHAD_TRS_single`, `GALAHAD_TRS_double`. **Uses:** `GALAHAD_CLOCK`, `GALAHAD_SYM-BOLS`, `GALAHAD_SPACE`, `GALAHAD_RAND`, `GALAHAD_NORMS`, `GALAHAD_ROOTS`, `GALAHAD_SPECFILE`, `GALAHAD_SLS`, `GALA-HAD_IR`, `GALAHAD_MOP` **Date:** October 2008. **Origin:** N. I. M. Gould, H. S. Thorne, Rutherford Appleton Laboratory, and D. P. Robinson, Oxford University. **Language:** Fortran 95 + TR 15581 or Fortran 2003. **Parallelism:** Some options may use OpenMP and its runtime library.

2   HOW TO USE THE PACKAGE

Access to the package requires a `USE` statement such as

Single precision version

```
USE GALAHAD_TRS_single
```

Double precision version

```
USE GALAHAD_TRS_double
```

If it is required to use both modules at the same time, the derived types `SMT_TYPE`, `TRS_control_type`, `TRS_history_type`, `TRS_inform_type`, `TRS_data_type`, (Section 2.4) and the subroutines `TRS_initialize`, `TRS_solve`, `TRS_terminate` (Section 2.5) and `TRS_read_specfile` (Section 2.7) must be renamed on one of the `USE` statements.

2.1   Matrix storage formats

The matrices $\mathbf{H}$ and (if required) $\mathbf{M}$ and $\mathbf{A}$ may be stored in a variety of input formats.

2.1.1 Dense storage format

The matrix $\mathbf{H}$ is stored as a compact dense matrix by rows, that is, the values of the entries of each row in turn are stored in order within an appropriate real one-dimensional array. Since $\mathbf{H}$ is symmetric, only the lower triangular part (that is the part h_{ij} for $1 \leq j \leq i \leq n$) need be held. In this case the lower triangle should be stored by rows, that is component $i*(i-1)/2+j$ of the storage array `H%val` will hold the value h_{ij} (and, by symmetry, h_{ji}) for $1 \leq j \leq i \leq n$. The same is true for $\mathbf{M}$ if it is used. If $\mathbf{A}$ is used, the entire matrix must be supplied, and component $n*(i-1)+j$ of the storage array `A%val` will hold the value a_{ij} for $i = 1, \ldots, m$, $j = 1, \ldots, n$.

2.1.2 Sparse co-ordinate storage format

Only the nonzero entries of the matrices are stored. For the l-th entry of $\mathbf{H}$, $1 \leq j \leq i \leq n$, its row index i, column index j and value h_{ij} are stored in the l-th components of the integer arrays `H%row`, `H%col` and real array `H%val`, respectively. Note that only the entries in the lower triangle should be stored. The same scheme may be used for $\mathbf{M}$ if it is required. If $\mathbf{A}$ is used, the entire matrix must be supplied using the same scheme in the integer arrays `A%row`, `A%col` and real array `A%val`.

2.1.3 Sparse row-wise storage format

Again only the nonzero entries are stored, but this time they are ordered so that those in row i appear directly before those in row $i+1$. For the i-th row of $\mathbf{H}$, the i-th component of the integer array `H%ptr` holds the position of the first entry in this row, while `H%ptr` $(m+1)$ holds the total number of entries plus one. The column indices j, $1 \leq j \leq i$, and values h_{ij} of the entries in the i-th row are stored in components $l =$ `H%ptr`$(i), \ldots,$`H%ptr` $(i+1)-1$ of the integer array `H%col`, and real array `H%val`, respectively. Note that as before only the entries in the lower triangle should be stored. For sparse matrices, this scheme almost always requires less storage than its predecessor. This scheme may also be used for $\mathbf{M}$ and $\mathbf{A}$ if they are required, excepting that for $\mathbf{A}$ the whole matrix must be stored.

2.1.4 Diagonal storage format

If $\mathbf{H}$ is diagonal (i.e., $h_{ij} = 0$ for all $1 \leq i \neq j \leq n$) only the diagonals entries h_{ii}, $1 \leq i \leq n$, need be stored, and the first n components of the array `H%val` may be used for the purpose. The same applies to $\mathbf{M}$ if it is required. This scheme is inappropriate and thus unavailable for $\mathbf{A}$.

2.2 Integer kinds

We use the term long `INTEGER` to denote `INTEGER(kind=long)`, where `long = selected_int_kind(18))`.

2.3 Parallel usage

OpenMP may be used by the `GALAHAD_TRS` package to provide parallelism for some solvers in shared memory environments. See the documentation for the GALAHAD package `SLS` for more details. To run in parallel, OpenMP must be enabled at compilation time by using the correct compiler flag (usually some variant of `-openmp`). The number of threads may be controlled at runtime by setting the environment variable `OMP_NUM_THREADS`.

MPI may also be used by the package to provide parallelism for some solvers in a distributed memory environment. To use this form of parallelism, MPI must be enabled at runtime by using the correct compiler flag (usually some variant of `-lmpi`). Although the MPI process will be started automatically when required, it should be stopped by the calling program once no further use of this form of parallelism is needed. Typically, this will be via statements of the form

```
CALL MPI_INITIALIZED( flag, ierr )
IF ( flag ) CALL MPI_FINALIZE( ierr )
```

The code may be compiled and run in serial mode.

2.4 The derived data types

Six derived data types are accessible from the package.

2.4.1 The derived data type for holding matrices

The derived data type SMT_TYPE is used to hold the matrices **H** and perhaps **M** and/or **A**. The components of SMT_TYPE used here are:

m   is a scalar component of type default INTEGER, that holds the number of rows in the matrix.

n   is a scalar component of type default INTEGER, that holds the number of columns in the matrix.

ne   is a scalar variable of type default INTEGER, that holds the number of matrix entries.

type   is a rank-one allocatable array of type default CHARACTER, that is used to indicate the matrix storage scheme used. Its precise length and content depends on the type of matrix to be stored.

val   is a rank-one allocatable array of type default REAL (double precision in GALAHAD_TRS_double) and dimension at least ne, that holds the values of the entries. Each pair of off-diagonal entries $h_{ij} = h_{ji}$ of the *symmetric* matrix **H** is represented as a single entry (see §2.1.1–2.1.3). Any duplicated entries that appear in the sparse co-ordinate or row-wise schemes will be summed.

row   is a rank-one allocatable array of type default INTEGER, and dimension at least ne, that may hold the row indices of the entries. (see §2.1.2).

col   is a rank-one allocatable array of type default INTEGER, and dimension at least ne, that may hold the column indices of the entries (see §2.1.2–2.1.3).

ptr   is a rank-one allocatable array of type default INTEGER, and dimension at least n + 1, that may holds the pointers to the first entry in each row (see §2.1.3).

2.4.2 The derived data type for holding control parameters

The derived data type TRS_control_type is used to hold controlling data. Default values may be obtained by calling TRS_initialize (see Section 2.5.1). The components of TRS_control_type are:

error is a scalar variable of type default INTEGER, that holds the stream number for error messages. Printing of error messages in TRS_solve and TRS_terminate is suppressed if error ≤ 0. The default is error = 6.

out   is a scalar variable of type default INTEGER, that holds the stream number for informational messages. Printing of informational messages in TRS_solve is suppressed if out < 0. The default is out = 6.

print_level is a scalar variable of type default INTEGER, that is used to control the amount of informational output which is required. No informational output will occur if print_level ≤ 0. If print_level = 1 a single line of output will be produced for each iteration of the process. If print_level ≥ 2 this output will be increased to provide significant detail of each iteration. The default is print_level = 0.

dense_factorization is a scalar variable of type default INTEGER, that is used to indicate whether the problem should be treated as dense and solved using dense-factorization methods. Possible values are:

0   the problem should be considered as sparse.

    `1` the problem should be considered as dense.

    `other` the algorithm will decide whether to treat the problem as dense or sparse depending on its dimension and the sparsity of the matrices involved.

    The default is `dense_factorization = 0`.

`new_h` is a scalar variable of type default `INTEGER`, that is used to indicate how **H** has changed (if at all) since the previous call to `TRS_solve`. Possible values are:

    `0` **H** is unchanged.

    `1` the values in **H** have changed, but its nonzero structure is as before.

    `2` both the values and structure of **H** have changed.

    The default is `new_h = 2`.

`new_m` is a scalar variable of type default `INTEGER`, that is used to indicate how **M** (if required) has changed (if at all) since the previous call to `TRS_solve`. Possible values are:

    `0` **M** is unchanged.

    `1` the values in **M** have changed, but its nonzero structure is as before.

    `2` both the values and structure of **M** have changed.

    The default is `new_m = 2`.

`new_a` is a scalar variable of type default `INTEGER`, that is used to indicate how **A** (if required) has changed (if at all) since the previous call to `TRS_solve`. Possible values are:

    `0` **A** is unchanged.

    `1` the values in **A** have changed, but its nonzero structure is as before.

    `2` both the values and structure of **A** have changed.

    The default is `new_a = 2`.

`max_factorizations` is a scalar variable of type default `INTEGER`, that holds the maximum number of factorizations which will be permitted. If `max_factorizations` is set to a negative number, there will be no limit on the number of factorizations allowed. The default is `max_factorizations = -1`.

`inverse_itmax` is a scalar variable of type default `INTEGER`, that holds the maximum number of inverse iterations which will be allowed per step when estimating the leftmost eigenvalue in `TRS_solve`. If `inverse_itmax` is set to a non-positive number, it will be reset by `TRS_solve` to 2. The default is `inverse_itmax = 2`.

`taylor_max_degree` is a scalar variable of type default `INTEGER`, that specifies the maximum degree of Taylor approximant that will be used to approximate the secular function when trying to improve λ; a first-degree approximant results in Newton's method. The higher the degree, the better in general the improvement, but the larger the cost. Thus there is a balance between many cheap low-degree approximants and a few more expensive higher-degree ones. Our experience favours higher-degree approximants. The default is `taylor_max_degree = 3`, which is the highest degree currently supported.

`initial_multiplier` is a scalar variables of type `REAL` (double precision in `GALAHAD_TRS_double`), that should be set to an initial estimate of the required multiplier λ_* (see Section 4). The algorithm will only use this value if `%use_initial_multiplier` is set `.TRUE.` (see below), and otherwise will be reset by `TRS_solve`. A good initial estimate may sometimes dramatically improve the performance of the package. The default is `initial_multiplier = 0.0`.

---

lower is a scalar variables of type REAL (double precision in GALAHAD_TRS_double), that holds the value of any known lower bound on the required multiplier λ_*. A good lower bound may sometimes dramatically improve the performance of the package, but an incorrect value might cause the method to fail. Thus resetting lower from its default should be used with caution. The default is lower = - HUGE(1.0) (-HUGE(1.0D0) in GALA-HAD_TRS_double).

upper is a scalar variables of type REAL (double precision in GALAHAD_TRS_double), that holds the value of any known upper bound on the required multiplier λ_*. A good upper bound may sometimes dramatically improve the performance of the package, but an incorrect value might cause the method to fail. Thus resetting upper from its default should be used with caution. The default is upper = HUGE(1.0) (HUGE(1.0D0) in GALAH-AD_TRS_double).

stop_normal, stop_absolute_normal and stop_hard are scalar variables of type REAL (double precision in GA-LAHAD_TRS_double), that hold values for the standard convergence tolerances of the method (see Section 4). In particular, the method is deemed to have converged when the computed solution $\mathbf{x}$ and its multiplier λ satisfy either $\lambda = 0$ and $\|\|\mathbf{x}\|_{\mathbf{M}} < \Delta$ or $\|\mathbf{x}\|_{\mathbf{M}} - \Delta| \leq \max($ stop_normal $* \Delta$, stop_absolute_normal$)$ or $\lambda_{\mathrm{U}} - \lambda_{\mathrm{L}} \leq$ stop_hard $* \max(1, |\lambda_{\mathrm{L}}|, |\lambda_{\mathrm{U}}|)$, where λ_{L} and λ_{U} are computed lower and upper bounds on the optimal multiplier λ_*. The defaults are stop_normal = stop_absolute_normal = stop_hard = $u^{0.75}$, where u is EPSILON(1.0) (EPSILON(1.0D0) in GALAHAD_TRS_double).

start_invit_tol is a scalar variable of type REAL (double precision in GALAHAD_TRS_double), that holds the value of the starting tolerance for inverse iteration. Specifically, inverse iteration is started as soon as $\lambda_{\mathrm{U}} - \lambda_{\mathrm{L}} \leq$ start_invit_tol $* \max(|\lambda_{\mathrm{L}}|, |\lambda_{\mathrm{U}}|)$, where λ_{L} and λ_{U} are computed lower and upper bounds on the optimal multiplier λ_*. The default is start_invit_tol = 0.5.

start_invitmax_tol is a scalar variables of type REAL (double precision in GALAHAD_TRS_double), that holds the value of the starting tolerance for full inverse iteration. Specifically, Specifically, inverse_itmax steps of inverse iteration are started as soon as $\lambda_{\mathrm{U}} - \lambda_{\mathrm{L}} \leq$ start_invitmax_tol $* \max(|\lambda_{\mathrm{L}}|, |\lambda_{\mathrm{U}}|)$, where λ_{L} and λ_{U} are computed lower and upper bounds on the optimal multiplier λ_*. The default is start_invitmax_tol = 0.1.

equality_problem is a scalar variable of type default LOGICAL, that may be set .TRUE. if the user requires that the solution occur on the constraint boundary (i.e., that the inequality constraint be replaced by $\|\mathbf{x}\|_{\mathbf{M}} = \Delta$), and .FALSE. otherwise. The default is equality_problem = .FALSE..

use_initial_multiplier is a scalar variable of type default LOGICAL, that may be set .TRUE. if the user wishes to use the value of initial multiplier supplied in %initial_multiplier, and .FALSE. if the initial value will be chosen automatically. The default is use_initial_multiplier = .FALSE..

initialize_approx_eigenvector is a scalar variable of type default LOGICAL, that be should set .TRUE. if the user wishes the package to choose an initial estimate of the eigenvector corresponding to the leftmost eigenvalue of the matrix pencil $(\mathbf{H}, \mathbf{M})$ in the null-space of $\mathbf{A}$. If the eigenvector corresponding to the previous problem (if any) might be useful, initial_approx_eigenvector should be set .FALSE.. The default is initialize_approx_eigenvector = .TRUE..

force_Newton is a scalar variable of type default LOGICAL, that may be set .TRUE. if the user wishes the package to ignore the trust-region bound when $\mathbf{H}$ is positive definite, and thus to compute the Newton step in this case, and .FALSE. otherwise. The default is force_Newton = .FALSE..

space_critical is a scalar variable of type default LOGICAL, that may be set .TRUE. if the user wishes the package to allocate as little internal storage as possible, and .FALSE. otherwise. The package may be more efficient if space_critical is set .FALSE.. The default is space_critical = .FALSE..

deallocate_error_fatal is a scalar variable of type default LOGICAL, that may be set .TRUE. if the user wishes the package to return to the user in the unlikely event that an internal array deallocation fails, and .FALSE. if the package should be allowed to try to continue. The default is deallocate_error_fatal = .FALSE..

---

symmetric_linear_solver is a scalar variable of type default CHARACTER and length 30, that specifies the external package to be used to solve any symmetric linear system that might arise. Current possible choices are 'sils', 'ma27', 'ma57', 'ma77', 'ma86', 'ma97', ssids, 'pardiso' and 'wsmp', although only 'sils' and, for OMP 4.0-compliant compilers, 'ssids' are installed by default. See the documentation for the GALAHAD package SLS for further details. The default is symmetric_linear_solver = 'sils'.

definite_linear_solver is a scalar variable of type default CHARACTER and length 30, that specifies the external package to be used to solve any symmetric positive-definite linear system that might arise. Current possible choices are 'sils', 'ma27', 'ma57', 'ma77', 'ma86', 'ma87', 'ma97', ssids, 'pardiso' and 'wsmp', although only 'sils' and, for OMP 4.0-compliant compilers, 'ssids' are installed by default. See the documentation for the GALAHAD package SLS for further details. The default is definite_linear_solver = 'sils'.

prefix is a scalar variable of type default CHARACTER and length 30, that may be used to provide a user-selected character string to preface every line of printed output. Specifically, each line of output will be prefaced by the string prefix(2:LEN(TRIM(prefix))-1), thus ignoring the first and last non-null components of the supplied string. If the user does not want to preface lines by such a string, they may use the default prefix = "".

SLS_control is a scalar variable of type SLS_control_type that is used to control various aspects of the factorization package SLS. See the documentation for GALAHAD_SLS for more details.

IR_control is a scalar variable of type IR_control_type that is used to control various aspects of the iterative refinement package IR. See the documentation for GALAHAD_IR for more details.

2.4.3 The derived data type for holding history information

The derived data type TRS_history_type is used to hold the value of $\|\mathbf{x}(\lambda)\|_{\mathbf{M}}$, where $\mathbf{x}(\lambda)$ satisfies $(\mathbf{H}+\lambda\mathbf{M})\mathbf{x}(\lambda) = -\mathbf{c}$ and $\mathbf{A}\mathbf{x}(\lambda) = \mathbf{0}$ for a specific λ arising during the computation. The components of TRS_history_type are:

lambda is a scalar variable of type default REAL (double precision in GALAHAD_TRS_double), that gives the value λ.

x_norm is a scalar variable of type default REAL, that gives the corresponding value $\|\mathbf{x}(\lambda)\|_{\mathbf{M}}$.

2.4.4 The derived data type for holding timing information

The derived data type TRS_time_type is used to hold elapsed CPU and system clock times for the various parts of the calculation. The components of TRS_time_type are:

total is a scalar variable of type default REAL (double precision in GALAHAD_TRS_double), that gives the total CPU time spent in the package.

assemble is a scalar variable of type default REAL (double precision in GALAHAD_TRS_double), that gives the CPU time spent assembling the matrix (1.1) from its constituent parts.

analyse is a scalar variable of type default REAL (double precision in GALAHAD_TRS_double), that gives the CPU time spent analysing required matrices prior to factorization.

factorize is a scalar variable of type default REAL (double precision in GALAHAD_TRS_double), that gives the CPU time spent factorizing the required matrices.

solve is a scalar variable of type default REAL (double precision in GALAHAD_TRS_double), that gives the CPU time spent using the factors to solve relevant linear equations.

clock_total is a scalar variable of type default REAL (double precision in GALAHAD_TRS_double), that gives the total elapsed system clock time spent in the package.

---

clock_assemble is a scalar variable of type default `REAL` (double precision in `GALAHAD_TRS_double`), that gives the elapsed system clock time spent assembling the matrix (1.1) from its constituent parts.

clock_analyse is a scalar variable of type default `REAL` (double precision in `GALAHAD_TRS_double`), that gives the elapsed system clock time spent analysing required matrices prior to factorization.

clock_factorize is a scalar variable of type default `REAL` (double precision in `GALAHAD_TRS_double`), that gives the elapsed system clock time spent factorizing the required matrices.

clock_solve is a scalar variable of type default `REAL` (double precision in `GALAHAD_TRS_double`), that gives the elapsed system clock time spent using the factors to solve relevant linear equations.

2.4.5 The derived data type for holding informational parameters

The derived data type `TRS_inform_type` is used to hold parameters that give information about the progress and needs of the algorithm. The components of `TRS_inform_type` are:

status is a scalar variable of type default `INTEGER`, that gives the current status of the algorithm. See Section 2.6 for details.

alloc_status is a scalar variable of type default `INTEGER`, that gives the status of the last internal array allocation or deallocation. This will be 0 if status = 0.

bad_alloc is a scalar variable of type default `CHARACTER` and length 80, that gives the name of the last internal array for which there were allocation or deallocation errors. This will be the null string if status = 0.

factorizations is a scalar variable of type default `INTEGER`, that gives the number of factorizations of the matrix (1.1) for different λ, performed during the calculation.

max_entries_factors is a scalar variable of type default `INTEGER`, that gives the maximum number of entries in any of the matrix factorizations performed during the calculation.

len_history is a scalar variable of type default `INTEGER`, that gives the number of $(\lambda, \|\mathbf{x}(\lambda)\|_{\mathbf{M}})$ pairs encountered during the calculation.

obj is a scalar variable of type default `REAL` (double precision in `GALAHAD_TRS_double`), that holds the value of the objective function $\frac{1}{2}\mathbf{x}^T\mathbf{H}\mathbf{x} + \mathbf{c}^T\mathbf{x} + f$.

multiplier is a scalar variable of type default `REAL` (double precision in `GALAHAD_TRS_double`), that holds the value of the Lagrange multiplier λ associated with the constraint.

x_norm is a scalar variable of type default `REAL` (double precision in `GALAHAD_TRS_double`), that holds the value of $\|\mathbf{x}\|_{\mathbf{M}}$.

pole is a scalar variable of type default `REAL` (double precision in `GALAHAD_TRS_double`), that holds a lower bound on $\max(0, -\lambda_1)$, where λ_1 is the left-most eigenvalue of the matrix pencil (H, M).

hard_case is a scalar variable of type default `LOGICAL`, that will be `.TRUE.` if the "hard-case" has occurred (see Section 4) and `.FALSE.` otherwise.

time is a scalar variable of type `TRS_time_type` whose components are used to hold elapsed CPU and system clock times for the various parts of the calculation (see Section 2.4.4).

history is an array argument of dimension `len_history` and type `TRS_history_type` that contains a list of pairs $(\lambda, \|\mathbf{x}(\lambda)\|_{\mathbf{M}})$ encountered during the calculation (see Section 2.4.3).

All use is subject to the conditions of the GNU Lesser General Public License version 3.
See `http://galahad.rl.ac.uk/galahad-www/cou.html` **for full details.**

SLS_inform is a scalar variable of type SLS_inform_type, that holds informational parameters concerning the analysis, factorization and solution phases performed by the GALAHAD sparse matrix factorization package SLS. See the documentation for the package SLS for details of the derived type SLS_inform_type.

IR_inform is a scalar variable of type IR_inform_type, that holds informational parameters concerning the iterative refinement subroutine contained in the GALAHAD refinement package IR. See the documentation for the package IR for details of the derived type IR_inform_type.

2.4.6 The derived data type for holding problem data

The derived data type TRS_data_type is used to hold all the data for a particular problem between calls of TRS procedures. This data should be preserved, untouched, from the initial call to TRS_initialize to the final call to TRS_terminate.

2.5 Argument lists and calling sequences

There are three procedures for user calls (see Section 2.7 for further features):

1. The subroutine TRS_initialize is used to set default values and initialize private data.

2. The subroutine TRS_solve is called to solve the problem.

3. The subroutine TRS_terminate is provided to allow the user to automatically deallocate array components of the private data, allocated by TRS_solve, at the end of the solution process.

We use square brackets [ ] to indicate OPTIONAL arguments.

2.5.1 The initialization subroutine

Default values are provided as follows:

CALL TRS_initialize(data, control, inform)

data is a scalar INTENT(INOUT) argument of type TRS_data_type (see Section 2.4.6). It is used to hold data about the problem being solved.

control is a scalar INTENT(OUT) argument of type TRS_control_type (see Section 2.4.2). On exit, control contains default values for the components as described in Section 2.4.2. These values should only be changed after calling TRS_initialize.

inform is a scalar INTENT(OUT) argument of type TRS_inform_type (see Section 2.4.5). A successful call to TRS_initialize is indicated when the component status has the value 0. For other return values of status, see Section 2.6.

2.5.2 The optimization problem solution subroutine

The optimization problem solution algorithm is called as follows:

CALL TRS_solve(n, radius, f, C, H, X, data, control, inform[, M, A])

n is a scalar INTENT(IN) argument of type default INTEGER, that must be set to the number of unknowns, n. **Restriction:** $n > 0$.

radius is a scalar INTENT(IN) variable of type default REAL (double precision in GALAHAD_TRS_double), that must be set on initial entry to the value of the radius of the trust-region constraint, Δ. **Restriction:** $\Delta > 0$.

All use is subject to the conditions of the GNU Lesser General Public License version 3.
See http://galahad.rl.ac.uk/galahad-www/cou.html **for full details.**

f   is a scalar `INTENT(IN)` variable of type default `REAL` (double precision in `GALAHAD_TRS_double`), that holds the scalar value f for the objective function.

C   is an array `INTENT(IN)` argument of dimension n and type default `REAL` (double precision in `GALAHAD_TRS_double`), whose i-th entry holds the component c_i of the vector **c** for the objective function.

H   is scalar `INTENT(IN)` argument of type `SMT_TYPE` that holds the Hessian matrix **H**. The following components are used here:

H%type is an allocatable array of rank one and type default `CHARACTER`, that is used to indicate the storage scheme used. If the dense storage scheme (see Section 2.1.1) is used, the first five components of `H%type` must contain the string `DENSE`. For the sparse co-ordinate scheme (see Section 2.1.2), the first ten components of `H%type` must contain the string `COORDINATE`, for the sparse row-wise storage scheme (see Section 2.1.3), the first fourteen components of `H%type` must contain the string `SPARSE_BY_ROWS`, and for the diagonal storage scheme (see Section 2.1.4), the first eight components of `H%type` must contain the string `DIAGONAL`.

For convenience, the procedure `SMT_put` may be used to allocate sufficient space and insert the required keyword into `H%type`. For example, if we wish to store **M** using the co-ordinate scheme, we may simply

```
CALL SMT_put( H%type, 'COORDINATE' )
```

See the documentation for the GALAHAD package `SMT` for further details on the use of `SMT_put`.

H%ne is a scalar variable of type default `INTEGER`, that holds the number of entries in the **lower triangular** part of **H** in the sparse co-ordinate storage scheme (see Section 2.1.2). It need not be set for any of the other three schemes.

H%val is a rank-one allocatable array of type default `REAL` (double precision in `GALAHAD_TRS_double`), that holds the values of the entries of the **lower triangular** part of the Hessian matrix **H** in any of the storage schemes discussed in Section 2.1.

H%row is a rank-one allocatable array of type default `INTEGER`, that holds the row indices of the **lower triangular** part of **H** in the sparse co-ordinate storage scheme (see Section 2.1.2). It need not be allocated for any of the other three schemes.

H%col is a rank-one allocatable array variable of type default `INTEGER`, that holds the column indices of the **lower triangular** part of **H** in either the sparse co-ordinate (see Section 2.1.2), or the sparse row-wise (see Section 2.1.3) storage scheme. It need not be allocated when the dense or diagonal storage schemes are used.

H%ptr is a rank-one allocatable array of dimension n+1 and type default `INTEGER`, that holds the starting position of each row of the **lower triangular** part of **H**, as well as the total number of entries plus one, in the sparse row-wise storage scheme (see Section 2.1.3). It need not be allocated when the other schemes are used.

X   is an array `INTENT(OUT)` argument of dimension n and type default `REAL` (double precision in `GALAHAD_TRS_double`), that holds an estimate of the solution **x** of the problem on exit.

data is a scalar `INTENT(INOUT)` argument of type `TRS_data_type` (see Section 2.4.6). It is used to hold data about the problem being solved. It must not have been altered **by the user** since the last call to `TRS_initialize`.

control is a scalar `INTENT(IN)` argument of type `TRS_control_type`. (see Section 2.4.2). Default values may be assigned by calling `TRS_initialize` prior to the first call to `TRS_solve`.

inform is a scalar `INTENT(INOUT)` argument of type `TRS_inform_type` (see Section 2.4.5) whose components need not be set on entry. A successful call to `TRS_solve` is indicated when the component `status` has the value 0. For other return values of `status`, see Section 2.6.

M   is an `OPTIONAL` scalar `INTENT(IN)` argument of type `SMT_TYPE` that holds the diagonally dominant scaling
        matrix **M**. It need only be set if $\mathbf{M} \neq \mathbf{I}$ and in this case the following components are used:

   `M%type` is an allocatable array of rank one and type default `CHARACTER`, that is used to indicate the storage
        scheme used. If the dense storage scheme (see Section 2.1.1) is used, the first five components of `M%type`
        must contain the string `DENSE`. For the sparse co-ordinate scheme (see Section 2.1.2), the first ten com-
        ponents of `M%type` must contain the string `COORDINATE`, for the sparse row-wise storage scheme (see
        Section 2.1.3), the first fourteen components of `M%type` must contain the string `SPARSE_BY_ROWS`, and for
        the diagonal storage scheme (see Section 2.1.4), the first eight components of `M%type` must contain the
        string `DIAGONAL`.

        For convenience, the procedure `SMT_put` may be used to allocate sufficient space and insert the required
        keyword into `M%type`. For example, if we wish to store **M** using the co-ordinate scheme, we may simply

```
CALL SMT_put( M%type, 'COORDINATE' )
```

        See the documentation for the GALAHAD package `SMT` for further details on the use of `SMT_put`.

   `M%ne` is a scalar variable of type default `INTEGER`, that holds the number of entries in the **lower triangular** part
        of **M** in the sparse co-ordinate storage scheme (see Section 2.1.2). It need not be set for any of the other
        three schemes.

   `M%val` is a rank-one allocatable array of type default `REAL` (double precision in `GALAHAD_TRS_double`), that
        holds the values of the entries of the **lower triangular** part of the scaling matrix **M** in any of the storage
        schemes discussed in Section 2.1.

   `M%row` is a rank-one allocatable array of type default `INTEGER`, that holds the row indices of the **lower triangu-
        lar** part of **M** in the sparse co-ordinate storage scheme (see Section 2.1.2). It need not be allocated for any
        of the other three schemes.

   `M%col` is a rank-one allocatable array variable of type default `INTEGER`, that holds the column indices of the
        lower triangular part of **M** in either the sparse co-ordinate (see Section 2.1.2), or the sparse row-wise (see
        Section 2.1.3) storage scheme. It need not be allocated when the dense or diagonal storage schemes are
        used.

   `M%ptr` is a rank-one allocatable array of dimension `n+1` and type default `INTEGER`, that holds the starting posi-
        tion of each row of the **lower triangular** part of **M**, as well as the total number of entries plus one, in the
        sparse row-wise storage scheme (see Section 2.1.3). It need not be allocated when the other schemes are
        used.

        If `M` is absent, the ℓ_2-norm, $\|\mathbf{x}\|_2 = \sqrt{\mathbf{x}^T\mathbf{x}}$, will be employed.

A   is an `OPTIONAL` scalar `INTENT(IN)` argument of type `SMT_TYPE` that holds the constraint matrix **A**. It need only
        be set if the constraints $\mathbf{A}\mathbf{x} = \mathbf{0}$ are required, and in this case the following components are used:

   `A%type` is an allocatable array of rank one and type default `CHARACTER`, that is used to indicate the storage
        scheme used. If the dense storage scheme (see Section 2.1.1) is used, the first five components of `A%type`
        must contain the string `DENSE`. For the sparse co-ordinate scheme (see Section 2.1.2), the first ten com-
        ponents of `A%type` must contain the string `COORDINATE`, for the sparse row-wise storage scheme (see
        Section 2.1.3), the first fourteen components of `A%type` must contain the string `SPARSE_BY_ROWS`.

        For convenience, the procedure `SMT_put` may be used to allocate sufficient space and insert the required
        keyword into `A%type`. For example, if we wish to store **A** using the co-ordinate scheme, we may simply

```
CALL SMT_put( A%type, 'COORDINATE' )
```

        See the documentation for the GALAHAD package `SMT` for further details on the use of `SMT_put`.

   `A%m` is a scalar variable of type default `INTEGER`, that holds the number of rows of **A**.

All use is subject to the conditions of the GNU Lesser General Public License version 3.
See `http://galahad.rl.ac.uk/galahad-www/cou.html` **for full details.**

A%ne is a scalar variable of type default INTEGER, that holds the number of entries in **A** in the sparse co-ordinate storage scheme (see Section 2.1.2). It need not be set for either of the other schemes.

A%val is a rank-one allocatable array of type default REAL (double precision in GALAHAD_TRS_double), that holds the values of the entries of the constraint matrix **A** in any of the storage schemes discussed in Section 2.1.

A%row is a rank-one allocatable array of type default INTEGER, that holds the row indices of **A** in the sparse co-ordinate storage scheme (see Section 2.1.2). It need not be allocated for either of the other schemes.

A%col is a rank-one allocatable array variable of type default INTEGER, that holds the column indices of **A** in either the sparse co-ordinate (see Section 2.1.2), or the sparse row-wise (see Section 2.1.3) storage scheme. It need not be allocated when the dense storage scheme is used.

A%ptr is a rank-one allocatable array of dimension A%m+1 and type default INTEGER, that holds the starting position of each row of **A**, as well as the total number of entries plus one, in the sparse row-wise storage scheme (see Section 2.1.3). It need not be allocated when the other schemes are used.

2.5.3 The termination subroutine

All previously allocated arrays are deallocated as follows:

```
CALL TRS_terminate( data, control, inform )
```

data is a scalar INTENT(INOUT) argument of type TRS_data_type exactly as for TRS_solve that must not have been altered **by the user** since the last call to TRS_initialize. On exit, array components will have been deallocated.

control is a scalar INTENT(IN) argument of type TRS_control_type exactly as for TRS_solve.

inform is a scalar INTENT(OUT) argument of type TRS_inform_type exactly as for TRS_solve. Only the component status will be set on exit, and a successful call to TRS_terminate is indicated when this component status has the value 0. For other return values of status, see Section 2.6.

2.6 Warning and error messages

A negative value of inform%status on exit from TRS_solve or TRS_terminate indicates that an error might have occurred. No further calls should be made until the error has been corrected. Possible values are:

-1. An allocation error occurred. A message indicating the offending array is written on unit control%error, and the returned allocation status and a string containing the name of the offending array are held in inform%alloc_status and inform%bad_alloc, respectively.

-2. A deallocation error occurred. A message indicating the offending array is written on unit control%error and the returned allocation status and a string containing the name of the offending array are held in inform%alloc_status and inform%bad_alloc, respectively.

-3. (TRS_solve only) One of the restrictions n > 0 or radius > 0 has been violated.

-9. (TRS_solve only) The analysis phase of the factorization of the matrix (1.1) failed.

-10. (TRS_solve only) The factorization of the matrix (1.1) failed.

-15. (TRS_solve only) The matrix **M** appears not to be diagonally dominant.

-16. (TRS_solve only) The problem is so ill-conditioned that further progress is impossible.

-18. (TRS_solve only) Too many factorizations have been required. This may happen if control%max_factorizations is too small, but may also be symptomatic of a badly scaled problem.

All use is subject to the conditions of the GNU Lesser General Public License version 3.
See http://galahad.rl.ac.uk/galahad-www/cou.html **for full details.**

2.7 Further features

In this section, we describe an alternative means of setting control parameters, that is components of the variable `control` of type `TRS_control_type` (see Section 2.4.2), by reading an appropriate data specification file using the subroutine `TRS_read_specfile`. This facility is useful as it allows a user to change `TRS` control parameters without editing and recompiling programs that call `TRS`.

A specification file, or specfile, is a data file containing a number of "specification commands". Each command occurs on a separate line, and comprises a "keyword", which is a string (in a close-to-natural language) used to identify a control parameter, and an (optional) "value", which defines the value to be assigned to the given control parameter. All keywords and values are case insensitive, keywords may be preceded by one or more blanks but values must not contain blanks, and each value must be separated from its keyword by at least one blank. Values must not contain more than 30 characters, and each line of the specfile is limited to 80 characters, including the blanks separating keyword and value.

The portion of the specification file used by `TRS_read_specfile` must start with a "BEGIN TRS" command and end with an "END" command. The syntax of the specfile is thus defined as follows:

```
( .. lines ignored by TRS_read_specfile .. )
  BEGIN TRS
     keyword    value
     .......    .....
     keyword    value
  END
( .. lines ignored by TRS_read_specfile .. )
```

where keyword and value are two strings separated by (at least) one blank. The "BEGIN TRS" and "END" delimiter command lines may contain additional (trailing) strings so long as such strings are separated by one or more blanks, so that lines such as

```
  BEGIN TRS SPECIFICATION
```

and

```
  END TRS SPECIFICATION
```

are acceptable. Furthermore, between the "BEGIN TRS" and "END" delimiters, specification commands may occur in any order. Blank lines and lines whose first non-blank character is ! or * are ignored. The content of a line after a ! or * character is also ignored (as is the ! or * character itself). This provides an easy manner to "comment out" some specification commands, or to comment specific values of certain control parameters.

The value of a control parameters may be of three different types, namely integer, logical or real. Integer and real values may be expressed in any relevant Fortran integer and floating-point formats (respectively). Permitted values for logical parameters are "ON", "TRUE", ".TRUE.", "T", "YES", "Y", or "OFF", "NO", "N", "FALSE", ".FALSE." and "F". Empty values are also allowed for logical control parameters, and are interpreted as "TRUE".

The specification file must be open for input when `TRS_read_specfile` is called, and the associated device number passed to the routine in device (see below). Note that the corresponding file is REWINDed, which makes it possible to combine the specifications for more than one program/routine. For the same reason, the file is not closed by `TRS_read_specfile`.

Control parameters corresponding to the components `SLS_control` and `IR_control` may be changed by including additional sections enclosed by "BEGIN SLS" and "END SLS", and "BEGIN IR" and "END IR", respectively. See the specification sheets for the packages `GALAHAD_SLS` and `GALAHAD_IR` for further details.

2.7.1 To read control parameters from a specification file

Control parameters may be read from a file as follows:

```
CALL TRS_read_specfile( control, device )
```

`control` is a scalar `INTENT(INOUT)` argument of type `TRS_control_type` (see Section 2.4.2). Default values should have already been set, perhaps by calling `TRS_initialize`. On exit, individual components of `control` may have been changed according to the commands found in the specfile. Specfile commands and the component (see Section 2.4.2) of `control` that each affects are given in Table 2.1.

command	component of control	value type
error-printout-device	%error	integer
printout-device	%out	integer
print-level	%print_level	integer
use-dense-factorization	%dense_factorization	integer
has-h-changed	%new_h	integer
has-m-changed	%new_m	integer
has-a-changed	%new_a	integer
factorization-limit	%max_factorizations	integer
inverse-iteration-limit	%inverse_itmax	integer
max-degree-taylor-approximant	%taylor_max_degree	integer
initial-multiplier	%initial_multiplier	real
lower-bound-on-multiplier	%lower	real
upper-bound-on-multiplier	%upper	real
stop-normal-case	%stop_normal	real
stop-absolute-normal-case	%stop_absolute_normal	real
stop-hard-case	%stop_hard	real
start-inverse-iteration-tolerance	%start_invit_tol	real
start-max-inverse-iteration-tolerance	%start_invitmax_tol	real
use-initial-multiplier	%use_initial_multiplier	logical
initialize-approximate-eigenvector	%initialize_approx_eigenvector	real
force-Newton-if-positive-definite	%force_Newton	logical
space-critical	%space_critical	logical
deallocate-error-fatal	%deallocate_error_fatal	logical
definite-linear-equation-solver	%definite_linear_solver	character
output-line-prefix	%prefix	character

Table 2.1: Specfile commands and associated components of `control`.

`device` is a scalar `INTENT(IN)` argument of type default `INTEGER`, that must be set to the unit number on which the specfile has been opened. If `device` is not open, `control` will not be altered and execution will continue, but an error message will be printed on unit `control%error`.

2.8 Information printed

If `control%print_level` is positive, information about the progress of the algorithm will be printed on unit `control%out`. If `control%print_level = 1`, a single line of output will be produced for each iteration of the process. In the first phase of the algorithm, this will include the current estimate of the multiplier and known brackets on its optimal value. In the second phase, the residual $\|\mathbf{x}\|_{\mathbf{M}} - \Delta$, the current estimate of the multiplier and the size of the correction will be printed. If `control%print_level` ≥ 2, this output will be increased to provide significant detail of each iteration. This extra output includes times for various phases.

3 GENERAL INFORMATION

Use of common: None.

Workspace: Provided automatically by the module.

Other routines called directly: None.

Other modules used directly: `TRS_solve` calls the GALAHAD packages `GALAHAD_CLOCK`, `GALAHAD_SYMBOLS`, `GALA-HAD_SPACE`, `GALAHAD_RAND`, `GALAHAD_NORMS`, `GALAHAD_ROOTS`, `GALAHAD_SPECFILE`, `GALAHAD_SLS`, `GALAHAD_IR` and `GALAHAD_MOP`.

Input/output: Output is under control of the arguments `control%error`, `control%out` and `control%print_level`.

Restrictions: $n > 0$, $\Delta > 0$.

Portability: ISO Fortran 95 + TR 15581 or Fortran 2003. The package is thread-safe.

4 METHOD

The required solution $\mathbf{x}_*$ necessarily satisfies the optimality condition $\mathbf{H}\mathbf{x}_* + \lambda_*\mathbf{M}\mathbf{x}_* + \mathbf{A}^T\mathbf{y}_* + \mathbf{c} = \mathbf{0}$ and $\mathbf{A}\mathbf{x}_* = 0$, where $\lambda_* \geq 0$ is a Lagrange multiplier corresponding to the constraint $\|\mathbf{x}\|_{\mathbf{M}} \leq \Delta$ and $\mathbf{y}_*$ are Lagrange multipliers for the linear constraints $\mathbf{A}\mathbf{x} = 0$, if any; for the equality-constrained problem $\|\mathbf{x}\|_{\mathbf{M}} = \Delta$, the multiplier is unconstrained. In addition in all cases, the matrix $\mathbf{H} + \lambda_*\mathbf{M}$ will be positive semi-definite on the null-space of $\mathbf{A}$; in most instances it will actually be positive definite, but in special "hard" cases singularity is a possibility.

The method is iterative, and proceeds in two phases. Firstly, lower and upper bounds, λ_{L} and λ_{U}, on λ_* are computed using Gershgorin's theorems and other eigenvalue bounds. The first phase of the computation proceeds by progressively shrinking the bound interval $[\lambda_{\mathrm{L}}, \lambda_{\mathrm{U}}]$ until a value λ for which $\|\mathbf{x}(\lambda)\|_{\mathbf{M}} \geq \Delta$ is found. Here $\mathbf{x}(\lambda)$ and its companion $\mathbf{y}(\lambda)$ are defined to be a solution of

$$(\mathbf{H} + \lambda\mathbf{M})\mathbf{x}(\lambda) + \mathbf{A}^T\mathbf{y}(\lambda) = -\mathbf{c} \ \text{ and } \ \mathbf{A}\mathbf{x}(\lambda) = \mathbf{0}; \tag{4.1}$$

along the way the possibility that $\mathbf{H}$ might be positive definite on the null-space of $\mathbf{A}$ and $\|\mathbf{x}(0)\|_{\mathbf{M}} \leq \Delta$ is examined, and if this transpires the process is terminated with $\mathbf{x}_* = \mathbf{x}(0)$. Once the terminating λ from the first phase has been discovered, the second phase consists of applying Newton or higher-order iterations to the nonlinear "secular" equation $\|\mathbf{x}(\lambda)\|_{\mathbf{M}} = \Delta$ with the knowledge that such iterations are both globally and ultimately rapidly convergent. It is possible in the "hard" case that the interval in the first-phase will shrink to the single point λ_*, and precautions are taken, using inverse iteration with Rayleigh-quotient acceleration to ensure that this too happens rapidly.

The dominant cost is the requirement that we solve a sequence of linear systems (4.1). In the absence of linear constraints, an efficient sparse Cholesky factorization with precautions to detect indefinite $\mathbf{H} + \lambda\mathbf{M}$ is used. If $\mathbf{A}\mathbf{x} = 0$ is required, a sparse symmetric, indefinite factorization of (1.1) is used rather than a Cholesky factorization.

Reference: The method is described in detail in

H. S. Dollar, N. I. M. Gould and D. P. Robinson. On solving trust-region and other regularised subproblems in optimization. *Mathematical Programming Computation* **2(1)** (2010) 21–57.

All use is subject to the conditions of the GNU Lesser General Public License version 3.
See `http://galahad.rl.ac.uk/galahad-www/cou.html` **for full details.**

5 EXAMPLE OF USE

Suppose we wish to solve a problem in 10,000 unknowns, whose data is

$$
\mathbf{H} = \begin{pmatrix} -2 & 1 & & & \\ 1 & -2 & . & & \\ & . & . & . & \\ & & . & -2 & 1 \\ & & & 1 & -2 \end{pmatrix}, \; \mathbf{M} = \begin{pmatrix} 2 & & & & \\ & 2 & & & \\ & & . & & \\ & & & 2 & \\ & & & & 2 \end{pmatrix}, \; \mathbf{c} = \begin{pmatrix} 1 \\ 1 \\ . \\ 1 \\ 1 \end{pmatrix} \text{ and } f = 1,
$$

with a radius $\Delta = 10$ but no other constraints. Then we may use the following code:

```
PROGRAM GALAHAD_TRS_EXAMPLE  !  GALAHAD 2.4 - 14/05/2010 AT 14:30 GMT.
USE GALAHAD_TRS_DOUBLE                     ! double precision version
IMPLICIT NONE
INTEGER, PARAMETER :: wp = KIND( 1.0D+0 )   ! set precision
INTEGER, PARAMETER :: n = 10000             ! problem dimension
REAL ( KIND = wp ), DIMENSION( n ) :: C, X
TYPE ( SMT_type ) :: H, M
TYPE ( TRS_data_type ) :: data
TYPE ( TRS_control_type ) :: control
TYPE ( TRS_inform_type ) :: inform
REAL ( KIND = wp ) :: f = 1.0_wp           ! constant term, f
REAL ( KIND = wp ) :: radius = 10.0_wp     ! trust-region radius
INTEGER :: i, s
C = 1.0_wp
CALL SMT_put( H%type, 'COORDINATE', s )     ! Specify co-ordinate for H
H%ne = 2 * n - 1
ALLOCATE( H%val( H%ne ), H%row( H%ne ), H%col( H%ne ) )
DO i = 1, n
 H%row( i ) = i ; H%col( i ) = i ; H%val( i ) = - 2.0_wp
END DO
DO i = 1, n - 1
 H%row( n + i ) = i + 1 ; H%col( n + i ) = i ; H%val( n + i ) = 1.0_wp
END DO
CALL SMT_put( M%type, 'DIAGONAL', s )       ! Specify diagonal for M
ALLOCATE( M%val( n ) ) ; M%val = 2.0_wp
CALL TRS_initialize( data, control, inform ) ! Initialize control parameters
! control%print_level = 1
CALL TRS_solve( n, radius, f, C, H, X, data, control, inform, M = M ) ! Solve
IF ( inform%status == 0 ) THEN !  Successful return
 WRITE( 6, "( 1X, I0, ' factorizations. Objective and Lagrange multiplier', &
&    ' =', 2ES12.4 )" ) inform%factorizations, inform%obj, inform%multiplier
ELSE  !  Error returns
 WRITE( 6, "( ' TRS_solve exit status = ', I0 ) " ) inform%status
END IF
CALL TRS_terminate( data, control, inform )  ! delete internal workspace
DEALLOCATE( H%row, H%col, H%val, M%val )
END PROGRAM GALAHAD_TRS_EXAMPLE
```

This produces the following output:

```
 4 factorizations. Objective and Lagrange multiplier = -7.0611E+02  7.0712E+00
```

If we now add the constraint $\sum\limits_{i=1}^{10000} x_i = 0$, for which $\mathbf{A} = (1 \; \dots \; 1)$, but revert to unit ($\mathbf{M} = \mathbf{I}$) trust-region, we may solve the resulting problem using the following code:

```
     PROGRAM GALAHAD_TRS_EXAMPLE2  !  GALAHAD 2.3 - 29/01/2009 AT 10:30 GMT.
     USE GALAHAD_TRS_DOUBLE                      ! double precision version
     IMPLICIT NONE
     INTEGER, PARAMETER :: wp = KIND( 1.0D+0 )    ! set precision
!    INTEGER, PARAMETER :: n = 10000              ! problem dimension
     INTEGER, PARAMETER :: n = 10                ! problem dimension
     REAL ( KIND = wp ), DIMENSION( n ) :: C, X
     TYPE ( SMT_type ) :: H, A
     TYPE ( TRS_data_type ) :: data
     TYPE ( TRS_control_type ) :: control
     TYPE ( TRS_inform_type ) :: inform
     REAL ( KIND = wp ) :: f = 1.0_wp            ! constant term, f
     REAL ( KIND = wp ) :: radius = 10.0_wp      ! trust-region radius
     INTEGER :: i, s
     C = 1.0_wp
     CALL SMT_put( H%type, 'COORDINATE', s )     ! Specify co-ordinate for H
     H%ne = 2 * n - 1
     ALLOCATE( H%val( H%ne ), H%row( H%ne ), H%col( H%ne ) )
     DO i = 1, n
      H%row( i ) = i ; H%col( i ) = i ; H%val( i ) = - 2.0_wp
     END DO
     DO i = 1, n - 1
      H%row( n + i ) = i + 1 ; H%col( n + i ) = i ; H%val( n + i ) = 1.0_wp
     END DO
     CALL SMT_put( A%type, 'DENSE', s )          ! Specify 1 by n matrix A
     ALLOCATE( A%val( n ) ) ; A%val = 1.0_wp ; A%m = 1 ; A%n = n
     DO i = 1, n
      A%val( i ) = REAL( i, KIND = wp )
     END DO
     CALL TRS_initialize( data, control, inform )  ! Initialize control parameters
     CALL TRS_solve( n, radius, f, C, H, X, data, control, inform, A = A ) ! Solve
     IF ( inform%status == 0 ) THEN !  Successful return
      WRITE( 6, "( 1X, I0, ' factorizations. Objective and Lagrange multiplier', &
     &   ' =', 2ES12.4 )" ) inform%factorizations, inform%obj, inform%multiplier
     ELSE  !  Error returns
      WRITE( 6, "( ' TRS_solve exit status = ', I0 ) " ) inform%status
     END IF
     CALL TRS_terminate( data, control, inform )  ! delete internal workspace
     DEALLOCATE( H%row, H%col, H%val, A%val )
     END PROGRAM GALAHAD_TRS_EXAMPLE2
```

This produces the following output:

```
 9 factorizations. Objective and Lagrange multiplier = -1.9570E+02  3.9226E+00
```